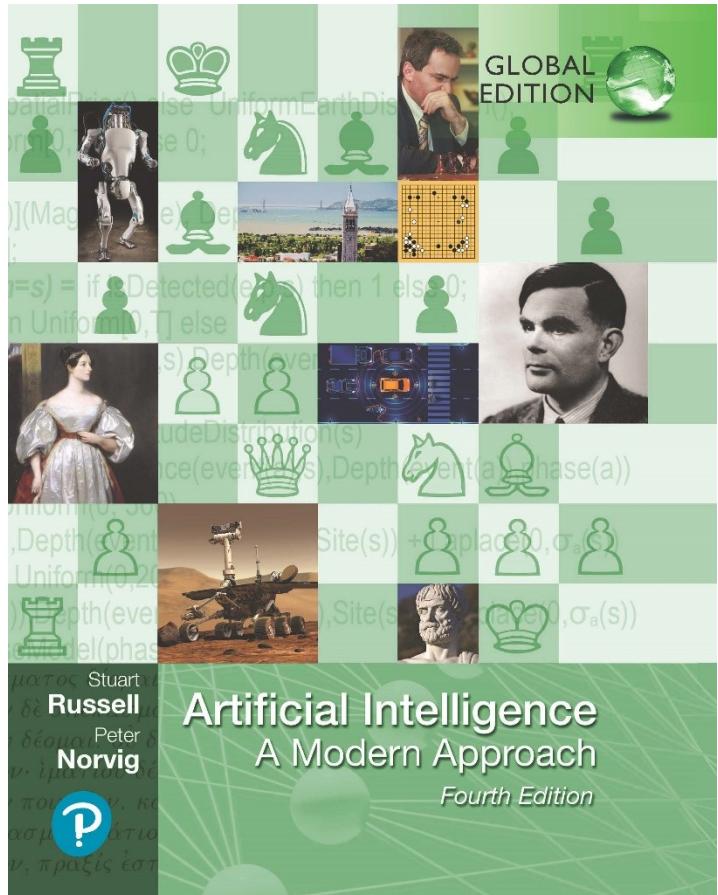


# Artificial Intelligence: A Modern Approach

Fourth Edition, Global Edition



## Chapter 1

### Introduction



# Lecture Presentations: Artificial Intelligence

**Adapted from:**

"Artificial Intelligence: A Modern Approach, Global Edition",  
4th Edition by Stuart Russell and Peter Norvig © 2021  
Pearson Education.

**Adapted for** educational use at ACE Engineering College.  
Some slides customized by Mr. Shafakhatullah Khan  
Mohammed, Assistant Professor @ ACE Engineering College.  
For instructional use only. Not for commercial distribution.

# Outline

- ◆ What is AI?
- ◆ A brief history
- ◆ The state of the art

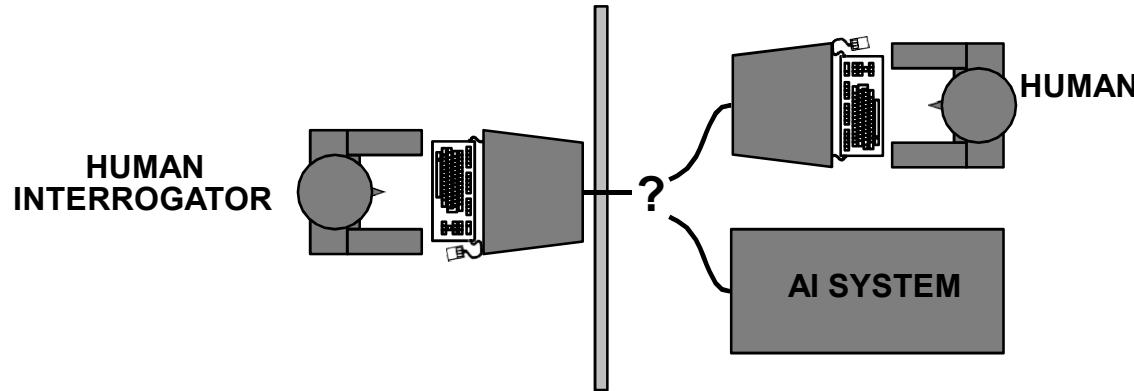
## What is AI?

Systems that think like humans	Systems that think rationally
Systems that act like humans	Systems that act rationally

## Acting humanly: The Turing test

Turing (1950) “Computing machinery and intelligence”:

- ◆ “**Can machines think?**” → “**Can machines behave intelligently?**”
- ◆ Operational test for intelligent behavior: the **Imitation Game**



- ◆ Predicted that by 2000, a machine might have a 30% chance of fooling a lay person for 5 minutes
- ◆ Anticipated all major arguments against AI in following 50 years
- ◆ Suggested major components of AI: knowledge, reasoning, language understanding, learning

Problem: Turing test is not **reproducible**, **constructive**, or amenable to **mathematical analysis**

## Thinking humanly: Cognitive Science

1960s “cognitive revolution”: information-processing psychology replaced prevailing orthodoxy of behaviorism

Requires scientific theories of internal activities of the brain

- What level of abstraction? “Knowledge” or “circuits”?
- How to validate? Requires
  - 1) Predicting and testing behavior of human subjects (top-down)
  - or 2) Direct identification from neurological data (bottom-up)

Both approaches (roughly, Cognitive Science and Cognitive Neuroscience) are now distinct from AI

Both share with AI the following characteristic:

the available theories do not explain (or engender) anything resembling human-level general intelligence

Hence, all three fields share one principal direction!

## Thinking rationally: Laws of Thought

Normative (or prescriptive) rather than descriptive

Aristotle: what are correct arguments/thought processes?

Several Greek schools developed various forms of logic:

notation and rules of derivation for thoughts;  
may or may not have proceeded to the idea of mechanization

Direct line through mathematics and philosophy to modern AI

Problems:

- 1) Not all intelligent behavior is mediated by logical deliberation
- 2) What is the purpose of thinking? What thoughts should I have out of all the thoughts (logical or otherwise) that I could have?

## Acting rationally

Rational behavior: doing the right thing

The right thing: that which is expected to maximize goal achievement, given the available information

Doesn't necessarily involve thinking—e.g., blinking reflex—but thinking should be in the service of rational action

Aristotle (Nicomachean Ethics):

Every art and every inquiry, and similarly every action and pursuit, is thought to aim at some good

## Rational agents

An **agent** is an entity that perceives and acts

This course is about designing **rational agents**

Abstractly, an agent is a function from percept histories to actions:

$$f : P^* \rightarrow A$$

For any given class of environments and tasks, we seek the agent (or class of agents) with the best performance

Caveat: **computational limitations make perfect rationality unachievable**

→ design best **program** for given machine resources

# AI prehistory

Philosophy	logic, methods of reasoning mind as physical system foundations of learning, language, rationality
Mathematics	formal representation and proof algorithms, computation, (un)decidability, (in)tractability probability
Psychology	adaptation phenomena of perception and motor control experimental techniques (psychophysics, etc.)
Economics	formal theory of rational decisions
Linguistics	knowledge representation grammar
Neuroscience	plastic physical substrate for mental activity
Control theory	homeostatic systems, stability simple optimal agent designs

## Potted history of AI

- 1943 McCulloch & Pitts: Boolean circuit model of brain
- 1950 Turing's "Computing Machinery and Intelligence"
- 1952–69 Look, Ma, no hands!
- 1950s Early AI programs, including Samuel's checkers program, Newell & Simon's Logic Theorist, Gelernter's Geometry Engine
- 1956 Dartmouth meeting: "Artificial Intelligence" adopted
- 1965 Robinson's complete algorithm for logical reasoning
- 1966–74 AI discovers computational complexity  
Neural network research almost disappears
- 1969–79 Early development of knowledge-based systems
- 1980–88 Expert systems industry booms
- 1988–93 Expert systems industry busts: "AI Winter"
- 1985–95 Neural networks return to popularity
- 1988– Resurgence of probability; general increase in technical depth  
"Nouvelle AI": ALife, GAs, soft computing
- 1995– Agents, agents, everywhere . . .
- 2003– Human-level AI back on the agenda

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ Drive safely along Telegraph Avenue

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ Buy a week's worth of groceries at Berkeley Bowl

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ Discover and prove a new mathematical theorem

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ **Discover and prove a new mathematical theorem**
- ◆ Design and execute a research program in molecular biology

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ **Discover and prove a new mathematical theorem**
- ◆ **Design and execute a research program in molecular biology**
- ◆ Write an intentionally funny story

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ Discover and prove a new mathematical theorem
- ◆ Design and execute a research program in molecular biology
- ◆ **Write an intentionally funny story**
- ◆ Give competent legal advice in a specialized area of law

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ Discover and prove a new mathematical theorem
- ◆ Design and execute a research program in molecular biology
- ◆ **Write an intentionally funny story**
- ◆ Give competent legal advice in a specialized area of law
- ◆ Translate spoken English into spoken Swedish in real time

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ **Discover and prove a new mathematical theorem**
- ◆ **Design and execute a research program in molecular biology**
- ◆ **Write an intentionally funny story**
- ◆ Give competent legal advice in a specialized area of law
- ◆ Translate spoken English into spoken Swedish in real time
- ◆ Converse successfully with another person for an hour

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ Discover and prove a new mathematical theorem
- ◆ Design and execute a research program in molecular biology
- ◆ Write an intentionally funny story
- ◆ Give competent legal advice in a specialized area of law
- ◆ Translate spoken English into spoken Swedish in real time
- ◆ Converse successfully with another person for an hour
- ◆ Perform a complex surgical operation

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ Discover and prove a new mathematical theorem
- ◆ Design and execute a research program in molecular biology
- ◆ Write an intentionally funny story
- ◆ Give competent legal advice in a specialized area of law
- ◆ Translate spoken English into spoken Swedish in real time
- ◆ Converse successfully with another person for an hour
- ◆ Perform a complex surgical operation
- ◆ Unload any dishwasher and put everything away

## State of the art

Which of the following can be done at present?

- ◆ Play a decent game of table tennis
- ◆ Drive safely along a curving mountain road
- ◆ **Drive safely along Telegraph Avenue**
- ◆ Buy a week's worth of groceries on the web
- ◆ **Buy a week's worth of groceries at Berkeley Bowl**
- ◆ Play a decent game of bridge
- ◆ Discover and prove a new mathematical theorem
- ◆ Design and execute a research program in molecular biology
- ◆ Write an intentionally funny story
- ◆ Give competent legal advice in a specialized area of law
- ◆ Translate spoken English into spoken Swedish in real time
- ◆ Converse successfully with another person for an hour
- ◆ Perform a complex surgical operation
- ◆ **Unload any dishwasher and put everything away**

## Risks and Benefits of AI

“First solve AI, then use AI to solve everything else.” Demis Hassabis,  
CEO of Google DeepMind

### Benefits:

- Decrease repetitive work
- Increase production of goods and services
- Accelerate scientific research (disease cures, climate change and resource shortages solutions)

### Risks:

- Lethal autonomous weapons
- Surveillance and persuasion
- Biased decision making
- Impact on employment
- Safety-critical applications
- Cybersecurity threats

## Risks and Benefits of AI

Development of an artificial superintelligence that surpasses human intelligence may pose a significant risk

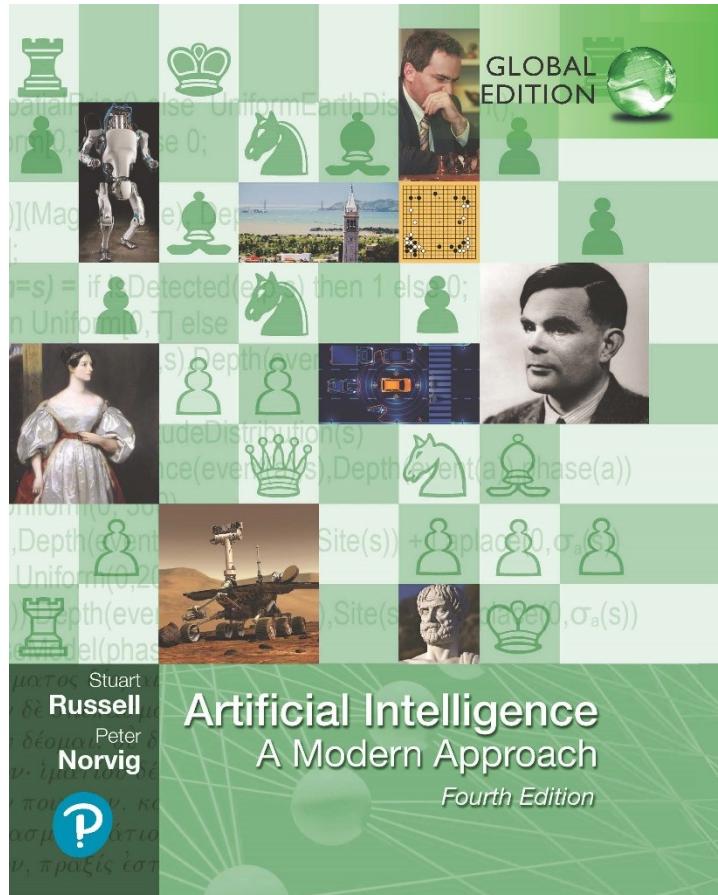
Analogous to the “[Gorilla problem](#)”

Humans and gorillas evolved from the same species, but humans have more control than other primates.

Thus, we should design AI systems in such a way that they do not end up taking control in the way that Turing suggests they might.

# Artificial Intelligence: A Modern Approach

Fourth Edition, Global Edition



## Chapter 2

## Intelligent Agents



# Lecture Presentations: Artificial Intelligence

**Adapted from:**

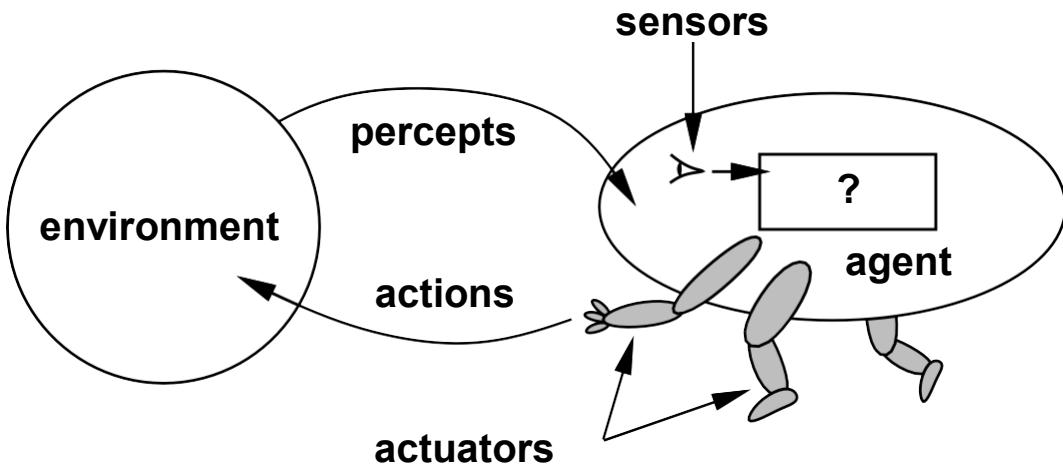
"Artificial Intelligence: A Modern Approach, Global Edition",  
4th Edition by Stuart Russell and Peter Norvig © 2021  
Pearson Education.

**Adapted for** educational use at ACE Engineering College.  
Some slides customized by Mr. Shafakhatullah Khan  
Mohammed, Assistant Professor @ ACE Engineering College.  
For instructional use only. Not for commercial distribution.

# Outline

- ◆ Agents and environments
- ◆ Rationality
- ◆ PEAS (Performance measure, Environment, Actuators, Sensors)
- ◆ Environment types
- ◆ Agent types

# Agents and environments



Agents include humans, robots, softbots, thermostats, etc.

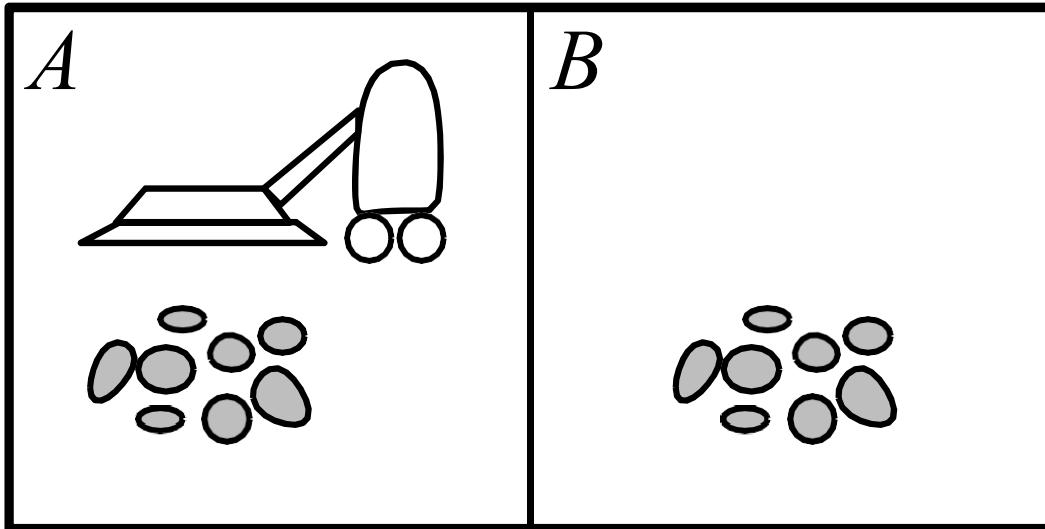
An agent can be anything that can be viewed as perceiving its environment through **sensors** and acting upon that environment through **actuators**

The **agent function** maps from percept histories to actions:

$$f : P^* \rightarrow A$$

The **agent program** runs on the physical **architecture** to produce  $f$

# Vacuum-cleaner world



Percepts: location and contents, e.g., [A, Dirty]

Actions: *Left, Right, Suck, NoOp*

# A vacuum-cleaner agent

Percept sequence	Action
$[A, Clean]$	<i>Right</i>
$[A, Dirty]$	<i>Suck</i>
$[B, Clean]$	<i>Left</i>
$[B, Dirty]$	<i>Suck</i>
$[A, Clean], [A, Clean]$	<i>Right</i>
$[A, Clean], [A, Dirty]$	<i>Suck</i>
.	.

```
function Reflex-Vacuum-Agent( [location,status] ) returns an action
    if status = Dirty then return Suck
    else if location = A then return Right
    else if location = B then return Left
```

What is the **right** function?  
 Can it be implemented in a small agent program?

# Rationality

Fixed **performance measure** evaluates the **environment sequence**

- one point per square cleaned up in time  $T$ ?
- one point per clean square per time step, minus one per move?
- penalize for  $> k$  dirty squares?

A **rational agent** chooses whichever action maximizes the **expected value** of the performance measure **given the percept sequence to date**

Rational  $\neq$  omniscient

- percepts may not supply all relevant information

Rational  $\neq$  clairvoyant

- action outcomes may not be as expected

Hence, rational  $\neq$  successful

Rational  $\Rightarrow$  exploration, learning, autonomy

# PEAS

To design a rational agent, we must specify the **task environment**

Consider, e.g., the task of designing an automated taxi:

Performance measure??

Environment??

Actuators??

Sensors??

# PEAS

To design a rational agent, we must specify the **task environment**

Consider, e.g., the task of designing an automated taxi:

Performance measure?? safety, destination, profits, legality, comfort, ...

Environment?? US streets/freeways, traffic, pedestrians, weather, ...

Actuators?? steering, accelerator, brake, horn, speaker/display, ...

Sensors?? video, accelerometers, gauges, engine sensors, keyboard, GPS, ...

# Internet shopping agent

Performance measure??

Environment??

Actuators??

Sensors??

## Internet shopping agent

Performance measure?? price, quality, appropriateness, efficiency

Environment?? current and future WWW sites, vendors, shippers

Actuators?? display to user, follow URL, fill in form

Sensors?? HTML pages (text, graphics, scripts)

# Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u> <u>Deterministic?</u> ? <u>Episodic??</u> <u>Static??</u> <u>Discrete??</u> <u>Single-agent??</u>				

## Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u> <u>Deterministic?</u> <u>? Episodic??</u> <u>Static??</u> <u>Discrete??</u> <u>Single-agent??</u>	Yes	Yes	No	No

## Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic?</u>	Yes	No	Partly	No
? <u>Episodic??</u>				
<u>Static??</u>				
<u>Discrete??</u>				
<u>Single-agent??</u>				

## Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic?</u>	Yes	No	Partly	No
? <u>Episodic??</u>	No	No	No	No
<u>Static??</u>				
<u>Discrete??</u>				
<u>Single-agent??</u>				

## Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic?</u>	Yes	No	Partly	No
? <u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>				
<u>Single-agent??</u>				

## Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic?</u>	Yes	No	Partly	No
? <u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>	Yes	Yes	Yes	No
<u>Single-agent??</u>				

## Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic?</u>	Yes	No	Partly	No
? <u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>	Yes	Yes	Yes	
<u>Single-agent??</u>	No Yes	No	Yes (except auctions)	No

The environment type largely determines the agent design

The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent

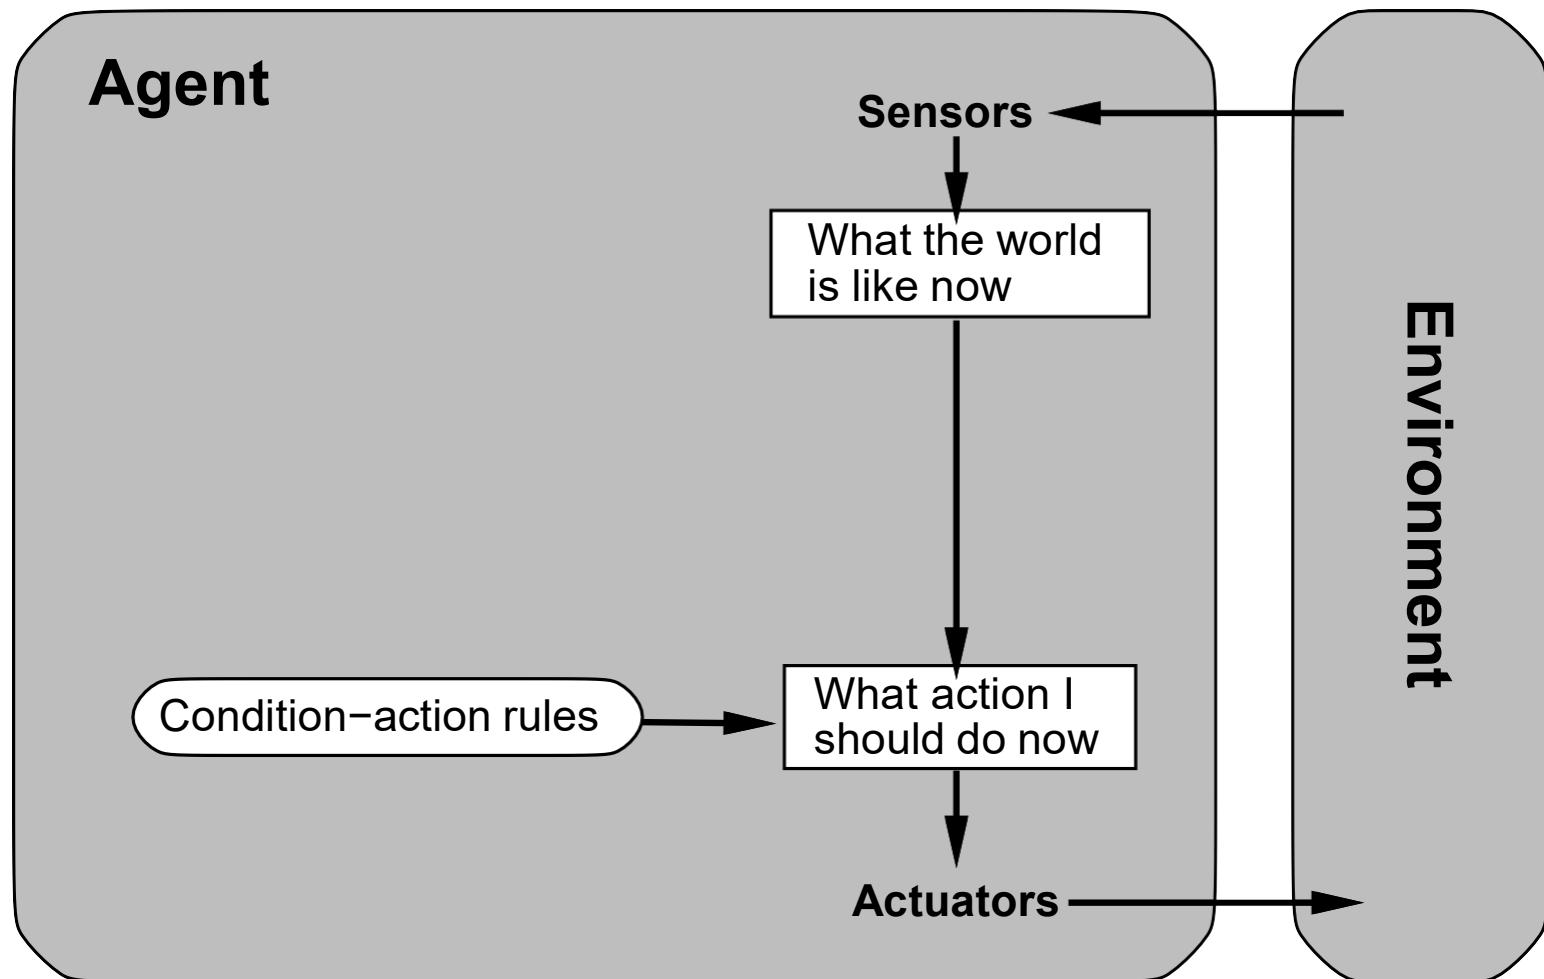
## Agent types

Four basic types in order of increasing generality:

- simple reflex agents
- reflex agents with state
- goal-based agents
- utility-based agents

All these can be turned into learning agents

# Simple reflex agents



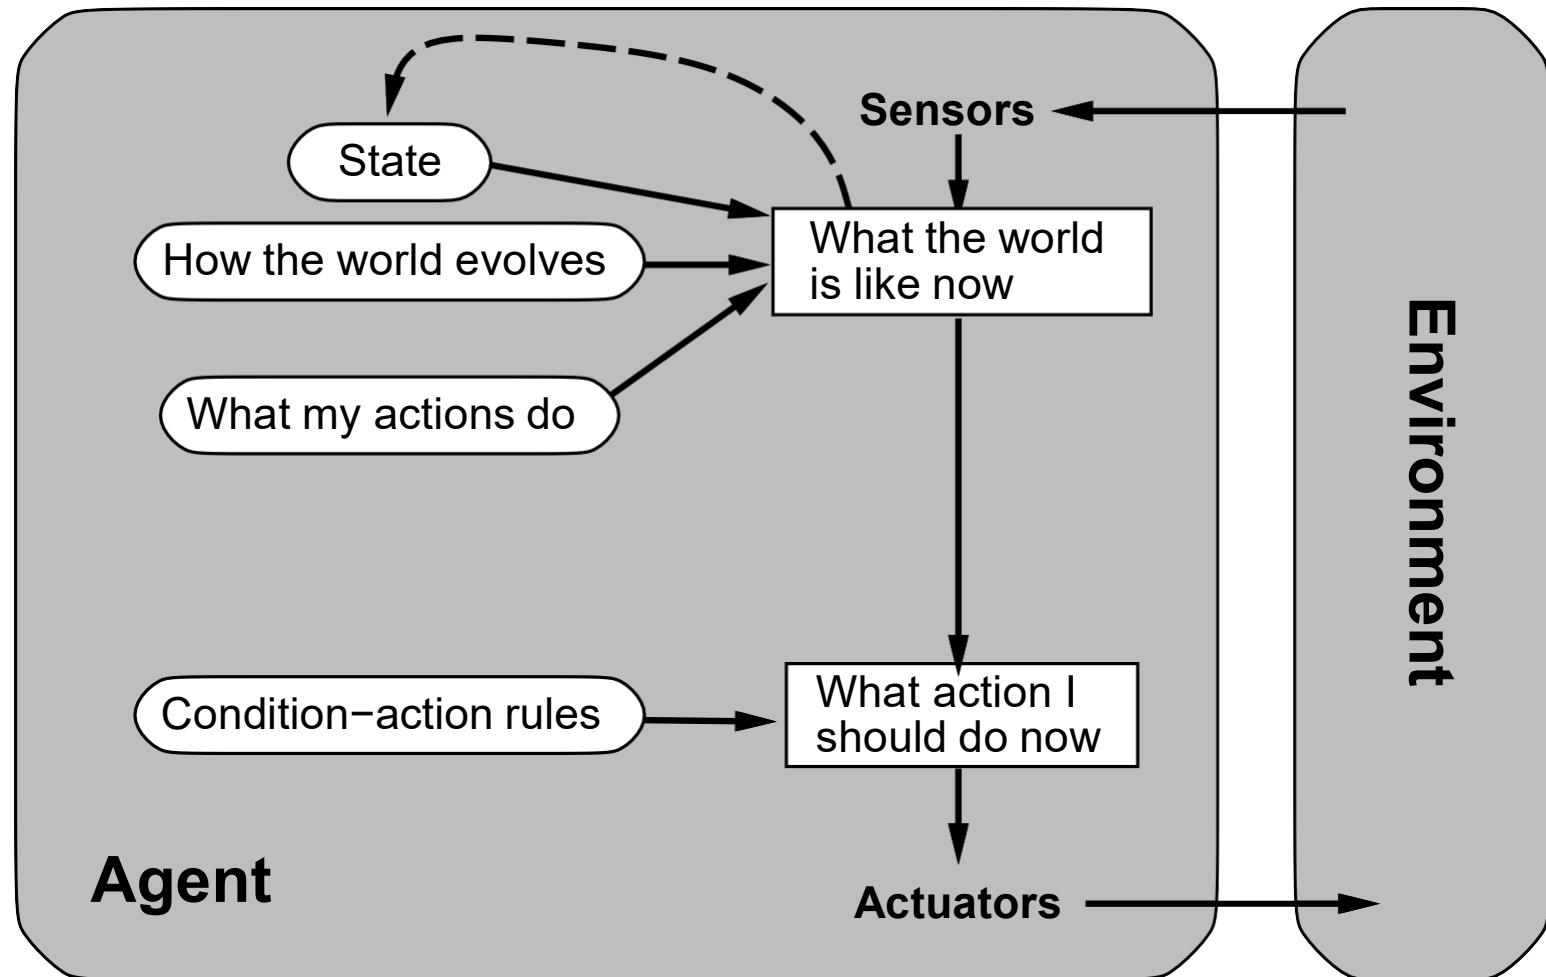
## Example

```
function Reflex-Vacuum-Agent( [location,status]) returns an action
    if status = Dirty then return Suck
    else if location = A then return Right
    else if location = B then return Left
```

```
(setq joe (make-agent :name 'joe :body (make-agent-body)
                      :program (make-reflex-vacuum-agent-program)))
```

```
(defun make-reflex-vacuum-agent-program ()
  #'(lambda (percept)
      (let ((location (first percept)) (status (second percept)))
        (cond ((eq status 'dirty) 'Suck)
              ((eq location 'A) 'Right)
              ((eq location 'B) 'Left))))
```

## Reflex agents with state

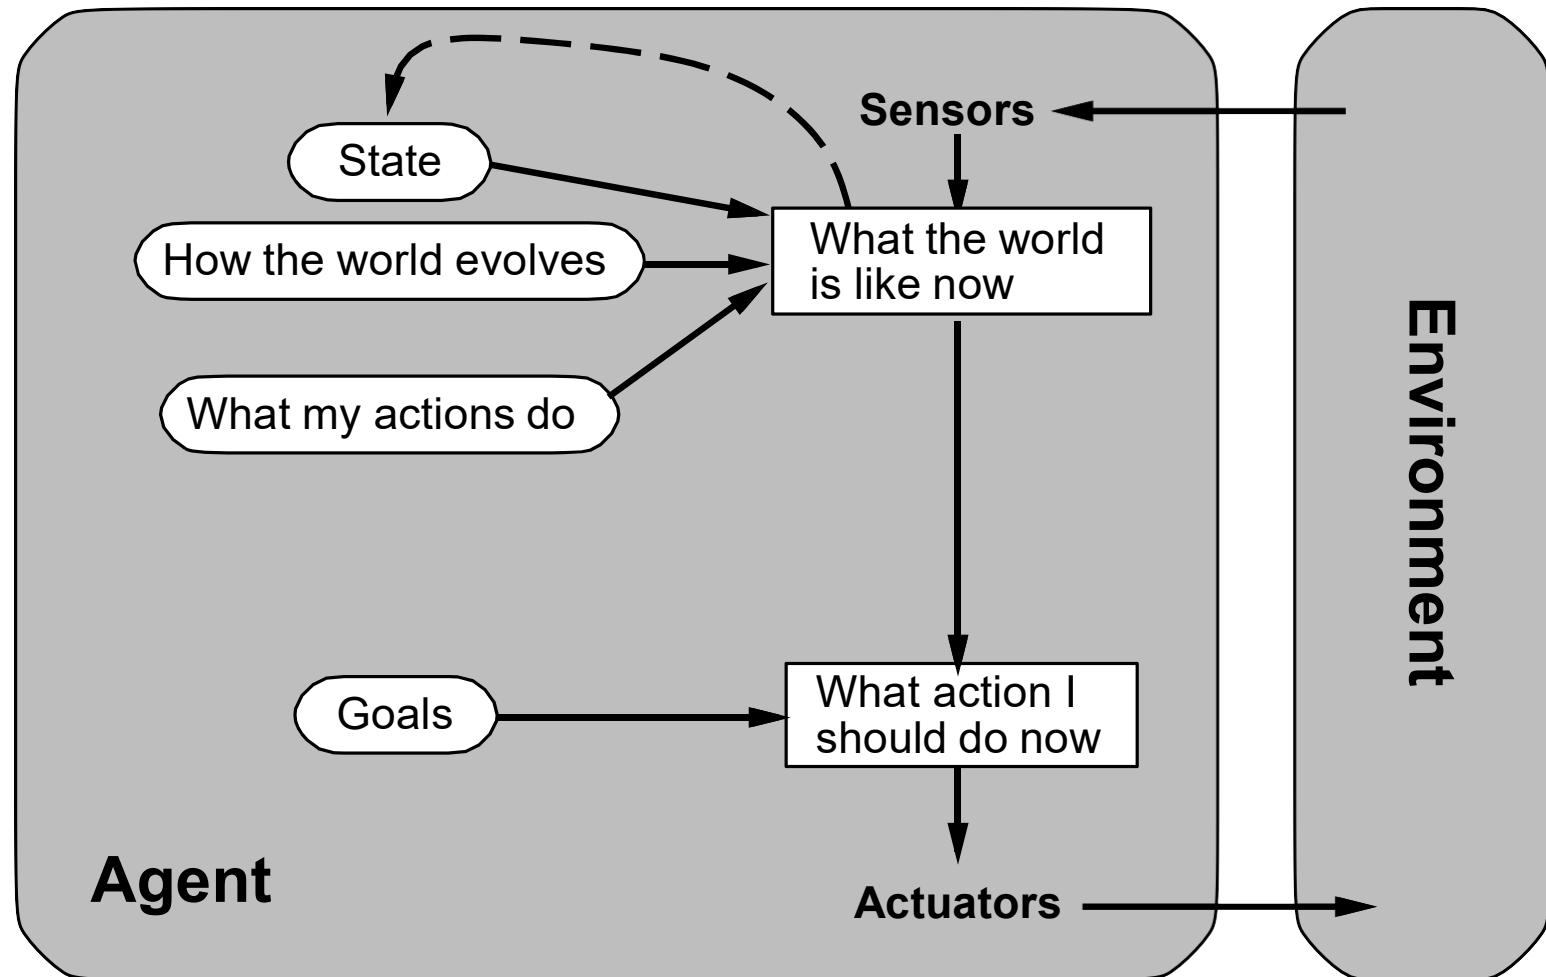


## Example

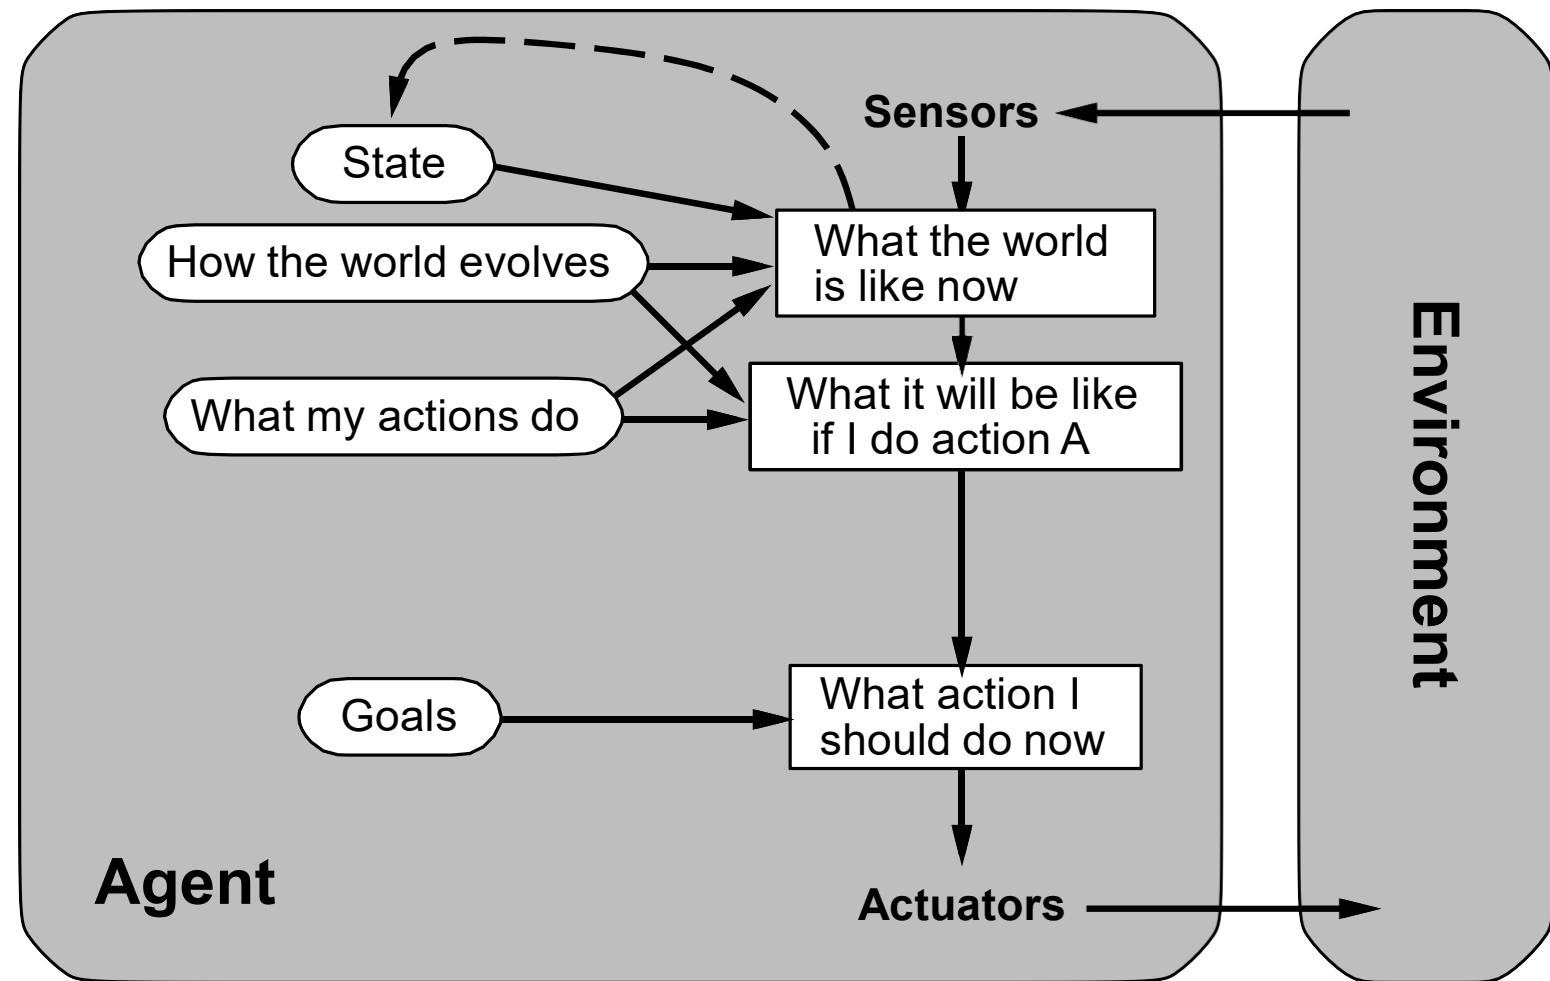
```
function Reflex-Vacuum-Agent( [location,status]) returns an action
static: last_A, last_B, numbers, initially ∞
    if status = Dirty then ...
```

```
(defun make-reflex-vacuum-agent-with-state-program ()
  (let ((last-A infinity) (last-B infinity))
    #'(lambda (percept)
        (let ((location (first percept)) (status (second percept)))
          (incf last-A) (incf last-B)
          (cond
            ((eq status 'dirty)
             (if (eq location 'A) (setq last-A 0) (setq last-B 0))
             'Suck)
            ((eq location 'A) (if (> last-B 3) 'Right 'NoOp))
            ((eq location 'B) (if (> last-A 3) 'Left 'NoOp)))))))
```

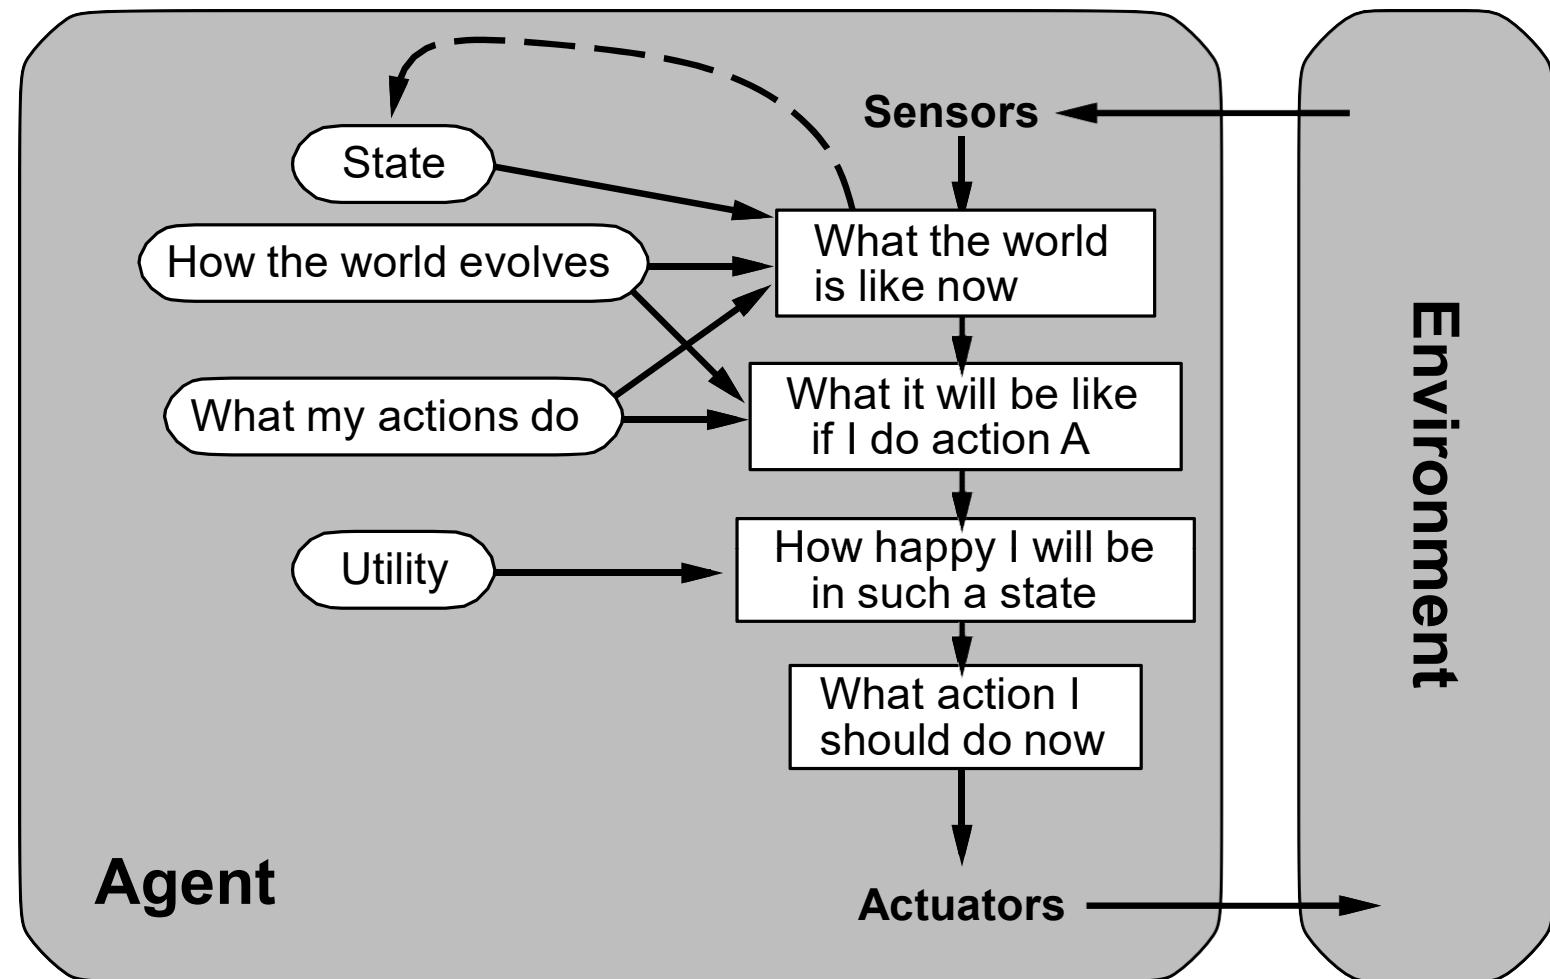
# Model-based agents



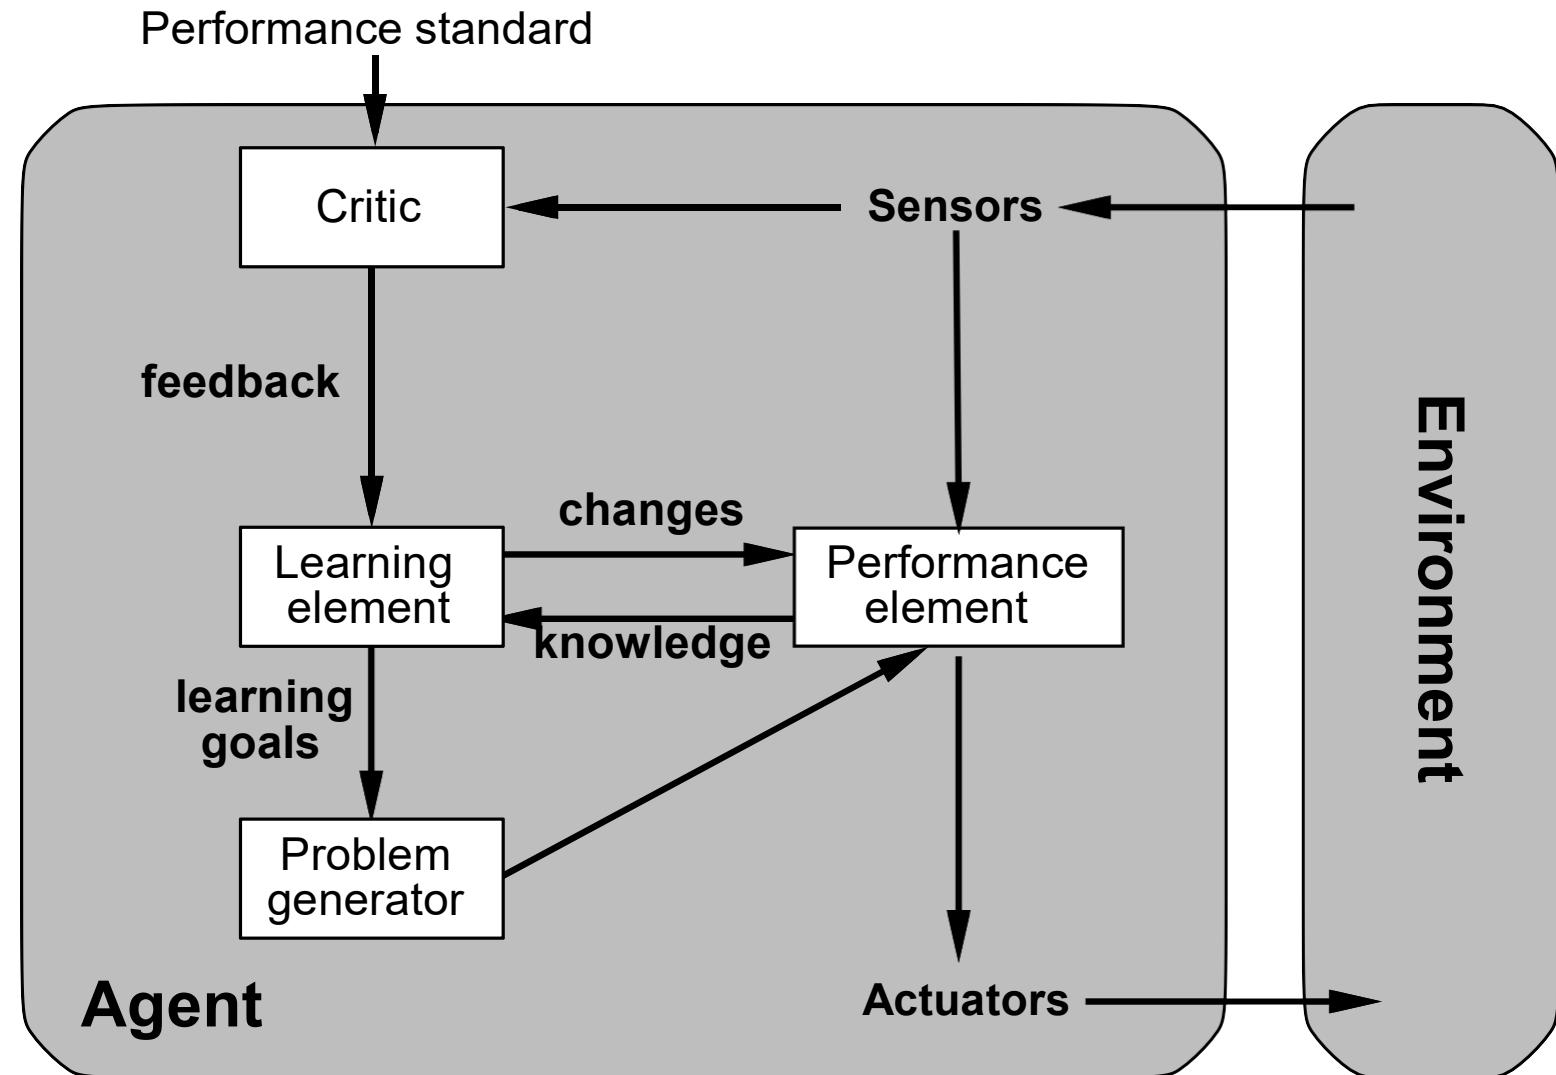
## Goal-based agents



# Utility-based agents



# Learning agents



# Summary

Agents interact with environments through actuators and sensors

The agent function describes what the agent does in all circumstances

The performance measure evaluates the environment sequence

A perfectly rational agent maximizes expected performance

Agent programs implement (some) agent functions

PEAS descriptions define task environments

Environments are categorized along several dimensions:

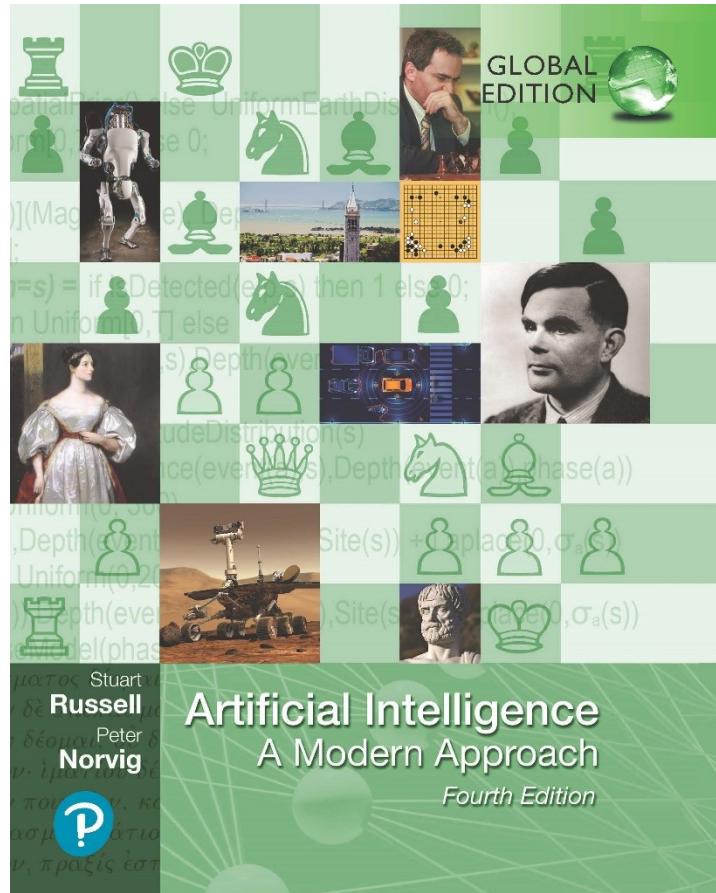
observable? deterministic? episodic? static? discrete? single-agent?

Several basic agent architectures exist:

reflex, reflex with state, goal-based, utility-based

# Artificial Intelligence: A Modern Approach

Fourth Edition, Global Edition



## Chapter 3

## Solving Problems By Searching



# Lecture Presentations: Artificial Intelligence

**Adapted from:**

"Artificial Intelligence: A Modern Approach, Global Edition",  
4th Edition by Stuart Russell and Peter Norvig © 2021  
Pearson Education.

**Adapted for** educational use at ACE Engineering College.  
Some slides customized by Mr. Shafakhatullah Khan  
Mohammed, Assistant Professor @ ACE Engineering College.  
For instructional use only. Not for commercial distribution.

# Outline

- ◆ Problem-solving agents
- ◆ Example Problems
- ◆ Problem formulation
- ◆ Search Algorithms
- ◆ Uninformed Search Strategies
- ◆ Informed (Heuristic) Search Strategies
- ◆ Heuristic Functions

# Problem-solving agents

Restricted form of general agent:

```

function Simple-Problem-Solving-Agent(percept) returns an action
    static: seq, an action sequence, initially empty
            state, some description of the current world state
            goal, a goal, initially null
            problem, a problem formulation

    state  $\leftarrow$  Update-State(state, percept)
    if seq is empty then
        goal  $\leftarrow$  Formulate-Goal(state)
        problem  $\leftarrow$  Formulate-Problem(state, goal)
        seq  $\leftarrow$  Search(problem)
    action  $\leftarrow$  Recommendation(seq, state)
    seq  $\leftarrow$  Remainder(seq, state)
    return action

```

Note: this is **offline** problem solving; solution executed “eyes closed.”  
**Online** problem solving involves acting without complete knowledge.

## Example: Romania

On holiday in Romania; currently in Arad.

Flight leaves tomorrow from Bucharest

**Formulate goal:**

be in Bucharest

**Formulate problem:**

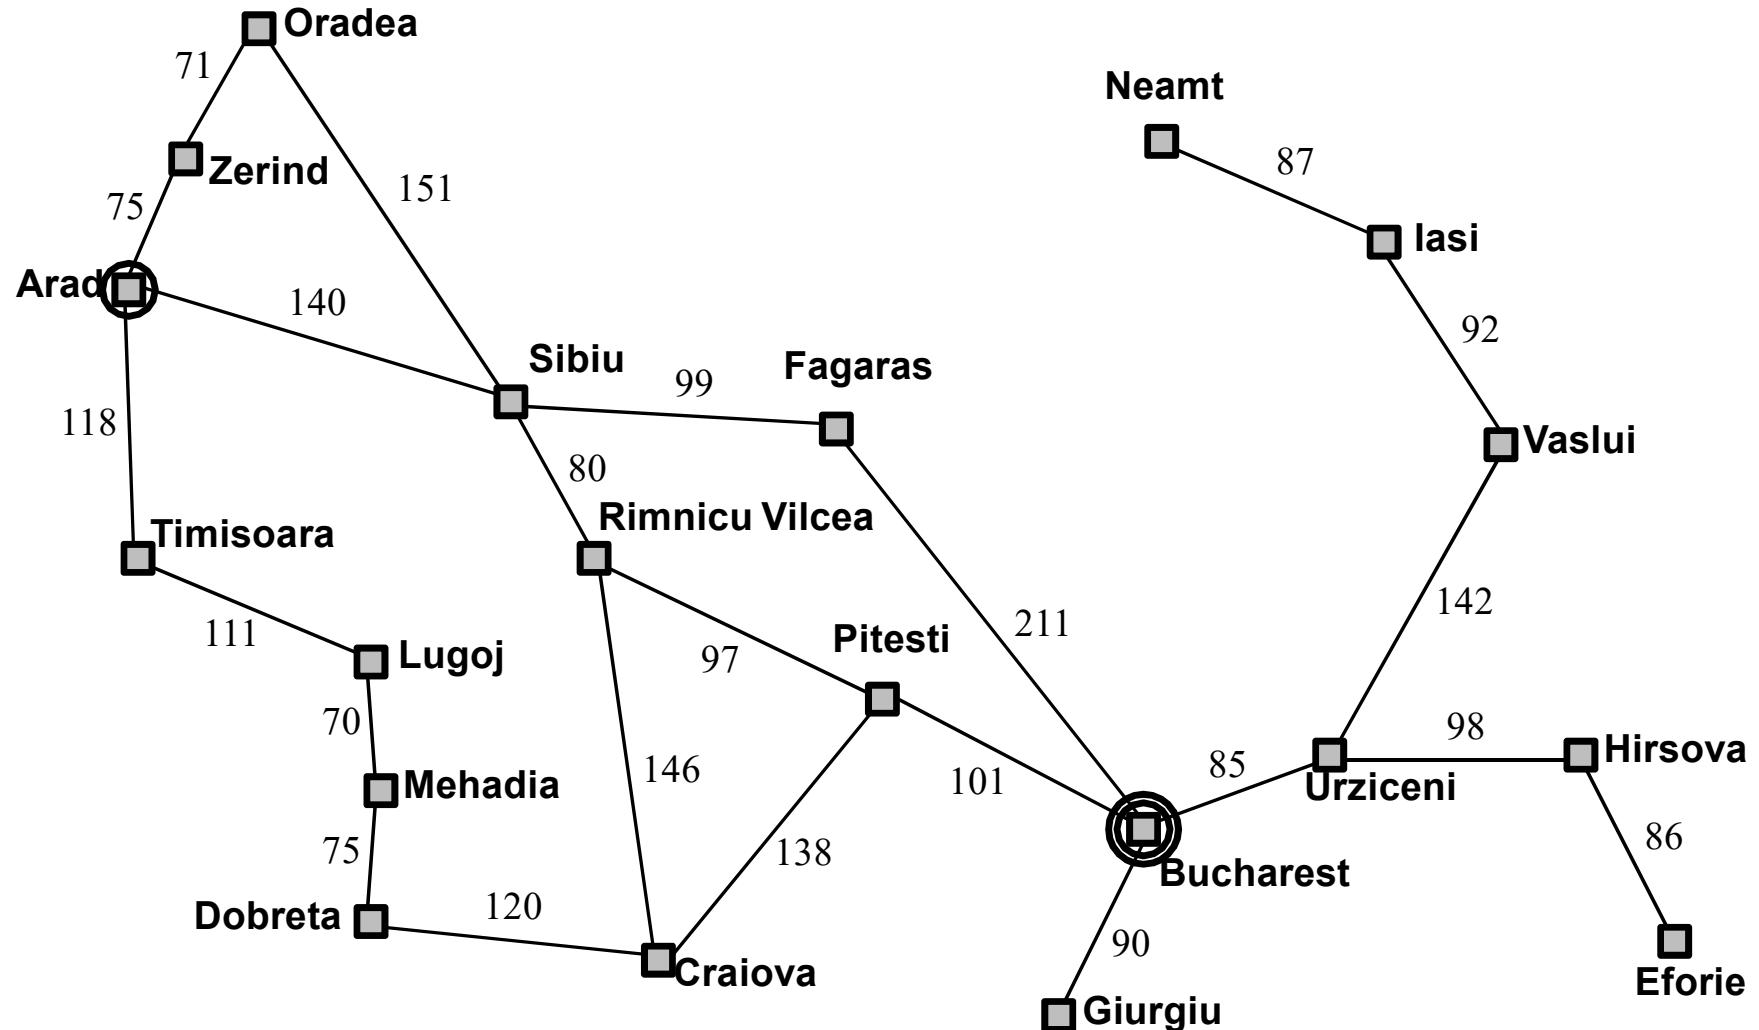
states: various cities

actions: drive between cities

**Find solution:**

sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

## Example: Romania



## Problem types

Deterministic, fully observable  $\Rightarrow$  single-state problem

Agent knows exactly which state it will be in; solution is a sequence

Non-observable  $\Rightarrow$  conformant problem

Agent may have no idea where it is; solution (if any) is a sequence

Nondeterministic and/or partially observable  $\Rightarrow$  contingency problem

percepts provide new information about current state

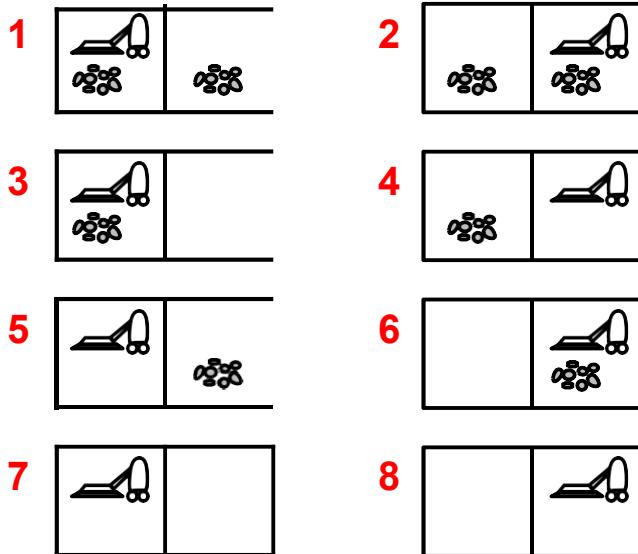
solution is a contingent plan or a policy

often interleave search, execution

Unknown state space  $\Rightarrow$  exploration problem ("online")

## Example: vacuum world

Single-state, start in #5. Solution??



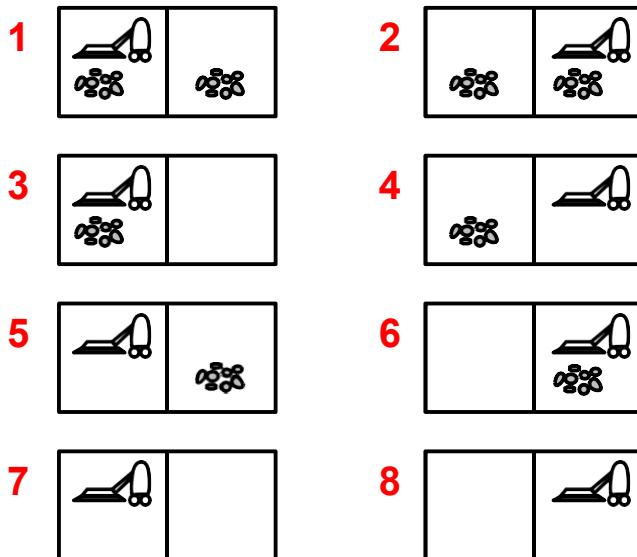
## Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in  $\{1, 2, 3, 4, 5, 6, 7, 8\}$

e.g., *Right* goes to  $\{2, 4, 6, 8\}$ . Solution??



## Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in  $\{1, 2, 3, 4, 5, 6, 7, 8\}$

e.g., *Right* goes to  $\{2, 4, 6, 8\}$ . Solution??

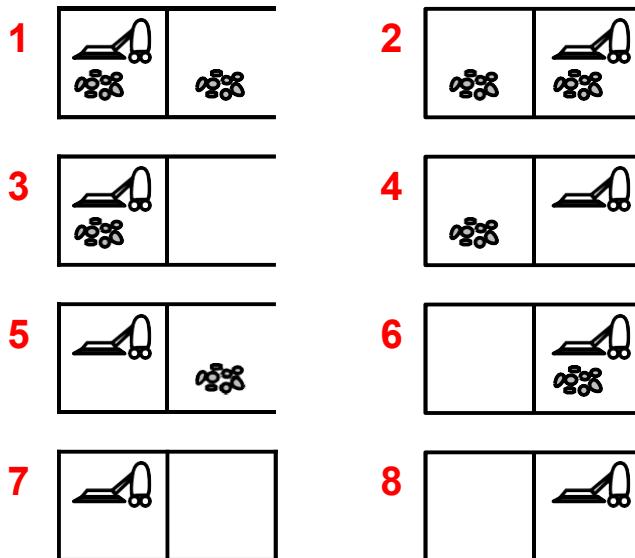
[*Right, Suck, Left, Suck*]

Contingency, start in #5

Murphy's Law: *Suck* can dirty a clean carpet

Local sensing: dirt, location only.

Solution??



## Example: vacuum world

Single-state, start in #5. Solution??

[*Right, Suck*]

Conformant, start in  $\{1, 2, 3, 4, 5, 6, 7, 8\}$

e.g., *Right* goes to  $\{2, 4, 6, 8\}$ . Solution??

[*Right, Suck, Left, Suck*]

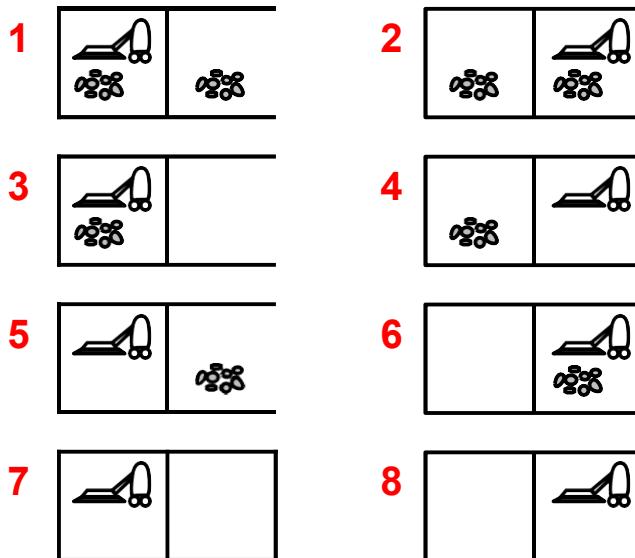
Contingency, start in #5

Murphy's Law: *Suck* can dirty a clean carpet

Local sensing: dirt, location only.

Solution??

[*Right, if dirt then Suck*]



## Single-state problem formulation

A **problem** is defined by four items:

initial state e.g., "at Arad"

successor function  $S(x)$  = set of action-state pairs

e.g.,  $S(Arad) = \{(Arad \rightarrow Zerind, Zerind), \dots\}$

goal test, can be

explicit, e.g.,  $x$  = "at Bucharest"

implicit, e.g.,  $NoDirt(x)$

path cost (additive)

e.g., sum of distances, number of actions executed, etc.

$c(x, a, y)$  is the **step cost**, assumed to be  $\geq 0$

A **solution** is a sequence of actions

leading from the initial state to a goal state

## Selecting a state space

Real world is absurdly complex

⇒ state space must be **abstracted** for problem solving

(Abstract) state = set of real states

(Abstract) action = complex combination of real actions

e.g., "Arad → Zerind" represents a complex set  
of possible routes, detours, rest stops, etc.

For guaranteed realizability, **any** real state "in Arad"

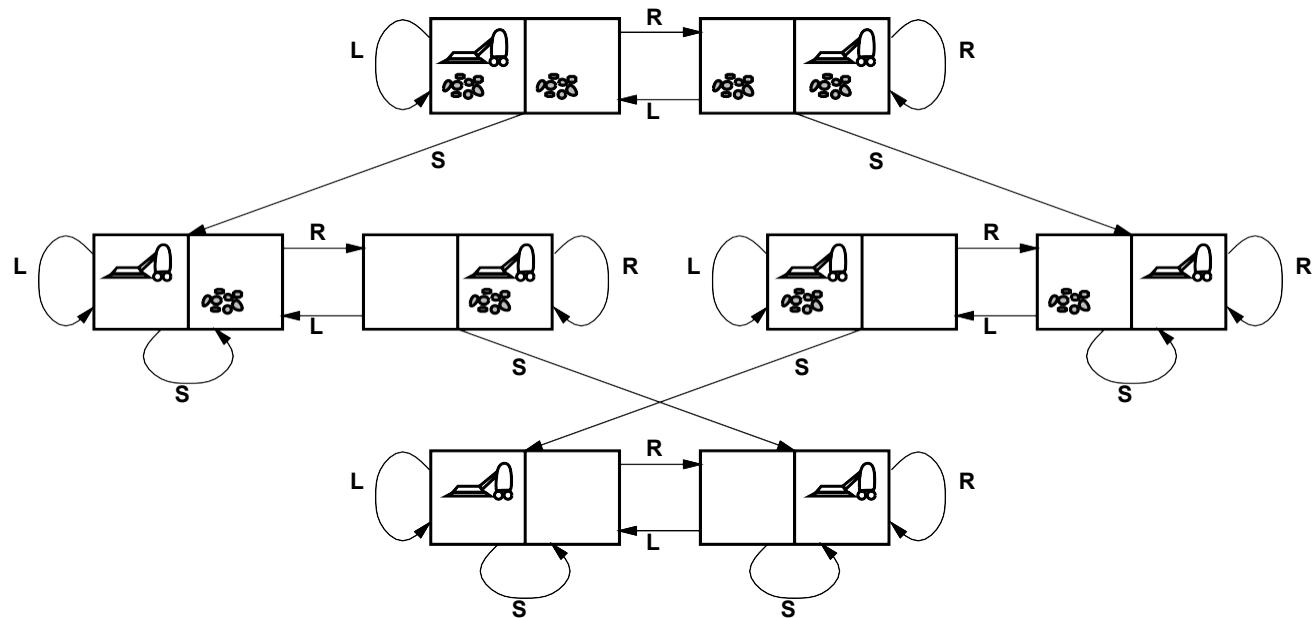
must get to **some** real state "in Zerind"

(Abstract) solution =

set of real paths that are solutions in the real world

Each abstract action should be "easier" than the original problem!

## Example: vacuum world state space graph



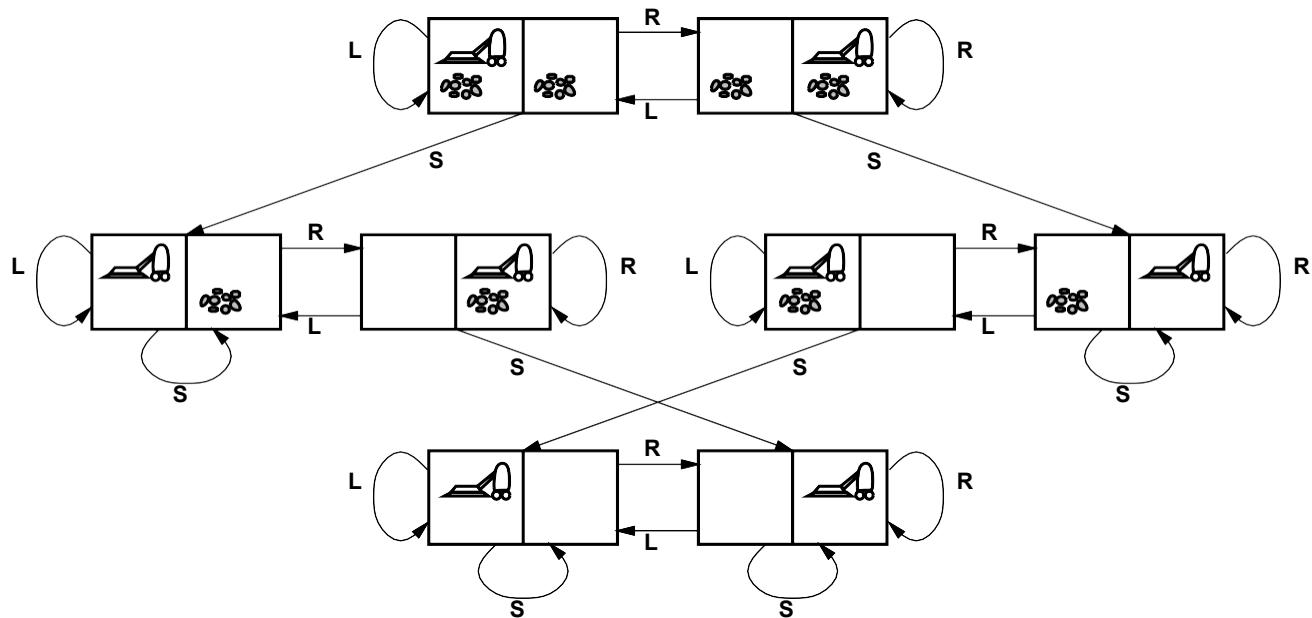
states??

actions??

goal test??

path cost??

## Example: vacuum world state space graph



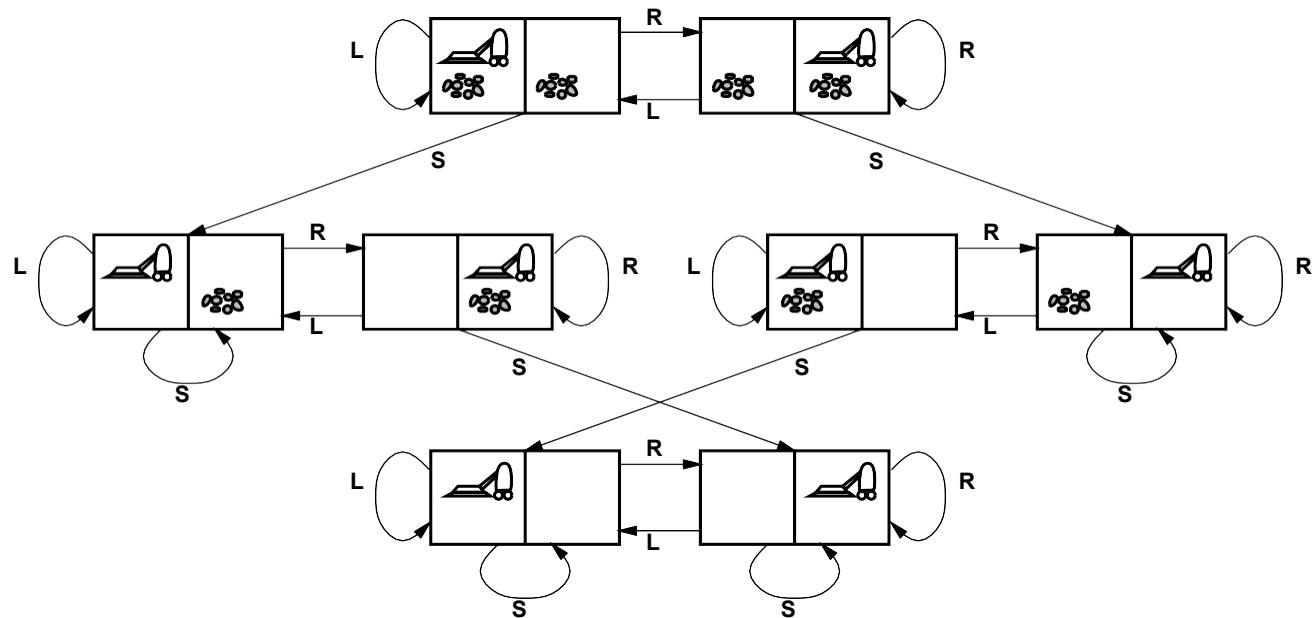
states??: integer dirt and robot locations (ignore dirt amounts etc.)

actions??

goal test??

path cost??

## Example: vacuum world state space graph



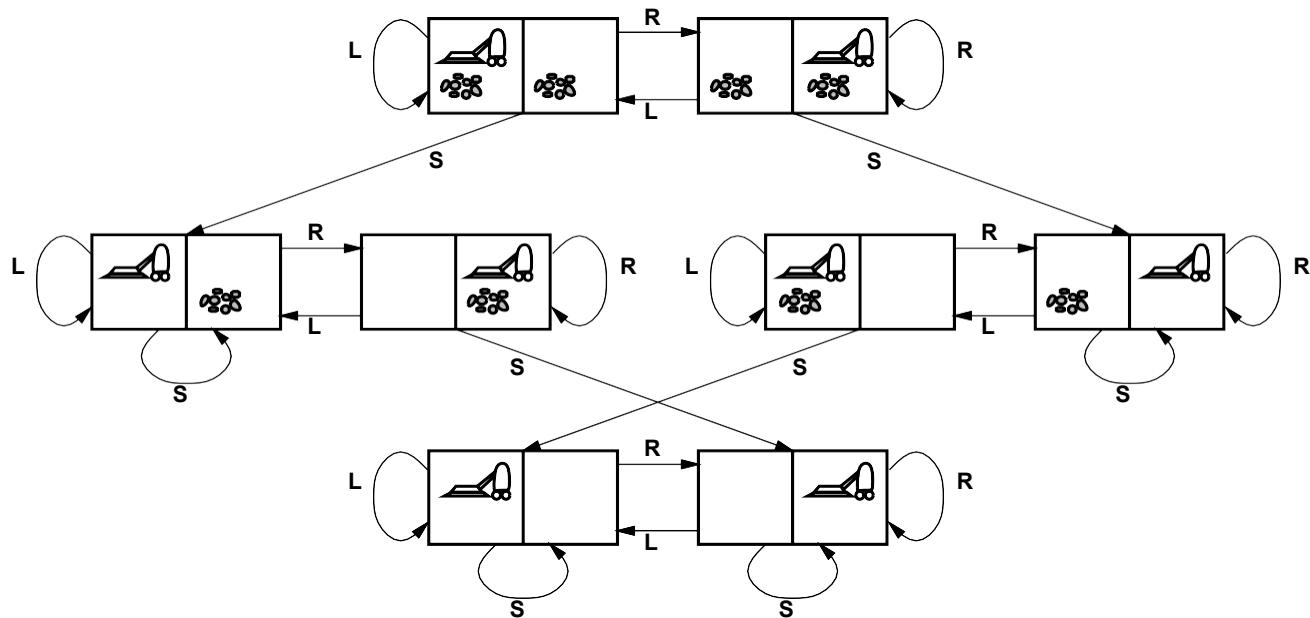
states??: integer dirt and robot locations (ignore dirt amounts etc.)

actions??: *Left*, *Right*, *Suck*, *NoOp*

goal test??

path cost??

## Example: vacuum world state space graph



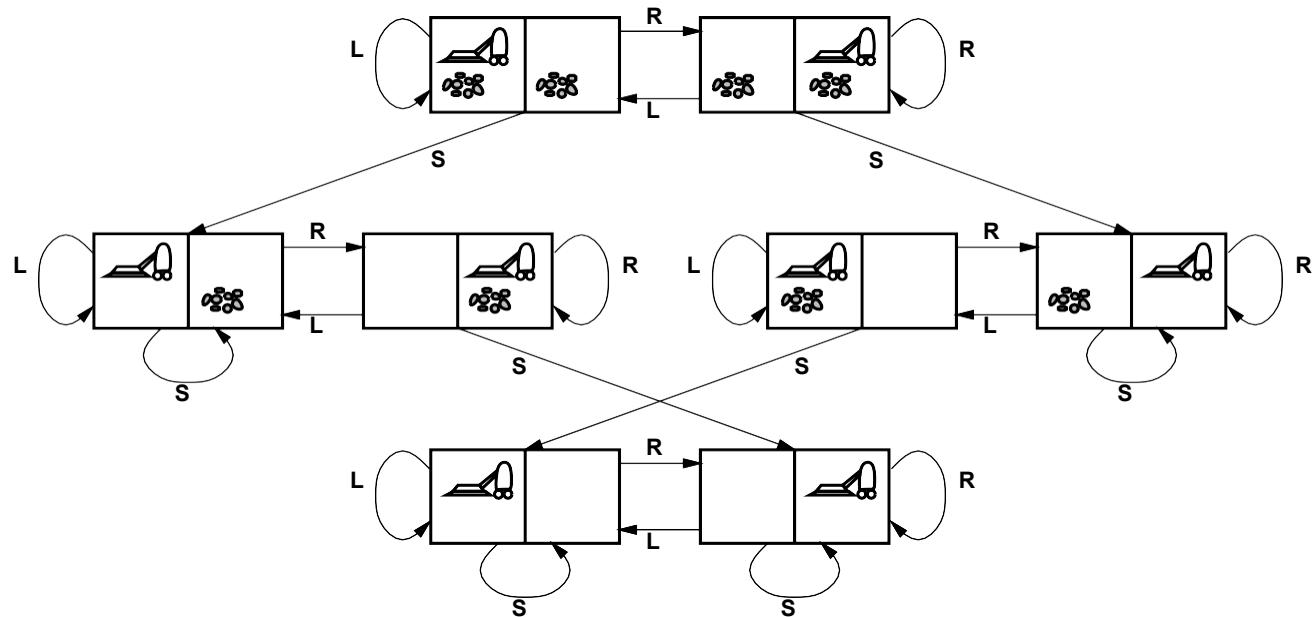
states??: integer dirt and robot locations (ignore dirt amounts etc.)

actions??: *Left*, *Right*, *Suck*, *NoOp*

goal test??: no dirt

path cost??

## Example: vacuum world state space graph



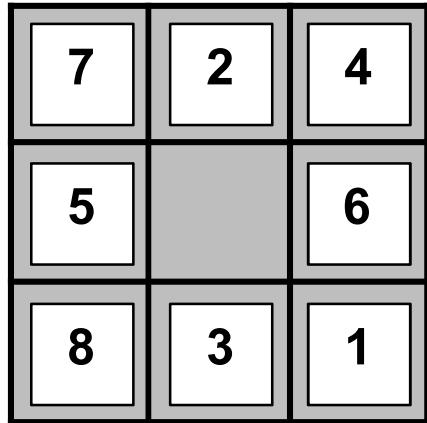
states??: integer dirt and robot locations (ignore dirt amounts etc.)

actions??: *Left*, *Right*, *Suck*, *NoOp*

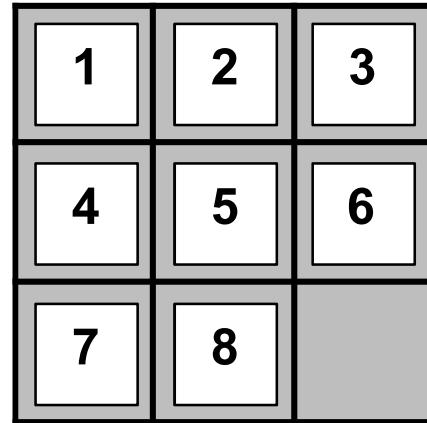
goal test??: no dirt

path cost??: 1 per action (0 for *NoOp*)

## Example: The 8-puzzle



Start State



Goal State

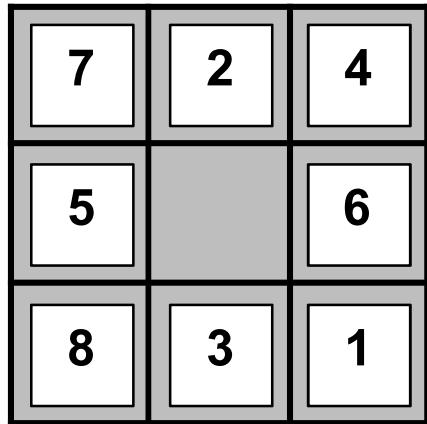
states??

actions??

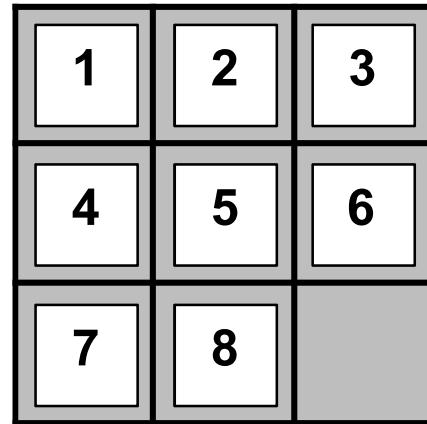
goal test??

path cost??

## Example: The 8-puzzle



Start State



Goal State

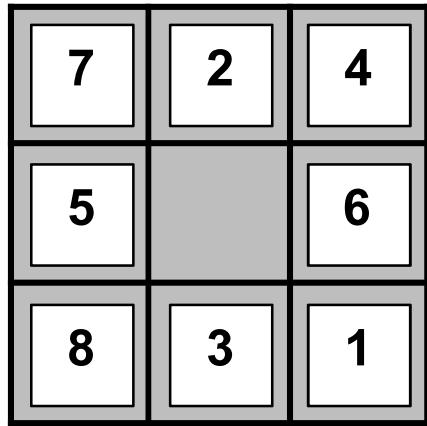
states??: integer locations of tiles (ignore intermediate positions)

actions??

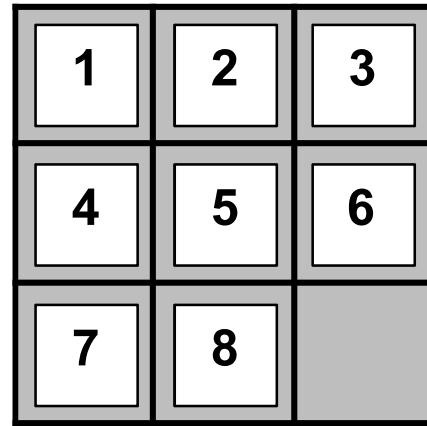
goal test??

path cost??

## Example: The 8-puzzle



Start State



Goal State

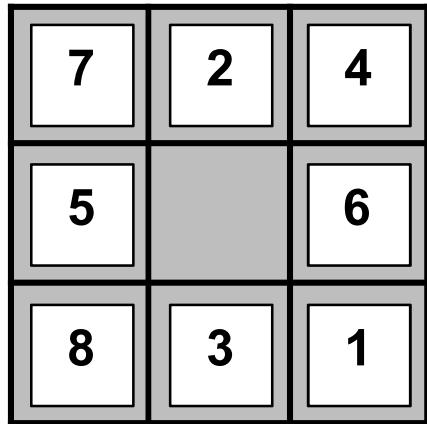
states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

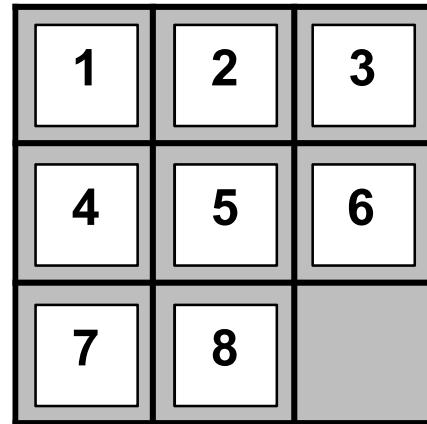
goal test??

path cost??

## Example: The 8-puzzle



Start State



Goal State

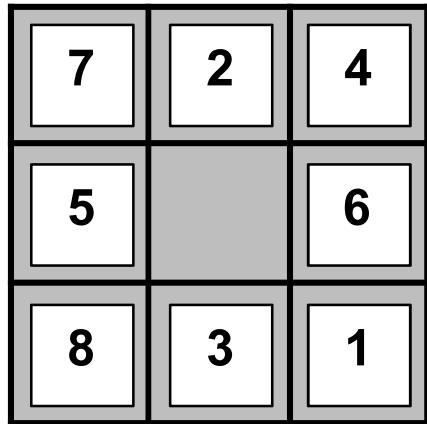
states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

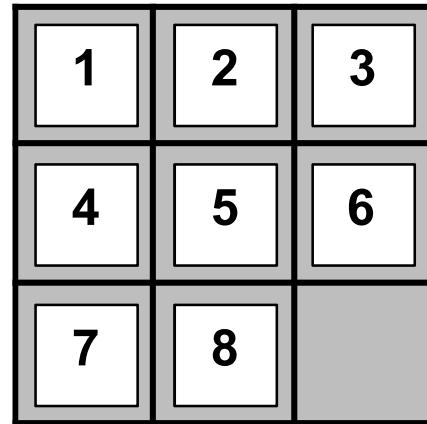
goal test??: = goal state (given)

path cost??

## Example: The 8-puzzle



Start State



Goal State

states??: integer locations of tiles (ignore intermediate positions)

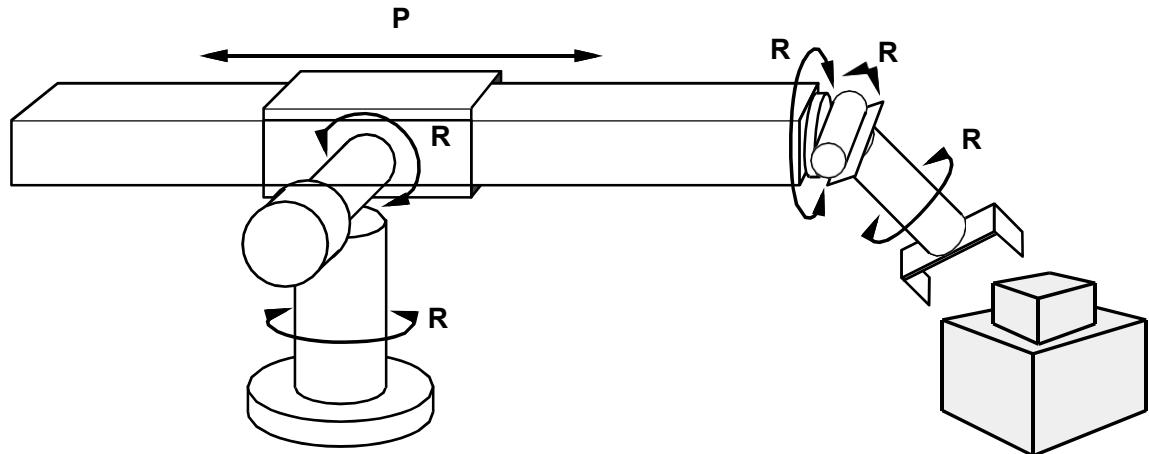
actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)

path cost??: 1 per move

[Note: optimal solution of  $n$ -Puzzle family is NP-hard]

## Example: robotic assembly



states??: real-valued coordinates of robot joint angles  
parts of the object to be assembled

actions??: continuous motions of robot joints

goal test??: complete assembly with no robot included!

path cost??: time to execute

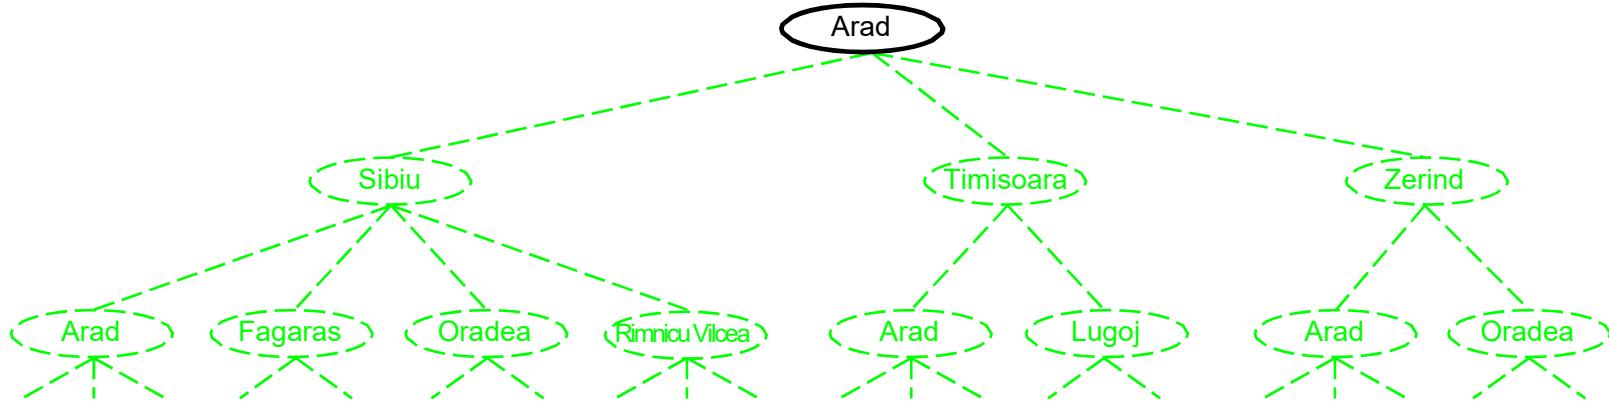
## Tree search algorithms

Basic idea:

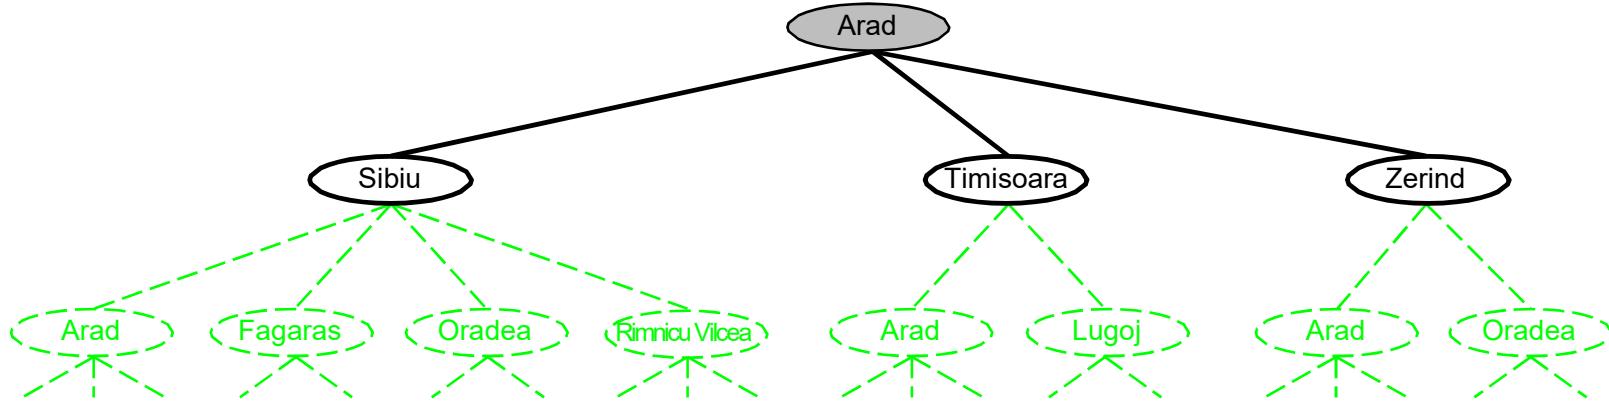
offline, simulated exploration of state space  
by generating successors of already-explored states  
(a.k.a. **expanding** states)

```
function Tree-Search(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
    end
```

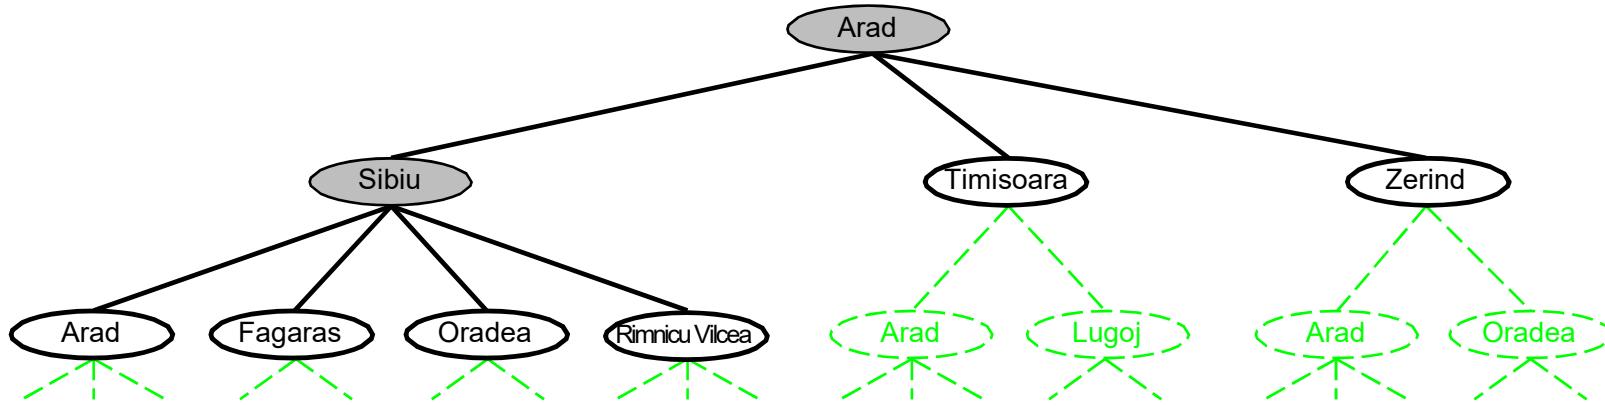
# Tree search example



# Tree search example



# Tree search example

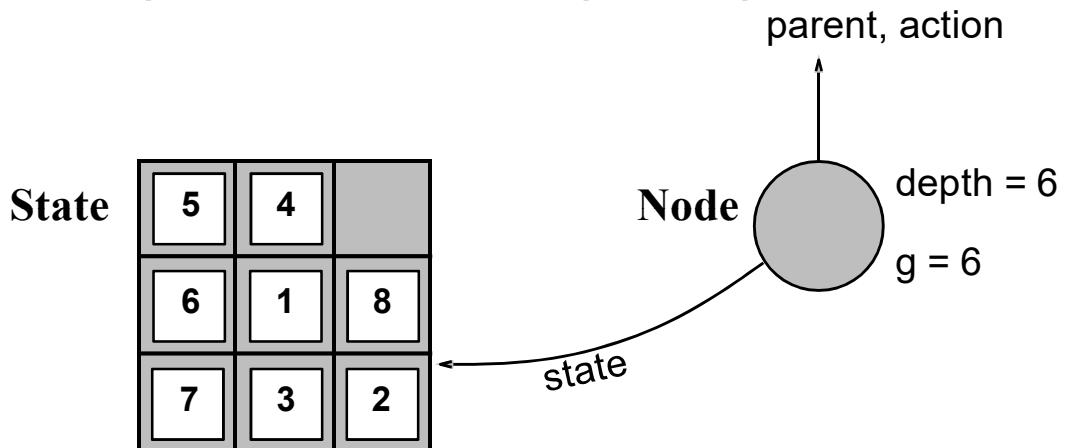


## Implementation: states vs. nodes

A **state** is a (representation of) a physical configuration

A **node** is a data structure constituting part of a search tree  
includes parent, children, depth, path cost  $g(x)$

States do not have parents, children, depth, or path cost!



The Expand function creates new nodes, filling in the various fields and using the SuccessorFn of the problem to create the corresponding states.

## Implementation: general tree search

```

function Tree-Search(problem,fringe) returns a solution, or failure
  fringe  $\leftarrow$  Insert(Make-Node(Initial-State[problem]),fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  Remove-Front(fringe)
    if Goal-Test(problem,State(node)) then return node
    fringe  $\leftarrow$  Insert All(Expand(node,problem),fringe)
  
```

---

```

function Expand(node,problem) returns a set of nodes
  successors  $\leftarrow$  the empty set
  for each action, result in Successor-Fn(problem,State[node]) do
    s  $\leftarrow$  a new Node
    Parent-Node[s]  $\leftarrow$  node; Action[s]  $\leftarrow$  action; State[s]  $\leftarrow$  result
    Path-Cost[s]  $\leftarrow$  Path-Cost[node] + Step-Cost(node,action,s)
    Depth[s]  $\leftarrow$  Depth[node] + 1
    add s to successors
  return successors

```

## Search strategies

A strategy is defined by picking the **order of node expansion**

Strategies are evaluated along the following dimensions:

**completeness**—does it always find a solution if one exists?

**time complexity**—number of nodes generated/expanded

**space complexity**—maximum number of nodes in memory

**optimality**—does it always find a least-cost solution?

Time and space complexity are measured in terms of

***b***—maximum branching factor of the search tree

***d***—depth of the least-cost solution

***m***—maximum depth of the state space (may be  $\infty$ )

## Uninformed search strategies

**Uninformed** strategies use only the information available in the problem definition

Breadth-first search

Uniform-cost search

Depth-first search

Depth-limited search

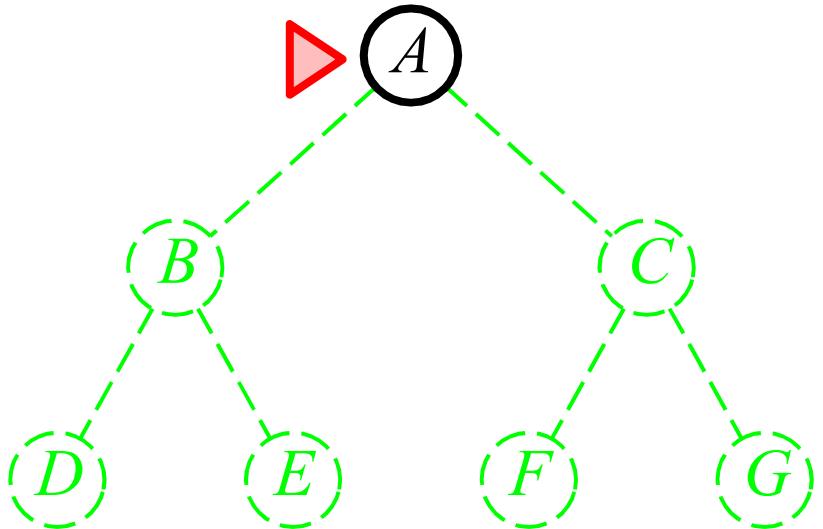
Iterative deepening search

## Breadth-first search

Expand shallowest unexpanded node

Implementation:

*fringe* is a FIFO queue, i.e., new successors go at end

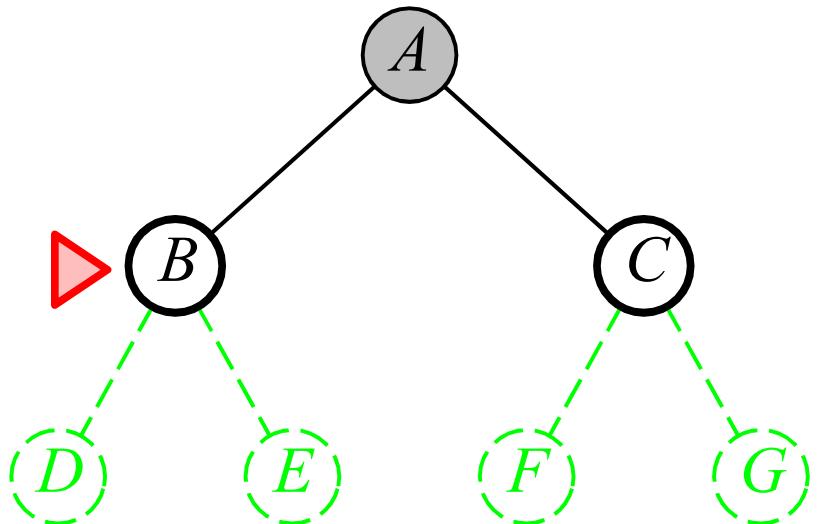


## Breadth-first search

Expand shallowest unexpanded node

Implementation:

*fringe* is a FIFO queue, i.e., new successors go at end

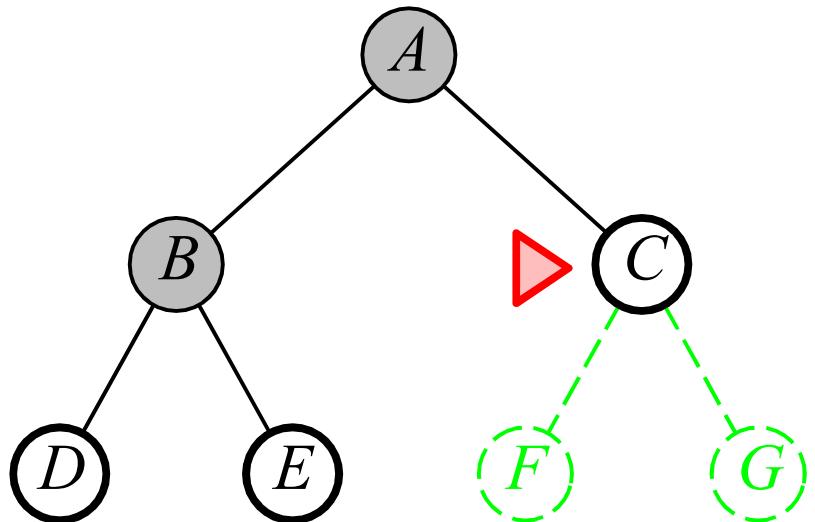


## Breadth-first search

Expand shallowest unexpanded node

Implementation:

*fringe* is a FIFO queue, i.e., new successors go at end

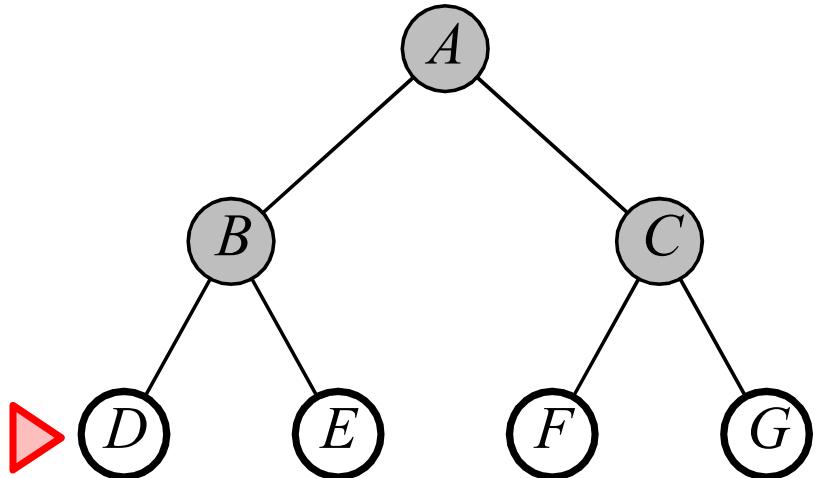


## Breadth-first search

Expand shallowest unexpanded node

Implementation:

*fringe* is a FIFO queue, i.e., new successors go at end



## Properties of breadth-first search

Complete??

## Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??

## Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??  $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$ , i.e., exp. in  $d$

Space??

## Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??  $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$ , i.e., exp. in  $d$

Space??  $O(b^{d+1})$  (keeps every node in memory)

Optimal??

## Properties of breadth-first search

Complete?? Yes (if  $b$  is finite)

Time??  $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$ , i.e., exp. in  $d$

Space??  $O(b^{d+1})$  (keeps every node in memory)

Optimal?? Yes (if cost = 1 per step); not optimal in general

**Space** is the big problem; can easily generate nodes at 100MB/sec  
so 24hrs = 8640GB.

## Uniform-cost search

Expand least-cost unexpanded node

Implementation:

*fringe* = queue ordered by path cost, lowest first

Equivalent to breadth-first if step costs all equal

Complete?? Yes, if step cost  $\geq c$

Time?? # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{fC^*/c})$   
where  $C^*$  is the cost of the optimal solution

Space?? # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{fC^*/c})$

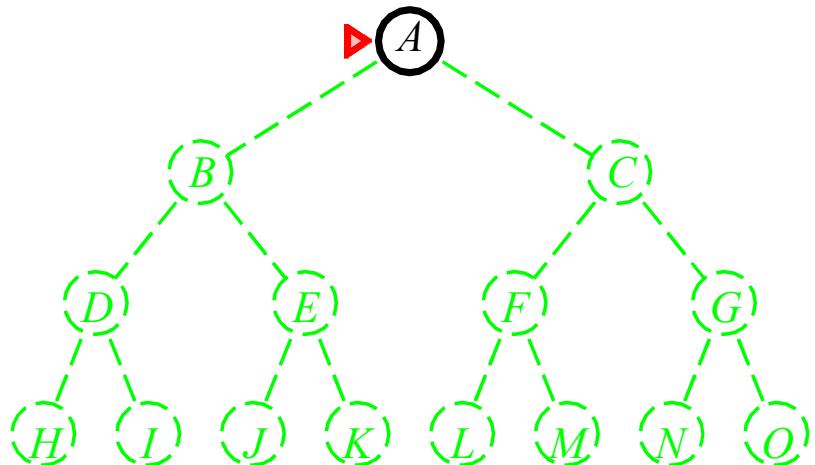
Optimal?? Yes—nodes expanded in increasing order of  $g(n)$

# Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

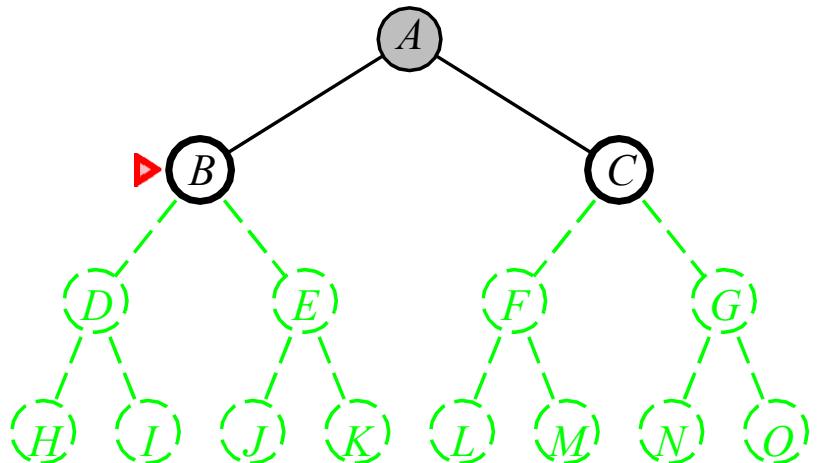


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

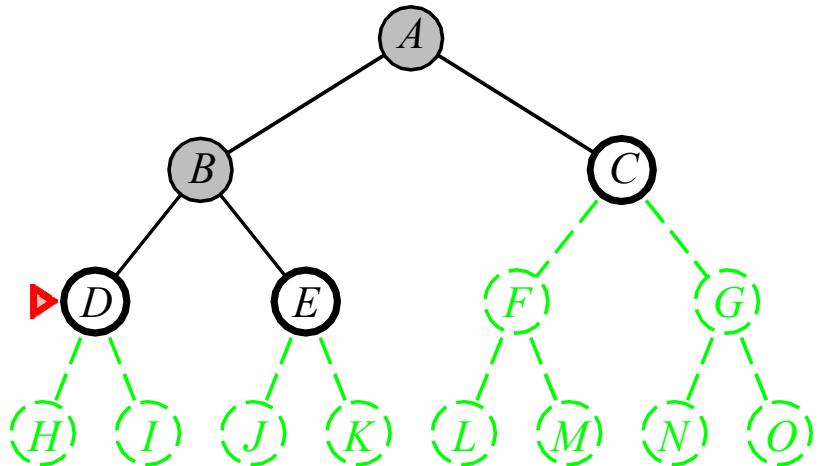


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

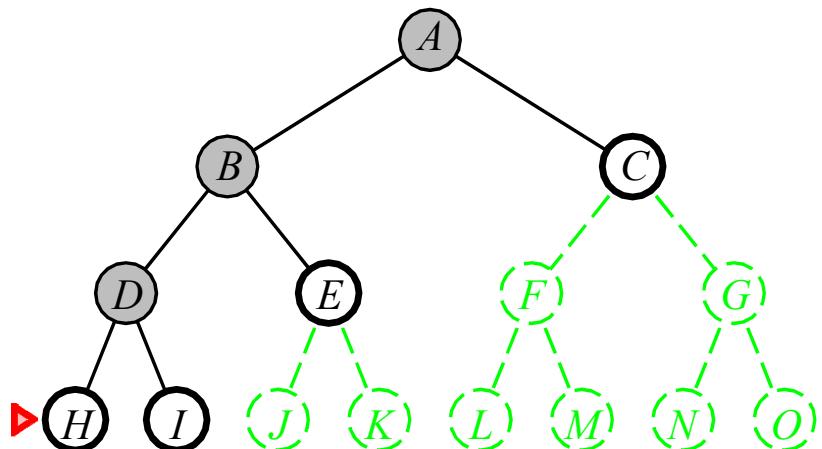


# Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

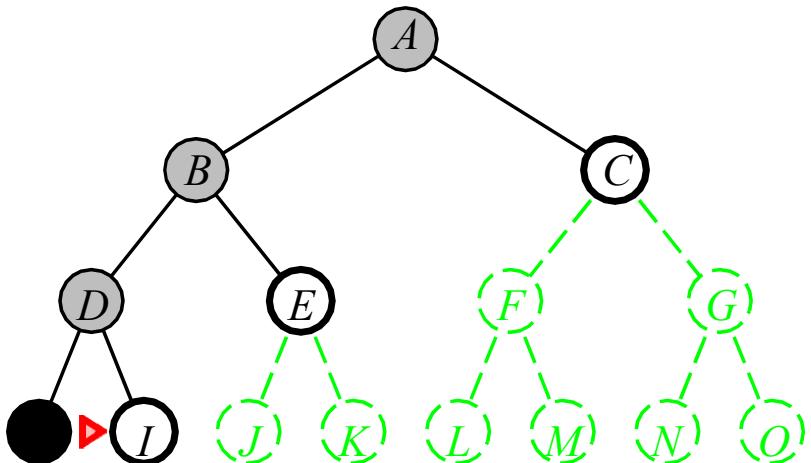


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

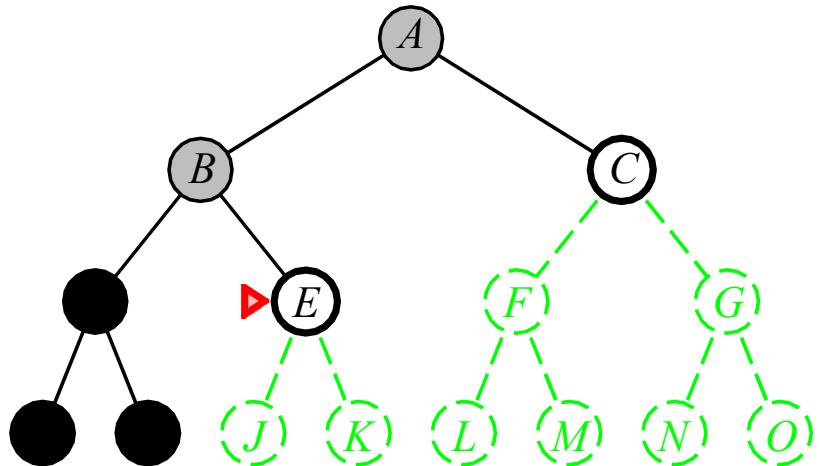


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

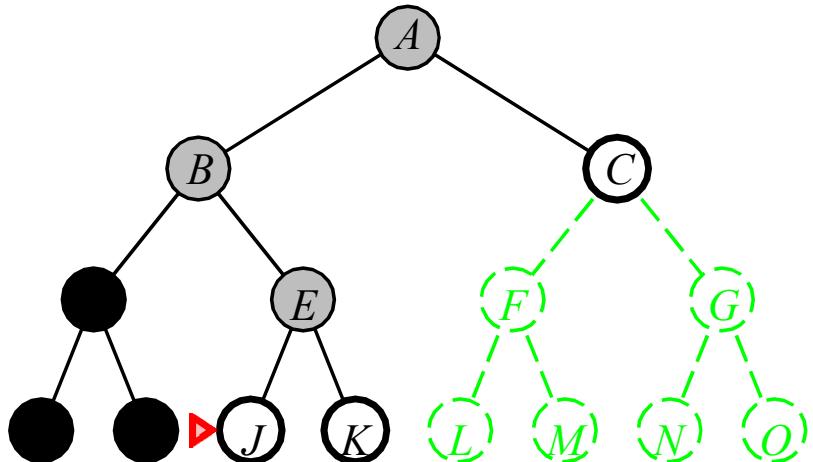


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

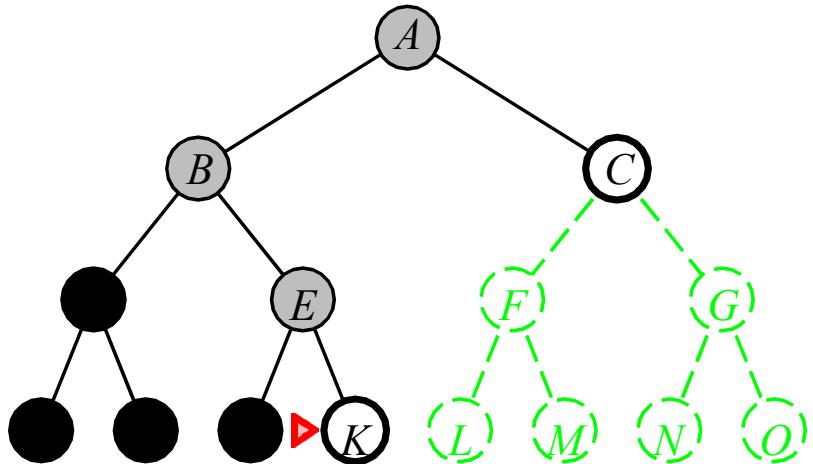


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

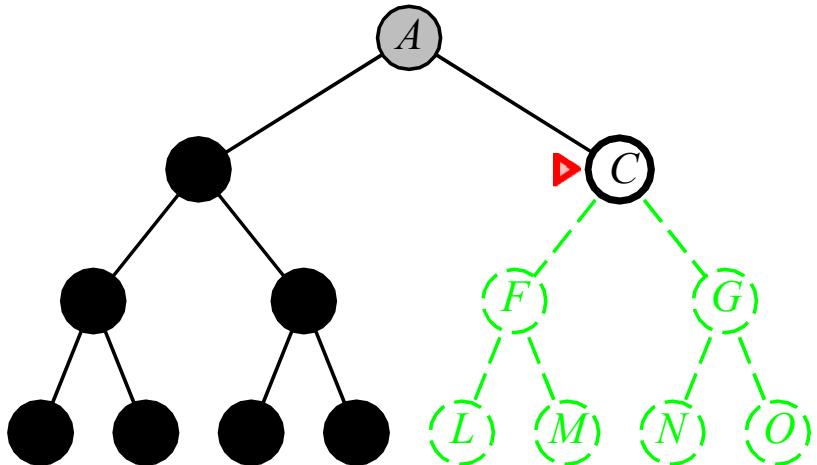


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

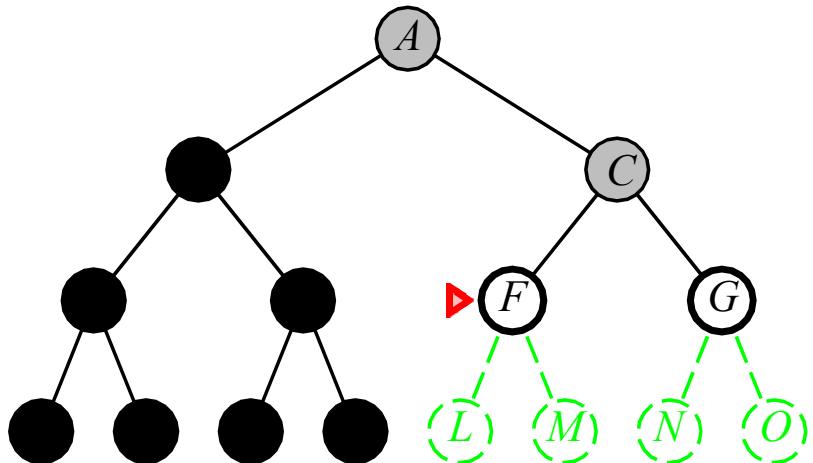


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

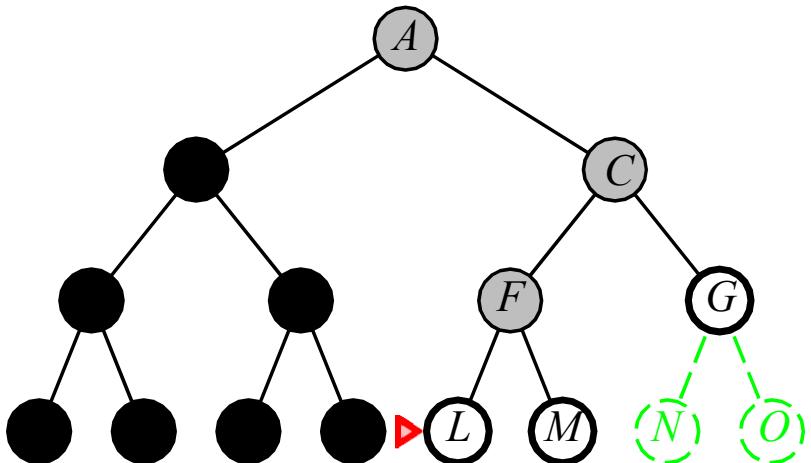


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front

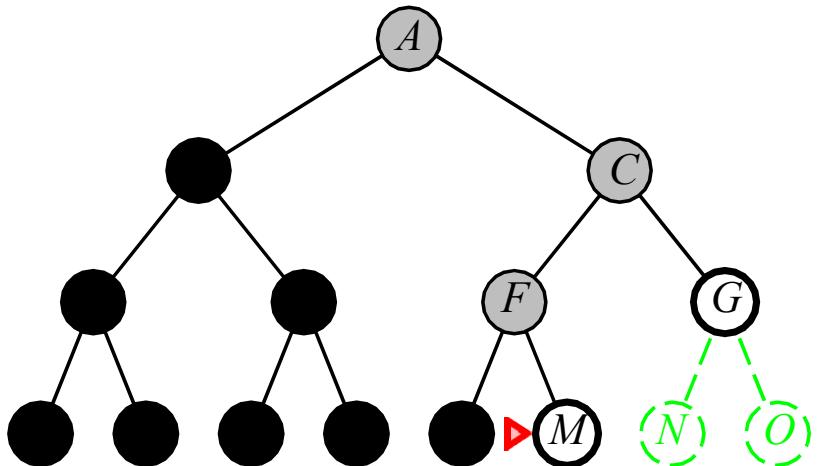


## Depth-first search

Expand deepest unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front



## Properties of depth-first search

Complete??

## Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path  
⇒ complete in finite spaces

Time??

## Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path  
⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$

but if solutions are dense, may be much faster than breadth-first

Space??

## Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path  
⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$

but if solutions are dense, may be much faster than breadth-first

Space??  $O(bm)$ , i.e., linear space!

Optimal??

## Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$

but if solutions are dense, may be much faster than breadth-first

Space??  $O(bm)$ , i.e., linear space!

Optimal?? No

## Depth-limited search

= depth-first search with depth limit  $l$ ,  
 i.e., nodes at depth  $l$  have no successors

Recursive implementation:

```

function Depth-Limited-Search(problem,limit) returns soln/fail/cutoff
  Recursive-DLS(Make-Node(Initial-State [problem]),problem,limit)
function Recursive-DLS(node,problem,limit) returns soln/fail/cutoff
  cutoff-occurred?  $\leftarrow$  false
  if Goal-Test(problem,State[node]) then return node
  else if Depth[node] = limit then return cutoff
  else for each successor in Expand(node,problem) do
    result  $\leftarrow$  Recursive-DLS(successor,problem,limit)
    if result = cutoff then cutoff-occurred?  $\leftarrow$  true
    else if result != failure then return result
  if cutoff-occurred? then return cutoff else return failure

```

## Iterative deepening search

```
function Iterative-Deepening-Search(problem) returns a solution
    inputs: problem, a problem
    for depth  $\leftarrow 0$  to  $\infty$  do
        result  $\leftarrow$  Depth-Limited-Search(problem, depth)
        if result  $\neq$  cutoff then return result
    end
```

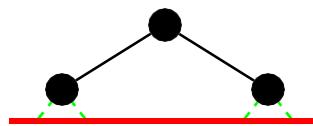
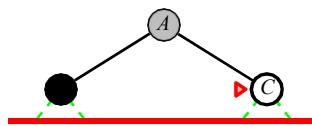
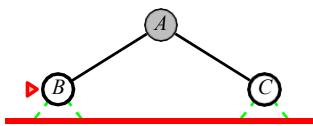
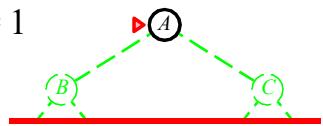
## Iterative deepening search $l = 0$

Limit = 0



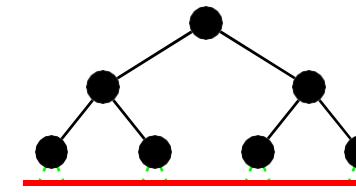
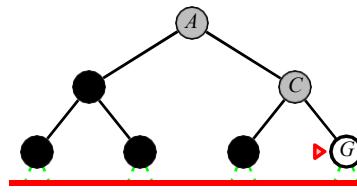
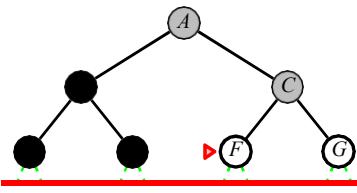
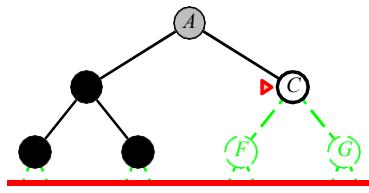
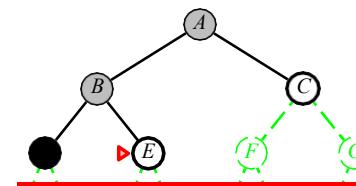
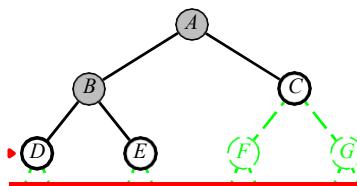
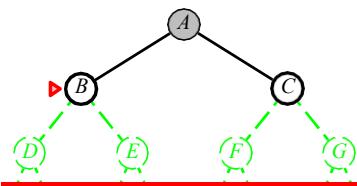
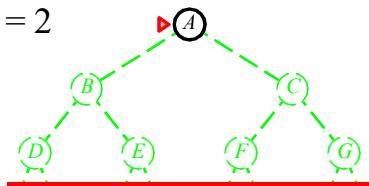
# Iterative deepening search $l = 1$

Limit = 1

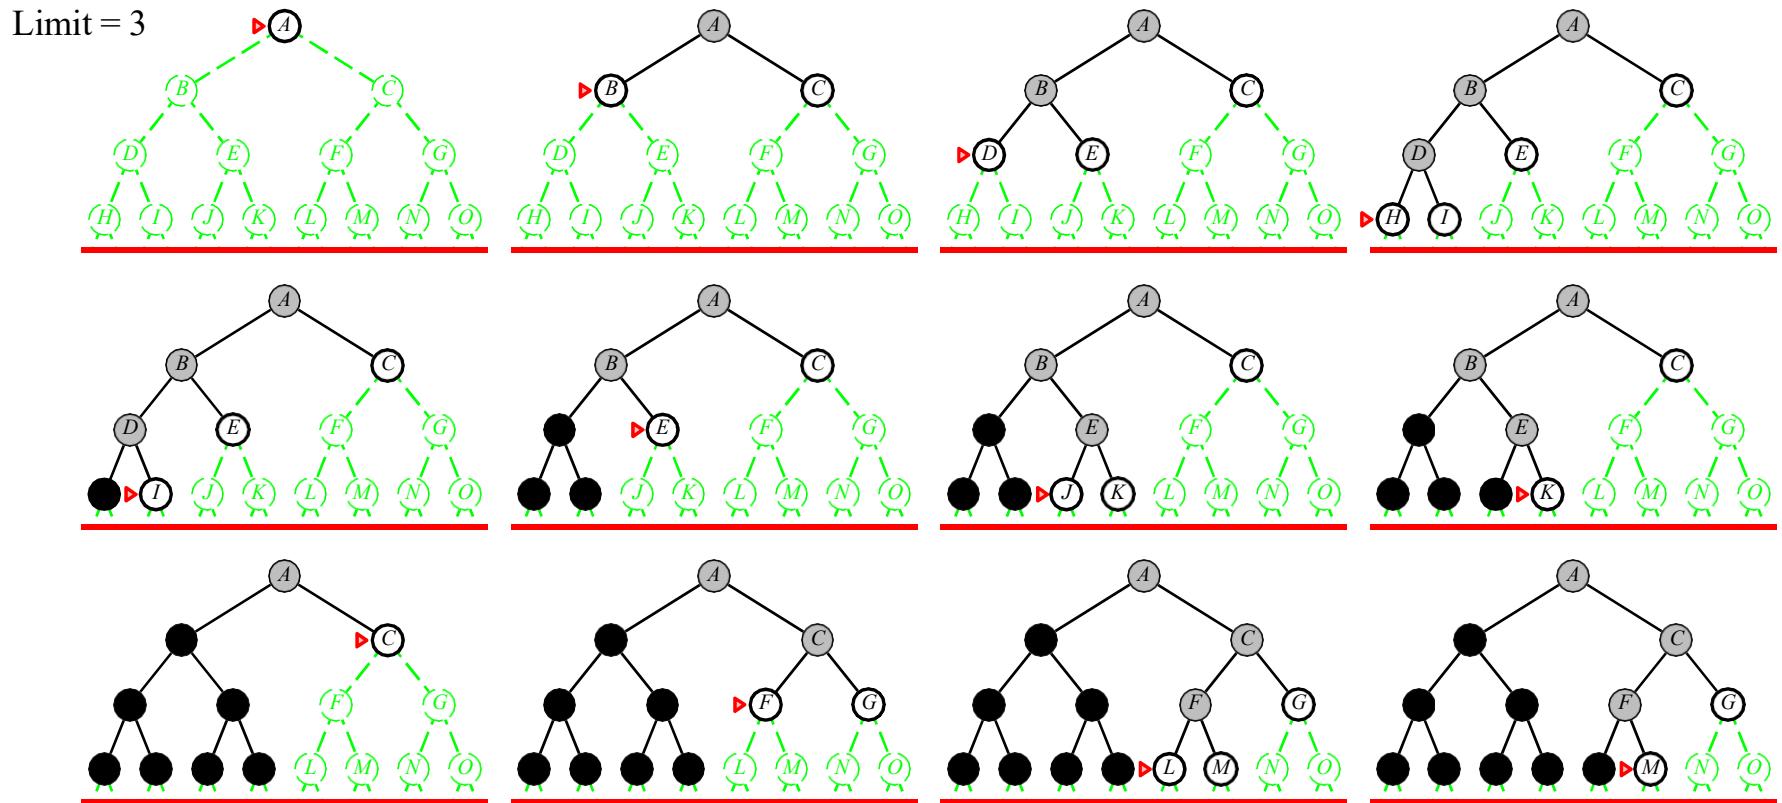


## Iterative deepening search $l = 2$

Limit = 2



## Iterative deepening search $l = 3$



# Properties of iterative deepening search

Complete??

## Properties of iterative deepening search

Complete?? Yes

Time??

## Properties of iterative deepening search

Complete?? Yes

Time??  $(d+1)b + db + (d-1)b + \dots + b^d = O(b^d)$

Space??

## Properties of iterative deepening search

Complete?? Yes

Time??  $(d+1)b + db + (d-1)b + \dots + b^d = O(b^d)$

Space??  $O(bd)$

Optimal??

## Properties of iterative deepening search

Complete?? Yes

Time??  $(d+1)b^0 + db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$

Space??  $O(bd)$

Optimal?? Yes, if step cost = 1

Can be modified to explore uniform-cost tree

Numerical comparison for  $b= 10$  and  $d= 5$ , solution at far right leaf:

$$N(\text{IDS}) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$

$$N(\text{BFS}) = 10 + 100 + 1,000 + 10,000 + 100,000 + 999,990 = 1,111,100$$

IDS does better because other nodes at depth  $d$  are not expanded

BFS can be modified to apply goal test when a node is generated

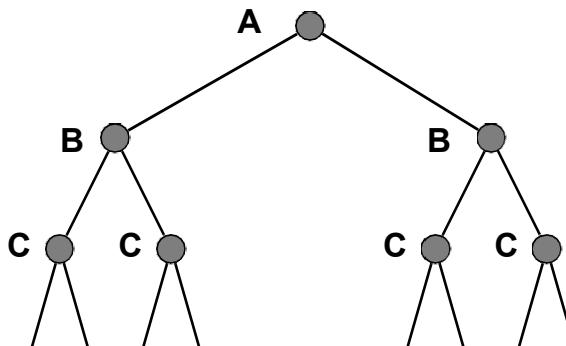
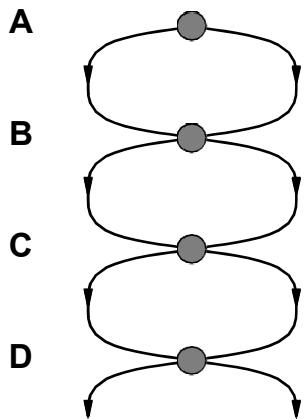
# Summary of Uninformed Search algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No	No	Yes <sup>1</sup>	Yes <sup>1,4</sup>
Optimal cost?	Yes <sup>3</sup>	Yes	No	No	Yes <sup>3</sup>	Yes <sup>3,4</sup>
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$

**Figure 3.15** Evaluation of search algorithms.  $b$  is the branching factor;  $m$  is the maximum depth of the search tree;  $d$  is the depth of the shallowest solution, or is  $m$  when there is no solution;  $\ell$  is the depth limit. Superscript caveats are as follows: <sup>1</sup> complete if  $b$  is finite, and the state space either has a solution or is finite. <sup>2</sup> complete if all action costs are  $\geq \epsilon > 0$ ; <sup>3</sup> cost-optimal if action costs are all identical; <sup>4</sup> if both directions are breadth-first or uniform-cost.

## Repeated states

Failure to detect repeated states can turn a linear problem into an exponential one!



## Graph search

```
function Graph-Search(problem,fringe) returns a solution, or failure
  closed  $\leftarrow$  an empty set
  fringe  $\leftarrow$  Insert (Make-Node(Initial-State[problem]),fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  Remove-Front(fringe)
    if Goal-Test(problem,State[node]) then return node
    if State[node] is not in closed then
      add State[node] to closed
      fringe  $\leftarrow$  Insert All(Expand(node,problem),fringe)
  end
```

## Review: Tree search

```
function Tree-Search(problem, fringe) returns a solution, or failure
  fringe  $\leftarrow$  Insert(Make-Node(Initial-State[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  Remove-Front(fringe)
    if Goal-Test[problem] applied to State(node) succeeds return node
    fringe  $\leftarrow$  InsertAll(Expand(node, problem), fringe)
```

A strategy is defined by picking the order of node expansion

## Best-first search

Idea: use an **evaluation function** for each node

- estimate of “desirability”

⇒ Expand most desirable unexpanded node

Implementation:

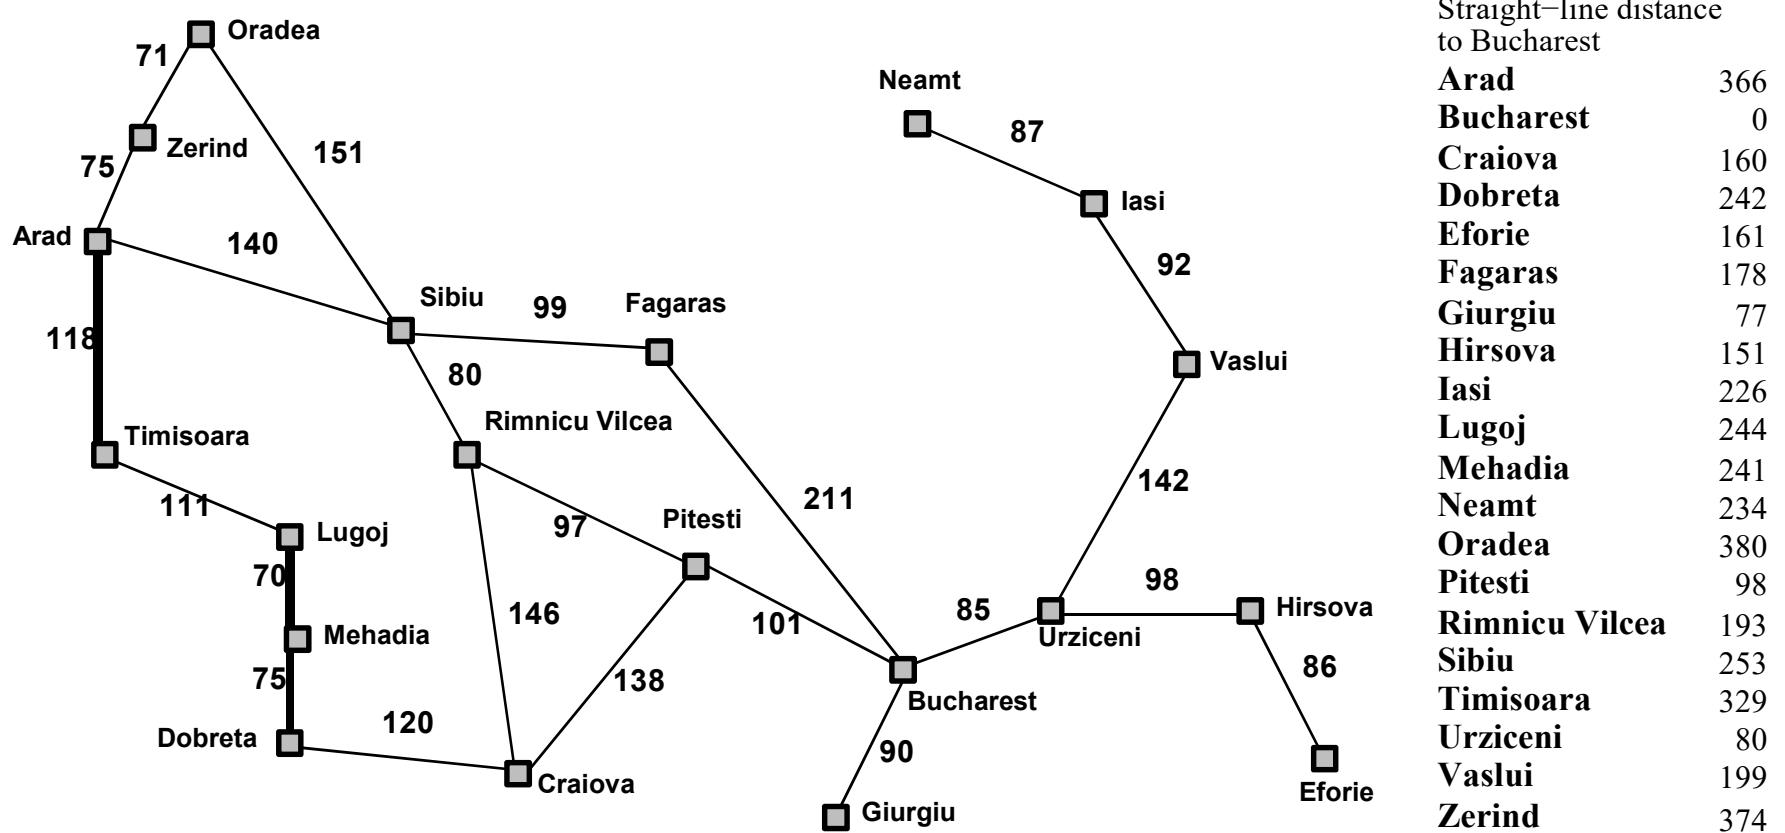
*fringe* is a queue sorted in decreasing order of desirability

Special cases:

- greedy search

- A\* search

# Romania with step costs in km



## Greedy search

Evaluation function  $h(n)$  (heuristic)

= estimate of cost from  $n$  to the closest goal

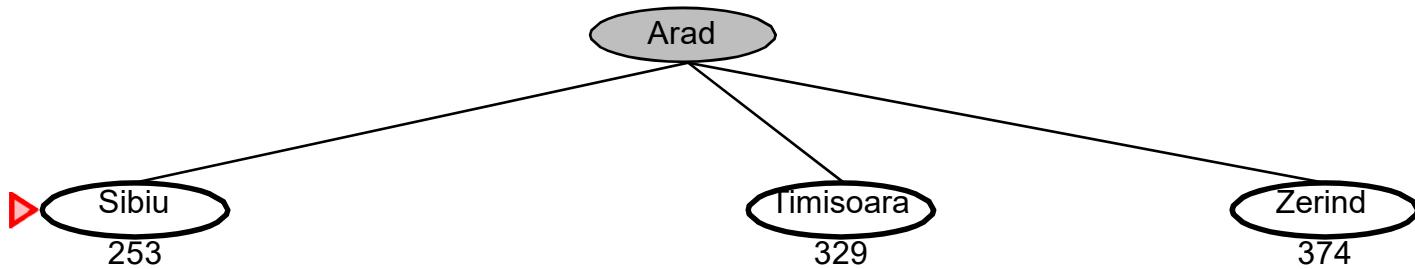
E.g.,  $h_{SLD}(n)$  = straight-line distance from  $n$  to Bucharest

Greedy search expands the node that appears to be closest to goal

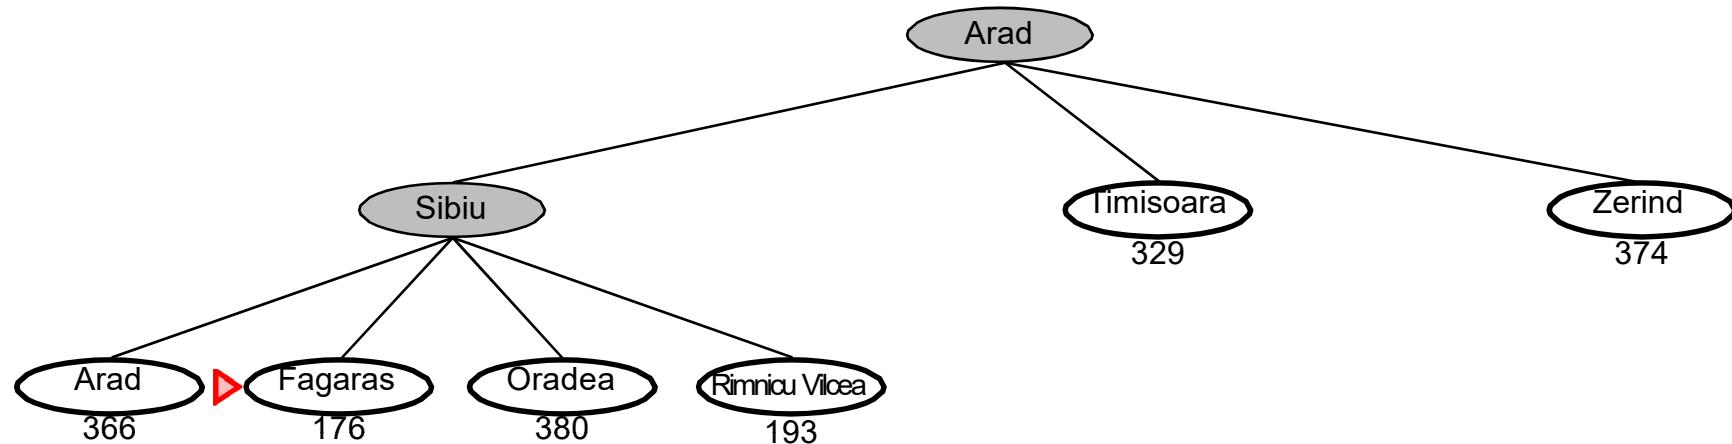
## Greedy search example



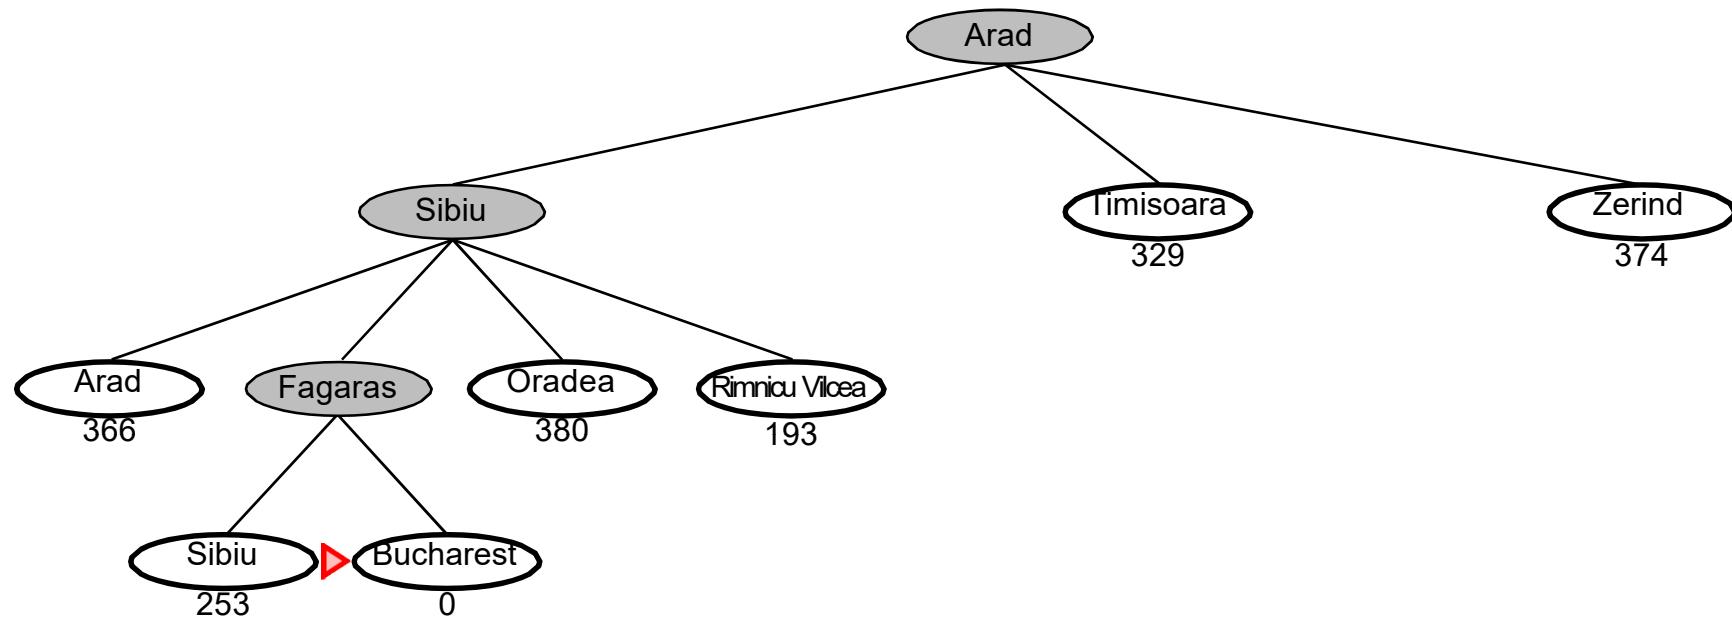
## Greedy search example



## Greedy search example



## Greedy search example



## Properties of greedy search

Complete??

## Properties of greedy search

Complete?? No—can get stuck in loops, e.g., with Oradea as goal,

Iasi → Neamt → Iasi → Neamt →

Complete in finite space with repeated-state checking

Time??

## Properties of greedy search

Complete?? No—can get stuck in loops, e.g.,

Iasi → Neamt → Iasi → Neamt →

Complete in finite space with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic improvement

Space??

## Properties of greedy search

Complete?? No—can get stuck in loops,

e.g., Iasi → Neamt → Iasi →  
Neamt →

Complete in finite space with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic

improvement Space??  $O(b^m)$ —keeps all nodes in memory

Optimal??

## Properties of greedy search

Complete?? No—can get stuck in loops,

e.g., Iasi → Neamt → Iasi →  
Neamt →

Complete in finite space with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic

improvement Space??  $O(b^m)$ —keeps all nodes in memory

Optimal?? No

## A\* search

Idea: avoid expanding paths that are already expensive

Evaluation function  $f(n) = g(n) + h(n)$

$g(n)$  = cost so far to reach  $n$

$h(n)$  = estimated cost to goal from  $n$

$f(n)$  = estimated total cost of path through  $n$  to goal

A\* search uses an **admissible** heuristic

i.e.,  $h(n) \leq h^*(n)$  where  $h^*(n)$  is the **true** cost from  $n$ . (Also require  $h(n) \geq 0$ , so  $h(G) = 0$  for any goal  $G$ .)

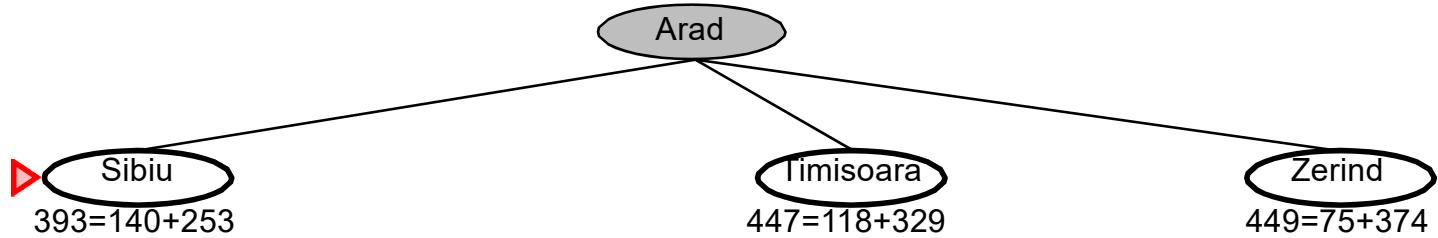
E.g.,  $h_{SLD}(n)$  never overestimates the actual road

distance Theorem: A\* search is optimal

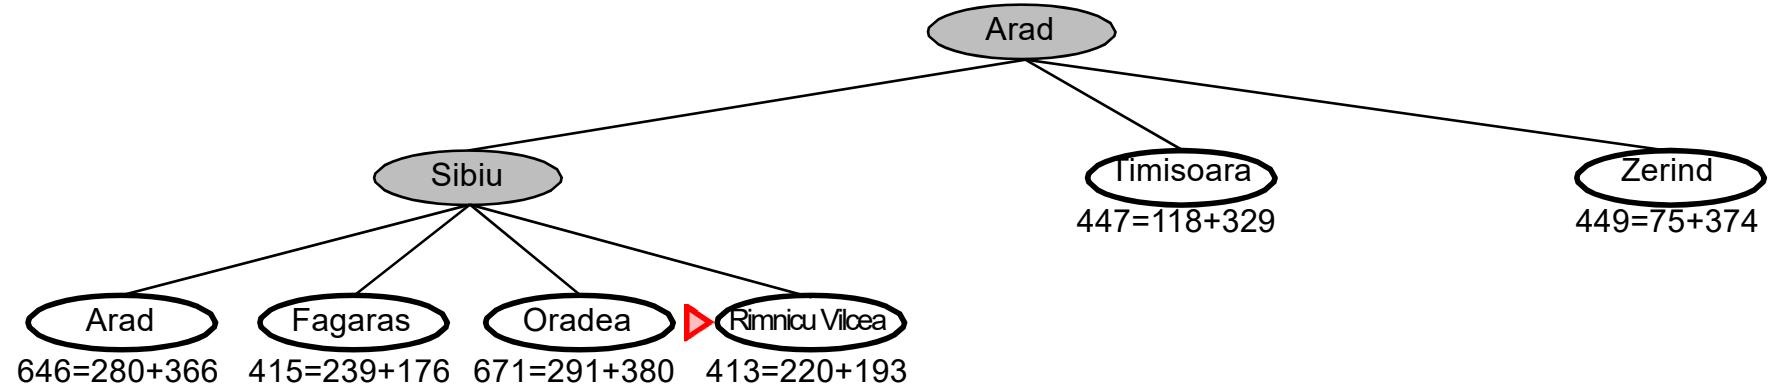
## A\* search example

► Arad  
 $366=0+366$

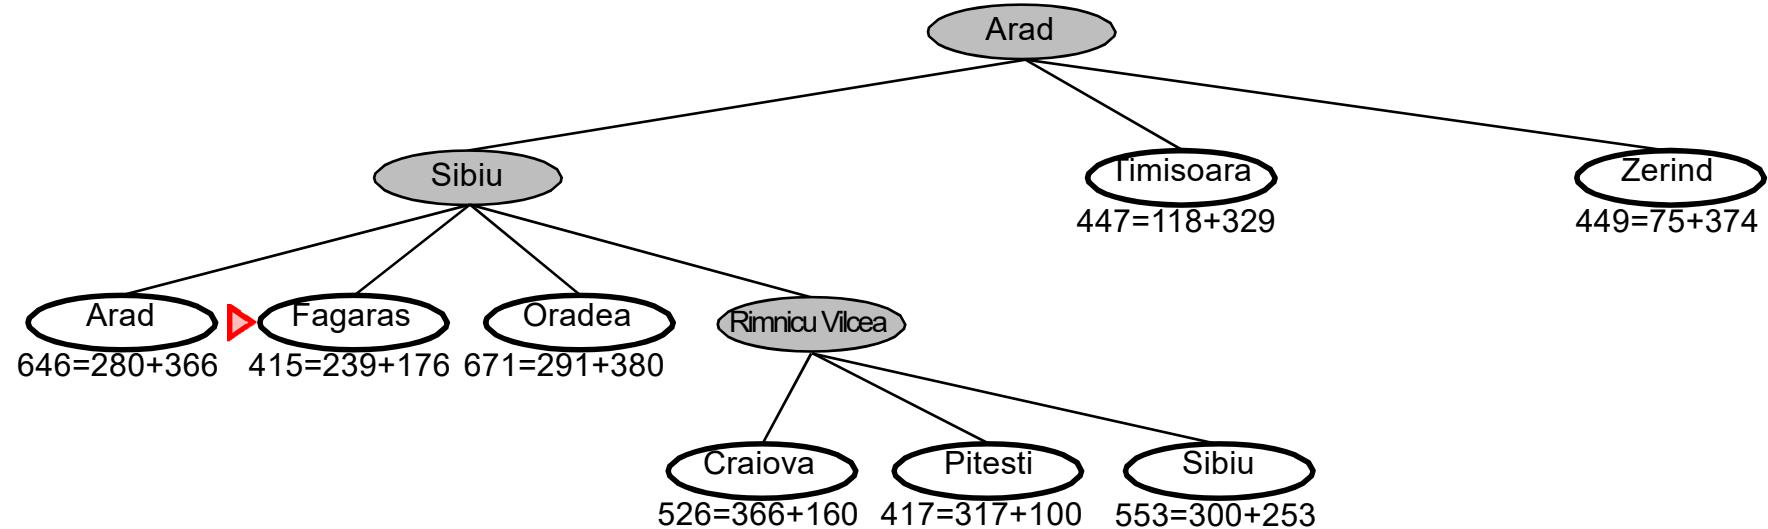
## A\* search example



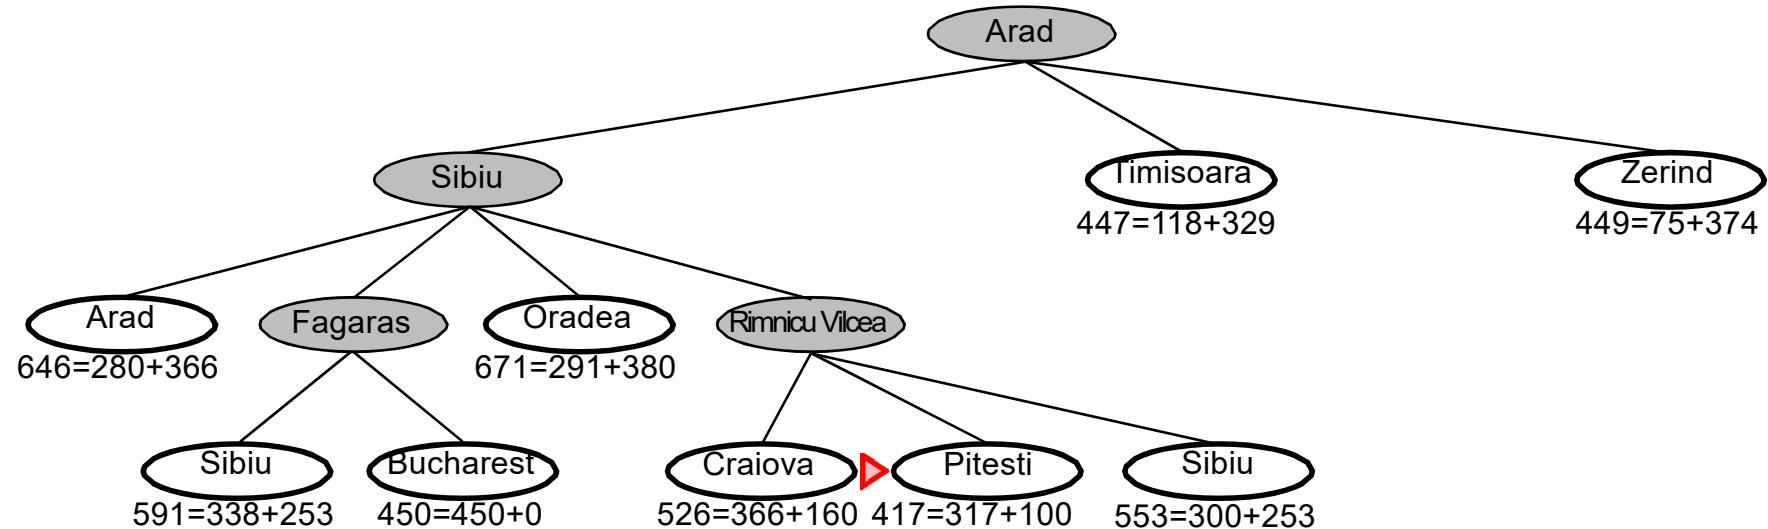
## A\* search example



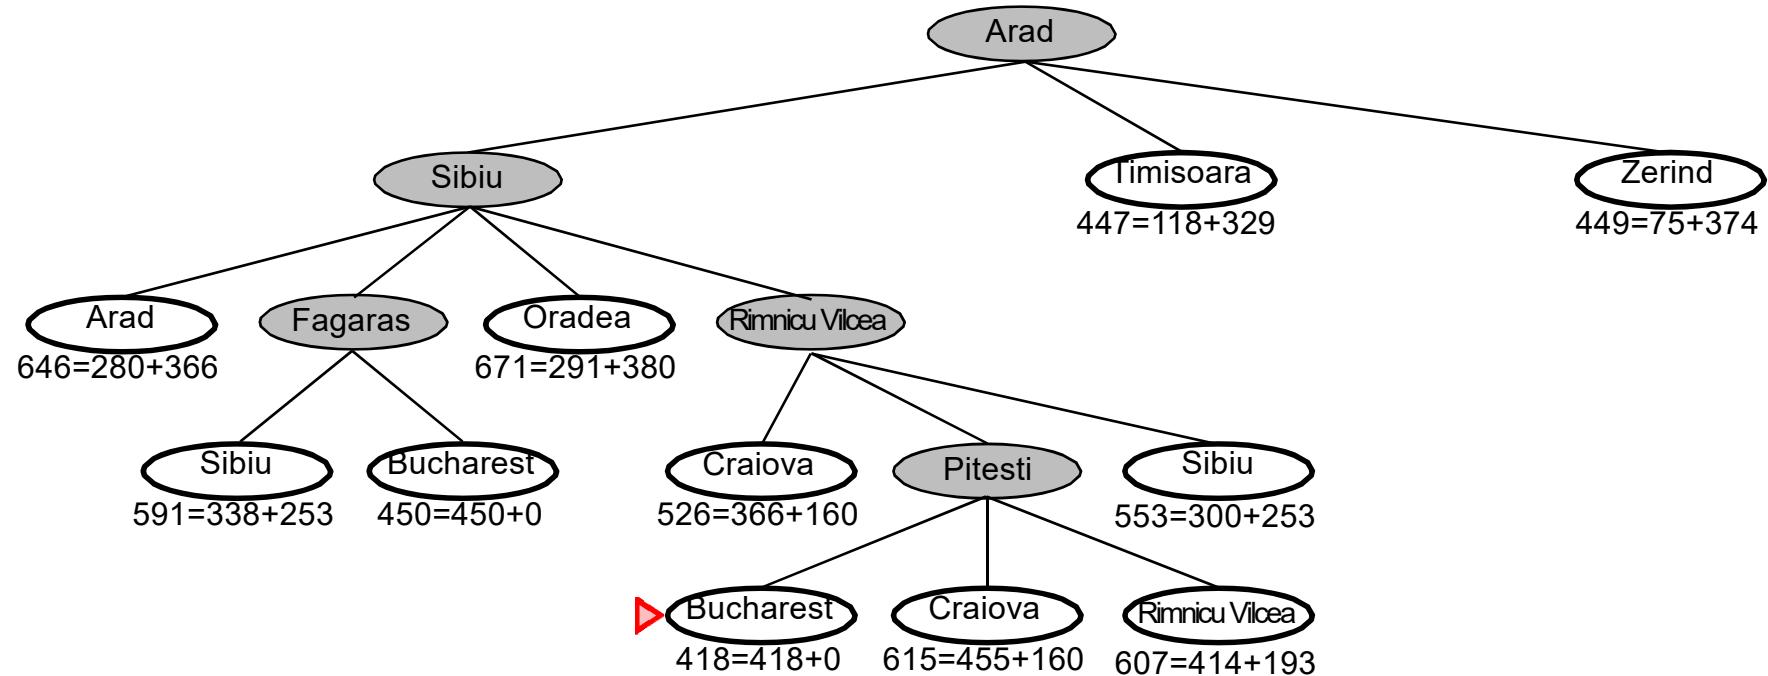
## A\* search example



## A\* search example

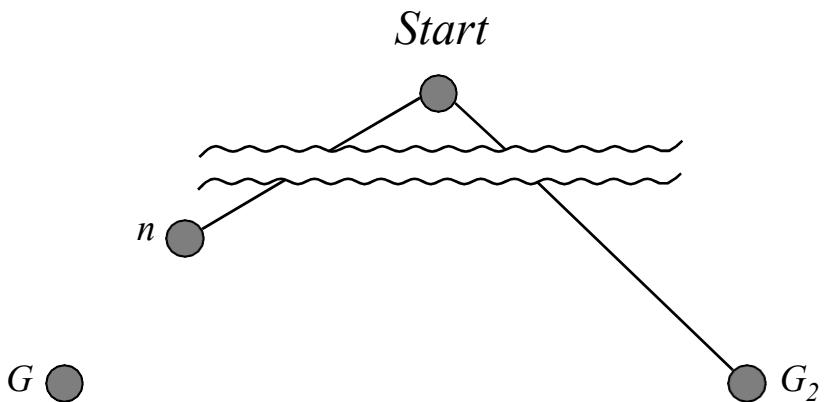


## A\* search example



## Optimality of A\* (standard proof)

Suppose some suboptimal goal  $G_2$  has been generated and is in the queue. Let  $n$  be an unexpanded node on a shortest path to an optimal goal  $G_1$ .



$$\begin{aligned}
 f(G_2) &= g(G_2) && \text{since } h(G_2) = 0 \\
 &> g(G_1) && \text{since } G_2 \text{ is suboptimal} \\
 &\geq f(n) && \text{since } h \text{ is admissible}
 \end{aligned}$$

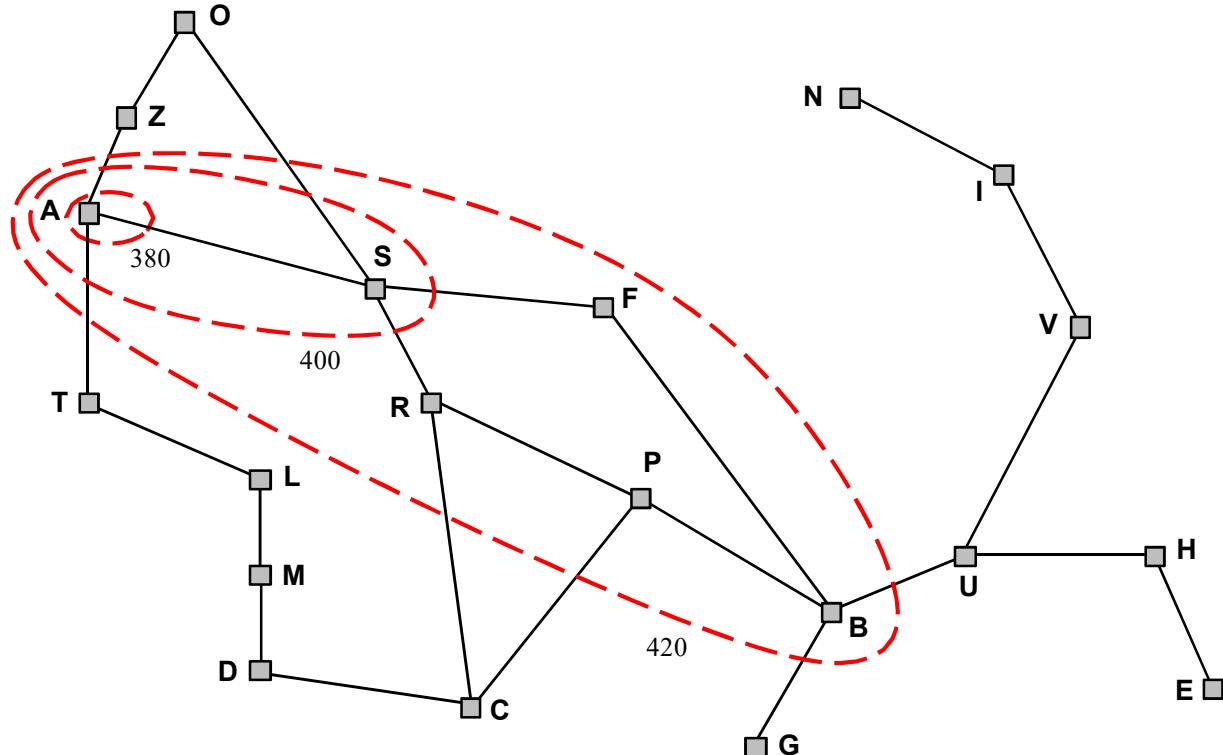
Since  $f(G_2) > f(n)$ , A\* will never select  $G_2$  for expansion

## Optimality of A\* (more useful)

**Lemma:** A\* expands nodes in order of increasing  $f$  value\*

Gradually adds “ $f$ -contours” of nodes (cf. breadth-first adds layers)

Contour  $i$  has all nodes with  $f = f_i$ , where  $f_i < f_{i+1}$



## Properties of A\*

Complete??

## Properties of A\*

Complete?? Yes, unless there are infinitely many nodes with  $f \leq f(G)$

Time??

## Properties of A\*

Complete?? Yes, unless there are infinitely many nodes with  $f \leq f(G)$

Time?? Exponential in [relative error in  $h \times$  length of soln.]

Space??

## Properties of A\*

Complete?? Yes, unless there are infinitely many nodes with  $f \leq f(G)$

Time?? Exponential in [relative error in  $h \times$  length of soln.]

Space?? Keeps all nodes in memory

Optimal??

## Properties of A\*

Complete?? Yes, unless there are infinitely many nodes with  $f \leq f(G)$

Time?? Exponential in [relative error in  $h \times$  length of soln.]

Space?? Keeps all nodes in memory

Optimal?? Yes—cannot expand  $f_{i+1}$  until  $f_i$  is finished

A\* expands all nodes with  $f(n) < C^*$

A\* expands some nodes with  $f(n) = C^*$

A\* expands no nodes with  $f(n) > C^*$

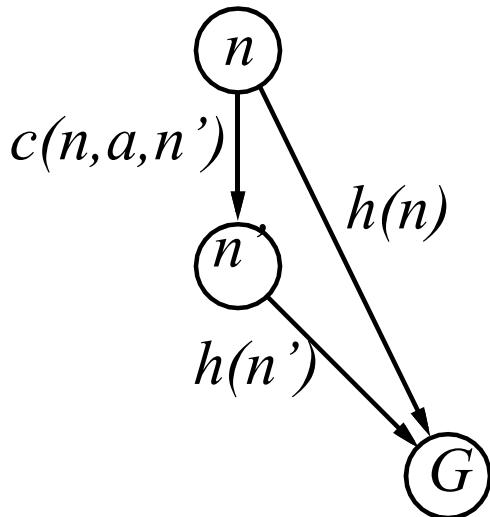
## Proof of lemma: Consistency

A heuristic is **consistent** if

$$h(n) \leq c(n, a, n') + h(n')$$

If  $h$  is consistent, we have

$$\begin{aligned} f(n) &= g(n) + h(n) \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$



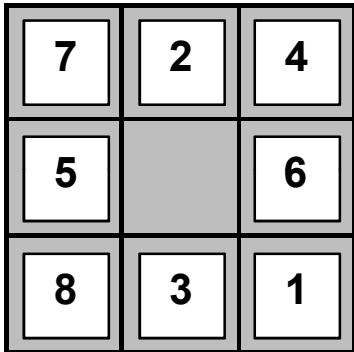
I.e.,  $f(n)$  is nondecreasing along any path.

## Admissible heuristics

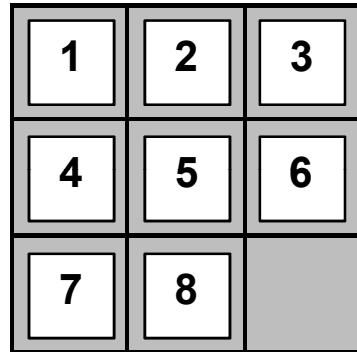
E.g., for the 8-puzzle:

$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total **Manhattan distance**  
(i.e., no. of squares from desired location of each tile)



Start State



Goal State

$$\underline{h_1(S) = ??}$$

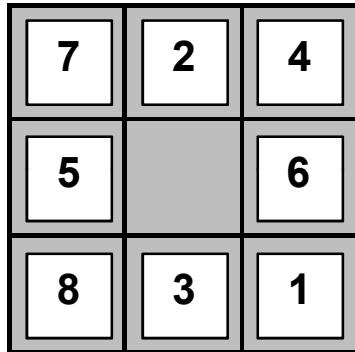
$$\underline{h_2(S) = ??}$$

## Admissible heuristics

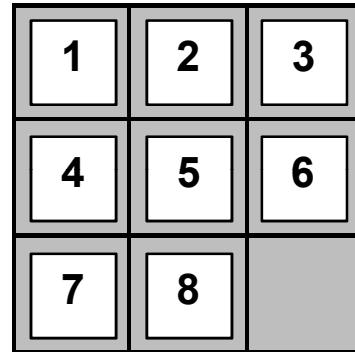
E.g., for the 8-puzzle:

$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total **Manhattan** distance  
(i.e., no. of squares from desired location of each tile)



Start State



Goal State

$$h_1(S) = ?? \ 6$$

$$\underline{h_2(S) = ??} \ 4+0+3+3+1+0+2+1 = 14$$

## Dominance

If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible), then  $h_2$  dominates  $h_1$  and is better for search

Typical search costs:

$d = 14$    IDS = 3,473,941 nodes

$A^*(h_1)$  = 539 nodes

$A^*(h_2)$  = 113 nodes

$d = 24$    IDS  $\approx$  54,000,000,000 nodes

$A^*(h_1)$  = 39,135 nodes

$A^*(h_2)$  = 1,641 nodes

Given any admissible heuristics  $h_a$ ,  $h_b$ ,

$$h(n) = \max(h_a(n), h_b(n))$$

is also admissible and dominates  $h_a$ ,  $h_b$

## Relaxed problems

Admissible heuristics can be derived from the **exact** solution cost of a **relaxed** version of the problem

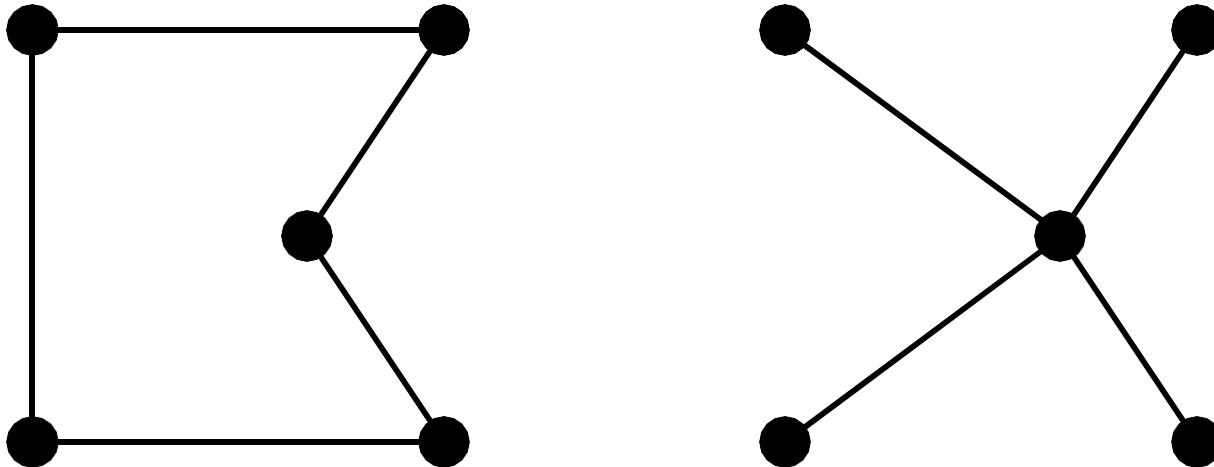
If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then  $h_1(n)$  gives the shortest solution

If the rules are relaxed so that a tile can move to **any adjacent square**, then  $h_2(n)$  gives the shortest solution

Key point: the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem

## Relaxed problems contd.

Well-known example: **travelling salesperson problem** (TSP) Find the shortest tour visiting all cities exactly once



Minimum spanning tree can be computed in  $O(n^2)$

and is a lower bound on the shortest (open) tour

## Summary

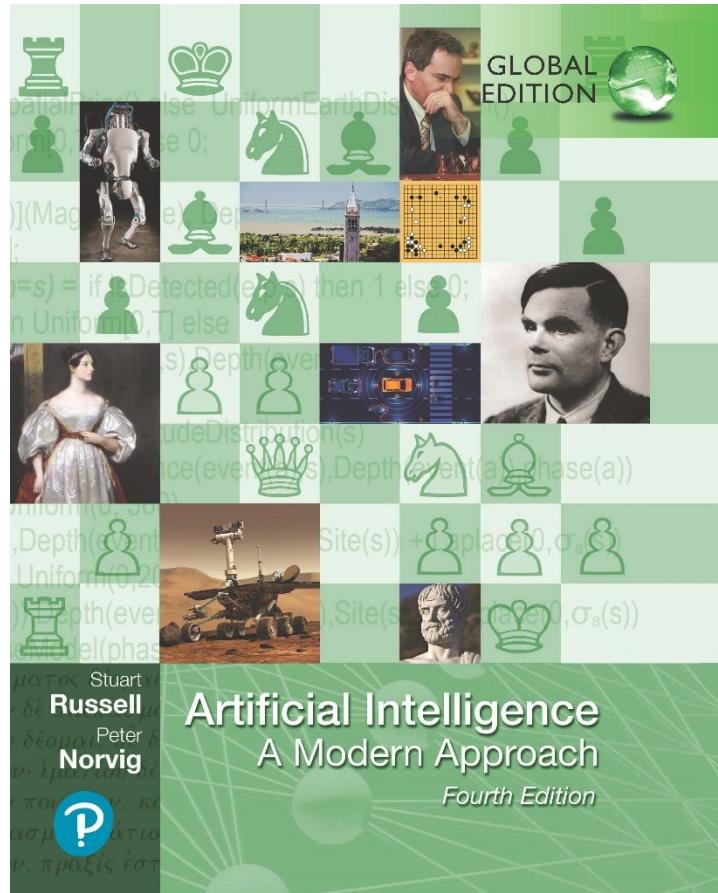
A problem consists of five parts: the **initial state**, a set of **actions**, a **transition model** describing the results of those actions, a set of **goal states**, and an **action cost function**.

**Uninformed search** methods have access only to the **problem definition**. Algorithms build a search tree in an attempt to find a solution.

**Informed search** methods have access to a **heuristic** function  $h(n)$  that estimates the cost of a solution from  $n$ .

# Artificial Intelligence: A Modern Approach

Fourth Edition, Global Edition



## Chapter 4

### Search in Complex Environments



# Lecture Presentations: Artificial Intelligence

**Adapted from:**

"Artificial Intelligence: A Modern Approach, Global Edition",  
4th Edition by Stuart Russell and Peter Norvig © 2021  
Pearson Education.

**Adapted for** educational use at ACE Engineering College.  
Some slides customized by Mr. Shafakhatullah Khan  
Mohammed, Assistant Professor @ ACE Engineering College.  
For instructional use only. Not for commercial distribution.

# Outline

- ◆ Local Search and Optimization Problems
  - ◆ Hill-climbing
  - ◆ Simulated annealing
  - ◆ Genetic algorithms
- ◆ Local search in continuous spaces
- ◆ Search with Nondeterministic Actions
- ◆ Search in Partially Observable Environments

## Local Search and Optimization Problems

In many optimization problems, **path** is irrelevant; the goal state itself is the solution

Then state space = set of “complete” configurations; find **optimal** configuration, e.g., TSP  
or, find configuration satisfying constraints, e.g., timetable

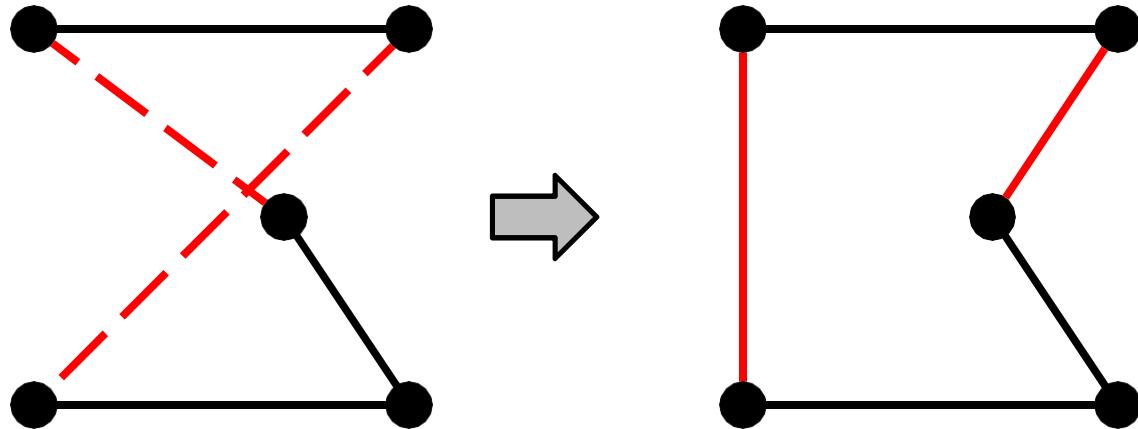
In such cases, one can use **iterative improvement** algorithms; keep a single “current” state, try to improve it

Local search algorithms operate by searching from a start state to neighboring states, without keeping track of the paths, nor the set of states that have been reached.

They are not systematic—they might never explore a portion of the search space where a solution actually resides.

## Example: Travelling Salesperson Problem

Start with any complete tour, perform pairwise exchanges

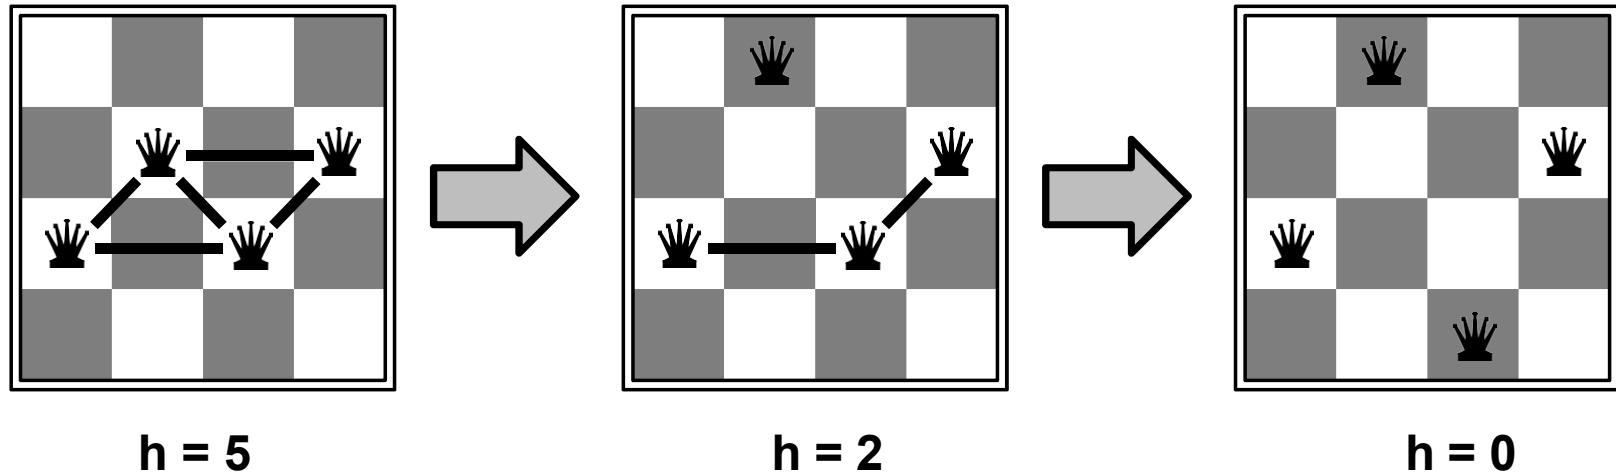


Variants of this approach get within 1% of optimal very quickly with thousands of cities

## Example: $n$ -queens

Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal

Move a queen to reduce number of conflicts



Almost always solves  $n$ -queens problems almost instantaneously for very large  $n$ , e.g.,  $n = 1\text{million}$

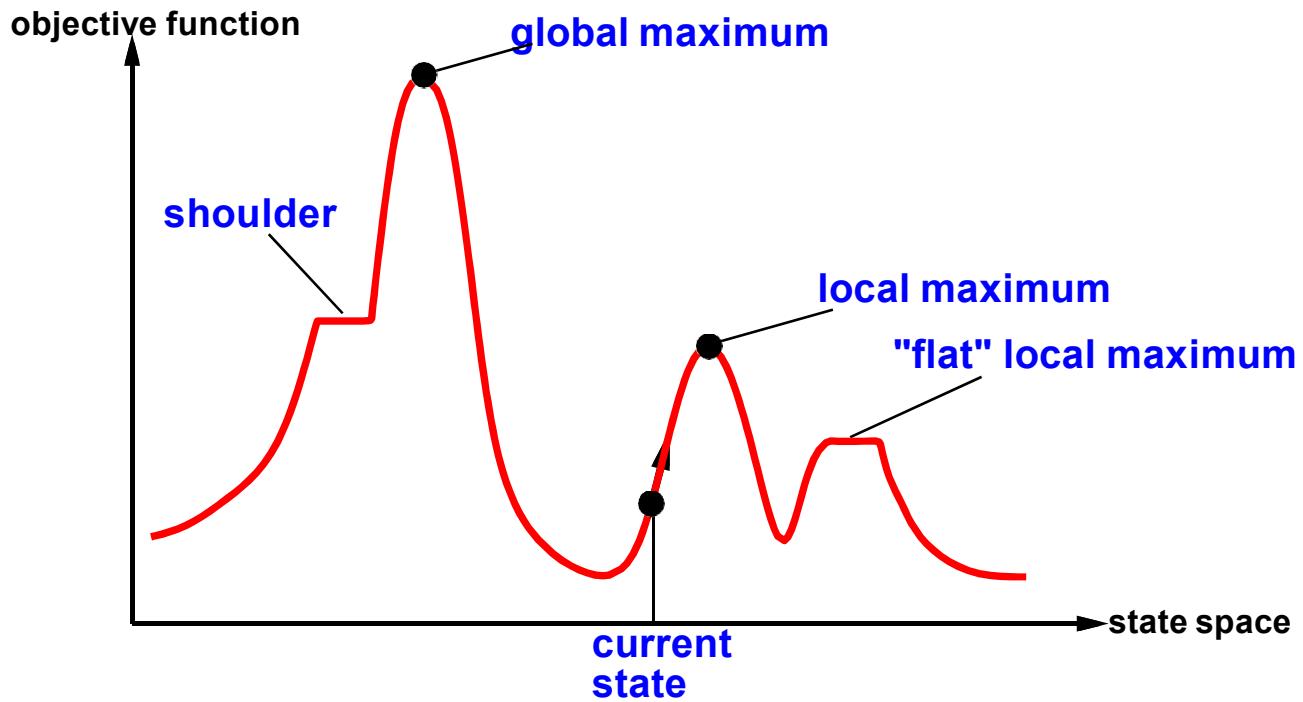
# Hill-climbing (or gradient ascent/descent)

“Like climbing Everest in thick fog with amnesia”

```
function Hill-Climbing(problem) returns a state that is a local maximum
    inputs: problem, a problem
    local variables: current, a node
                  neighbor, a node
    current  $\leftarrow$  Make-Node(Initial-State [problem])
    loop do
        neighbor  $\leftarrow$  a highest-valued successor of current
        if Value[neighbor]  $\leq$  Value[current] then return State[current]
        current  $\leftarrow$  neighbor
    end
```

## Hill-climbing contd.

Useful to consider state space landscape



Random-restart hill climbing overcomes local maxima—trivially complete

Random sideways moves 😊 escape from shoulders 😟 loop on flat maxima

# Simulated annealing

Idea: escape local maxima by allowing some “bad” moves  
 but gradually decrease their size and frequency

```

function Simulated-Annealing(problem, schedule) returns a solution state
    inputs: problem, a problem
                schedule, a mapping from time to “temperature”
    local variables: current, a node
                    next, a node
                    T, a “temperature” controlling prob. of downward steps

    current  $\leftarrow$  Make-Node(Initial-State [problem])
    for t  $\leftarrow$  1 to  $\infty$  do
        T  $\leftarrow$  schedule[t]
        if T = 0 then return current
        next  $\leftarrow$  a randomly selected successor of current
         $\Delta E \leftarrow$  Value[next] – Value[current]
        if  $\Delta E > 0$  then current  $\leftarrow$  next
        else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 

```

## Properties of simulated annealing

At fixed “temperature”  $T$ , state occupation probability reaches Boltzman distribution

$$p(x) = ae^{\frac{E(x)}{kT}}$$

$T$  decreased slowly enough  $\Rightarrow$  always reach best state  $x^*$   
because  $e^{\frac{E(x^*)}{kT}} / e^{\frac{E(x)}{kT}} = e^{\frac{E(x^*) - E(x)}{kT}} \diamondsuit 1$  for small  $T$

Is this necessarily an interesting guarantee??

Devised by Metropolis et al., 1953, for physical process modelling

Widely used in VLSI layout, airline scheduling, etc.

## Local beam search

Idea: keep  $k$  states instead of 1; choose top  $k$  of all their successors

Not the same as  $k$  searches run in parallel!

Searches that find good states recruit other searches to join them

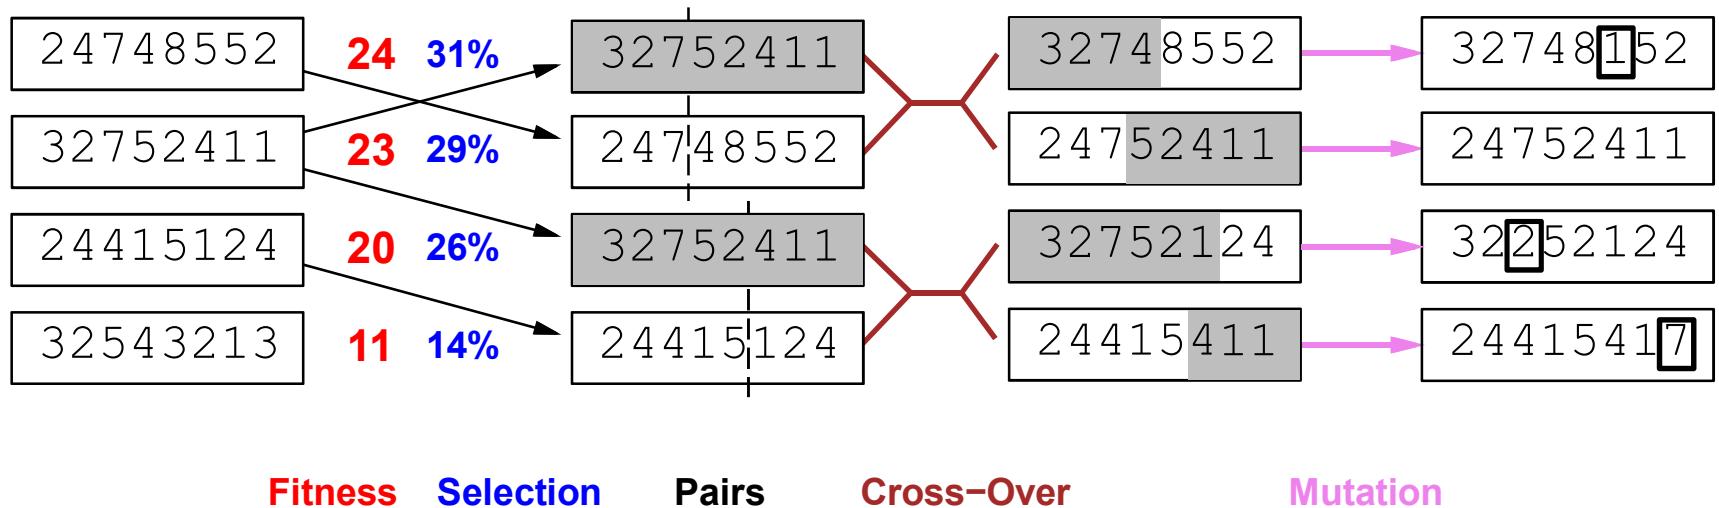
Problem: quite often, all  $k$  states end up on same local hill

Idea: choose  $k$  successors randomly, biased towards good ones

Observe the close analogy to natural selection!

# Genetic algorithms

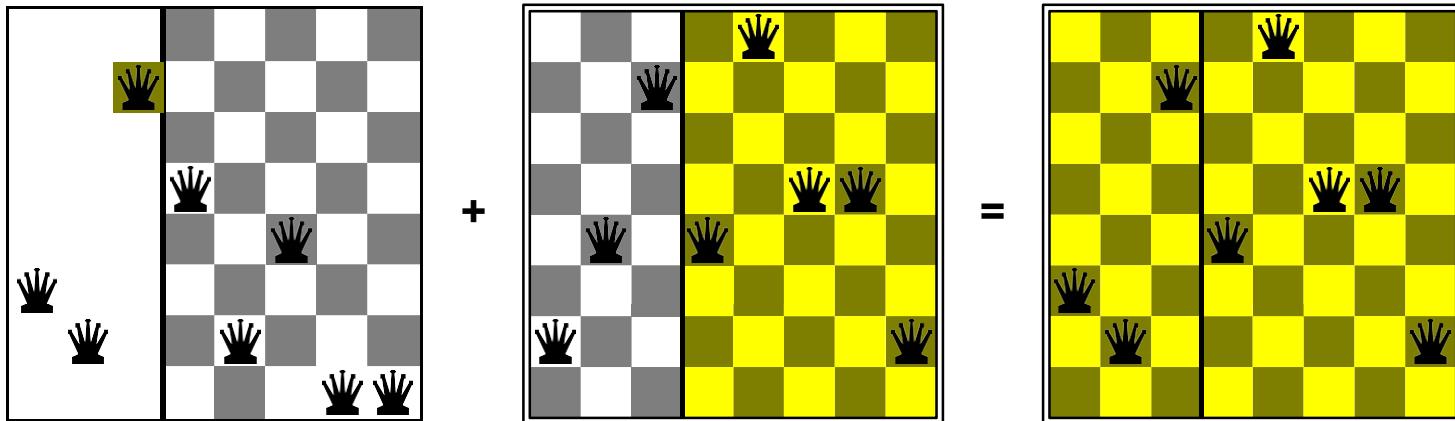
= stochastic local beam search + generate successors from **pairs** of states



## Genetic algorithms contd.

GAs require states encoded as strings (GPs use programs)

Crossover helps iff substrings are meaningful components



GAs / = evolution: e.g., real genes encode replication machinery!

## Continuous state spaces

Suppose we want to site three airports in Romania:

- 6-D state space defined by  $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- objective function  $f(x_1, y_1, x_2, y_2, x_3, y_3) =$   
sum of squared distances from each city to nearest airport

Discretization methods turn continuous space into discrete space,  
e.g., empirical gradient considers  $\pm \delta$  change in each coordinate

Gradient methods compute

$$\nabla f = \left[ \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right]$$

to increase/reduce  $f$ , e.g., by  $x \leftarrow x + \alpha \nabla f(x)$

Sometimes can solve for  $\nabla f(x) = 0$  exactly (e.g., with one city).

Newton–Raphson (1664, 1690) iterates  $x \leftarrow x - H_f^{-1}(x) \nabla f(x)$   
to solve  $\nabla f(x) = 0$ , where  $H_{ij} = \partial^2 f / \partial x_i \partial x_j$

## Search with Nondeterministic Actions

**Agent** doesn't know the state its transitioned to **after action**, the environment is **nondeterministic**.

Rather, it will know the possible states it will be in, which is called "**belief state**"

Examples:

- The **erratic vacuum world** (if-then-else) steps. If statement tests to know the current state.
- **AND-OR** search trees. Two possible actions (**OR nodes**). Branching that happens from a choice (**AND nodes**).
- **Try, try again.** A cyclic plan where minimum condition (every leaf = goal state & reachable from other points in the plan)

## Search in Partially Observable Environments

**Problem of partial observability**, where the agent's percepts are not enough to pin down the exact state.

**Searching with no observation:** Agent's percepts provide **no information at all**, sensorless problem (or a conformant problem).

- Solution: sequence of actions, not a conditional plan

**Searching in partially observable environments** requires a function that **monitors** or **estimates** the environment to maintain the belief state.

## Summary

**Local search methods** keep only a **small number of states** in memory that are useful for optimization.

In **nondeterministic environments**, agents can apply **AND–OR search** to generate contingency plans that reach the goal regardless of which outcomes occur during execution.

**Belief-state** is the set of **possible states** that the agent is in for **partially observable environments**.

**Standard search algorithms** can be applied directly to belief-state space to solve **sensorless** problems.