


# Artificial Intelligence: A Modern Approach

Fourth Edition, Global Edition



## Chapter 7

### Logical agents

# Lecture Presentations: Artificial Intelligence

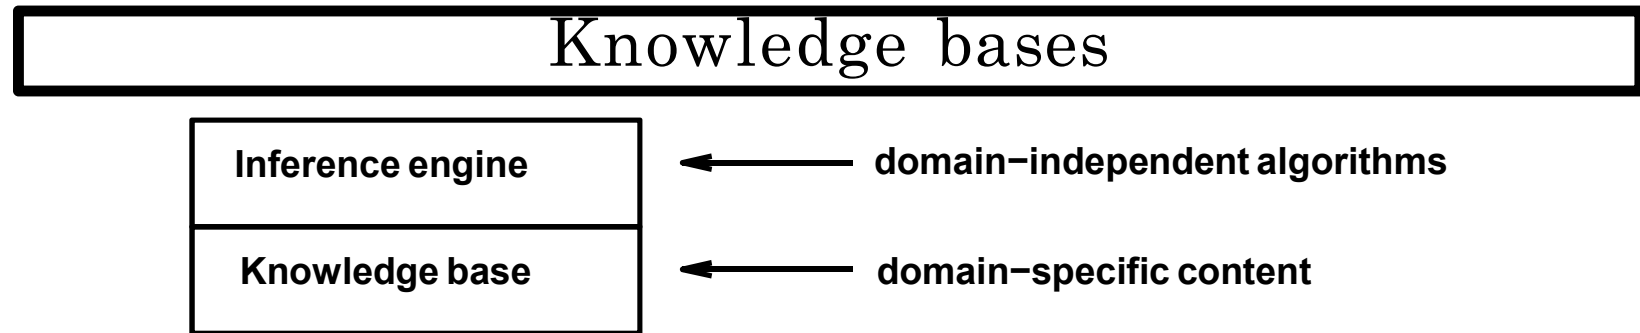
## Adapted from:

"Artificial Intelligence: A Modern Approach, Global Edition",  
4th Edition by Stuart Russell and Peter Norvig © 2021  
Pearson Education.

**Adapted for** educational use at ACE Engineering College.  
Some slides customized by Mr. Shafakhatullah Khan  
Mohammed, Assistant Professor @ ACE Engineering College.  
For instructional use only. Not for commercial distribution.

# Outline

- ◆ Knowledge-based agents
- ◆ Wumpus world
- ◆ Logic in general—models and entailment
- ◆ Propositional (Boolean) logic
- ◆ Equivalence, validity, satisfiability
- ◆ Inference rules and theorem proving
  - resolution
  - forward chaining
  - backward chaining
- ◆ Effective Propositional Model Checking



Knowledge base = set of sentences in a formal language

Declarative approach to building an agent (or other system):

**Tell** it what it needs to know

Then it can **Ask** itself what to do—answers should follow from the KB

Agents can be viewed at the knowledge level

i.e., **what they know**, regardless of how implemented

Or at the implementation level

i.e., data structures in KB and algorithms that manipulate them

## A simple knowledge-based agent

```
function KB-Agent(percept) returns an action
  static: KB, a knowledge base
          t, a counter, initially 0, indicating time
  Tell(KB, Make-Percept-Sentence(percept, t))
  action ← Ask(KB, Make-Action-Query(t))
  Tell(KB, Make-Action-Sentence(action, t))
  t ← t + 1
  return action
```

The agent must be able to:

- Represent states, actions, etc.

- Incorporate new percepts

- Update internal representations of the world

- Deduce hidden properties of the world

- Deduce appropriate actions

# Wumpus World PEAS description

## Performance measure

gold +1000, death -1000

-1 per step, -10 for using the arrow

## Environment

Squares adjacent to wumpus are smelly

Squares adjacent to pit are breezy


Glitter iff gold is in the same square

Shooting kills wumpus if you are facing it

Shooting uses up the only arrow

Grabbing picks up gold if in same square

Releasing drops the gold in same square



**Actuators** Left turn, Right turn,

Forward, Grab, Release, Shoot

**Sensors** Breeze, Glitter, Smell

# Wumpus world characterization

Observable??

# Wumpus world characterization

Observable?? No—only **local** perception

Deterministic??



# Wumpus world characterization

Observable?? No—only **local** perception

Deterministic?? Yes—outcomes exactly specified

Episodic??

# Wumpus world characterization

Observable?? No—only **local** perception

Deterministic?? Yes—outcomes exactly specified

Episodic?? No—sequential at the level of actions

Static??

# Wumpus world characterization

Observable?? No—only **local** perception

Deterministic?? Yes—outcomes exactly specified

Episodic?? No—sequential at the level of actions

Static?? Yes—Wumpus and Pits do not move

Discrete??

# Wumpus world characterization

Observable?? No—only **local** perception

Deterministic?? Yes—outcomes exactly specified

Episodic?? No—sequential at the level of actions

Static?? Yes—Wumpus and Pits do not move

Discrete?? Yes

Single-agent??

# Wumpus world characterization

Observable?? No—only **local** perception

Deterministic?? Yes—outcomes exactly specified

Episodic?? No—sequential at the level of actions

Static?? Yes—Wumpus and Pits do not move


Discrete?? Yes

Single-agent?? Yes—Wumpus is essentially a natural feature


## Exploring a wumpus world

OK			
OK <div>A</div>	OK		

# Exploring a wumpus world




# Exploring a wumpus world







# Exploring a wumpus world




# Exploring a wumpus world




# Exploring a wumpus world




# Exploring a wumpus world



# Exploring a wumpus world




## Other tight spots



Breeze in (1,2) and (2,1)  
 $\Rightarrow$  no safe actions

Assuming pits uniformly distributed,  
 (2,2) has pit w/ prob 0.86, vs. 0.31

Smell in (1,1)  
 $\Rightarrow$  cannot move



Can use a strategy of **coercion**:

shoot straight ahead

wumpus was there  $\Rightarrow$  dead  $\Rightarrow$  safe

wumpus wasn't there  $\Rightarrow$  safe

# Logic in general

**Logics** are formal languages for representing information such that conclusions can be drawn

**Syntax** defines the sentences in the language

**Semantics** define the “meaning” of sentences;  
i.e., define **truth** of a sentence in a world

E.g., the language of arithmetic

$x + 2 \geq y$  is a sentence;  $x^2 + y >$  is not a sentence

$x + 2 \geq y$  is true iff the number  $x + 2$  is no less than the number  $y$

$x + 2 \geq y$  is true in a world where  $x = 7$ ,  $y = 1$

$x + 2 \geq y$  is false in a world where  $x = 0$ ,  $y = 6$

# Entailment

Entailment means that one thing follows from another:

$$KB \models a$$

Knowledge base  $KB$  entails sentence  $a$

if and only if

$a$  is true in all worlds where  $KB$  is true

E.g., the KB containing “the Giants won” and “the Reds won” entails “Either the Giants won or the Reds won”

E.g.,  $x + y = 4$  entails  $4 = x + y$

Entailment is a relationship between sentences (i.e., syntax) that is based on semantics

Note: brains process syntax (of some sort)



# Models


Logicians typically think in terms of **models**, which are formally structured worlds with respect to which truth can be evaluated

We say  **$m$**  is a **model** of a sentence  **$a$**  if  **$a$**  is true in  **$m$**

**$M(a)$**  is the set of all models of  **$a$**

Then  **$KB \models a$**  if and only if  **$M(KB) \subseteq M(a)$**

E.g.  **$KB$**  = Giants won and Reds won  
 **$a$**  = Giants won




# Entailment in the wumpus world









Situation after detecting nothing in [1,1],  
moving right, breeze in [2,1]

Consider possible models for ?s  
assuming only pits


3 Boolean choices  $\Rightarrow$  8 possible models



# Wumpus models




# Wumpus models



*KB* = wumpus-world rules + observations


# Wumpus models



$KB$  = wumpus-world rules + observations


$\alpha_1$  = "[1,2] is safe",  $KB \models \alpha_1$ , proved by model checking

# Wumpus models



*KB* = wumpus-world rules + observations

# Wumpus models



$KB$  = wumpus-world rules + observations

$a_2$  = "[2,2] is safe",  $KB \models a_2$

# Inference

$KB \vdash_i a$  = sentence  $a$  can be derived from  $KB$  by procedure  $i$

Consequences of  $KB$  are a haystack;  $a$  is a needle.

Entailment = needle in haystack; inference = finding it

Soundness:  $i$  is sound if

whenever  $KB \vdash_i a$ , it is also true that  $KB \models a$

Completeness:  $i$  is complete if

whenever  $KB \models a$ , it is also true that  $KB \vdash_i a$

Preview: we will define a logic (first-order logic) which is expressive enough to say almost anything of interest, and for which there exists a sound and complete inference procedure.

That is, the procedure will answer any question whose answer follows from what is known by the  $KB$ .



## Propositional logic: Syntax

Propositional logic is the simplest logic—illustrates basic ideas

The proposition symbols  $P_1, P_2$  etc are sentences

If  $S$  is a sentence,  $\neg S$  is a sentence (negation)

If  $S_1$  and  $S_2$  are sentences,  $S_1 \wedge S_2$  is a sentence (conjunction)

If  $S_1$  and  $S_2$  are sentences,  $S_1 \vee S_2$  is a sentence (disjunction)

If  $S_1$  and  $S_2$  are sentences,  $S_1 \Rightarrow S_2$  is a sentence (implication)

If  $S_1$  and  $S_2$  are sentences,  $S_1 \Leftrightarrow S_2$  is a sentence (biconditional)

# Propositional logic: Semantics

Each model specifies true/false for each proposition symbol

E.g.  $P_{1,2}$   $P_{2,2}$   $P_{3,1}$   
*true true false*

(With these symbols, 8 possible models, can be enumerated automatically.)

Rules for evaluating truth with respect to a model  $m$ :

$\neg S$ is true iff	$S$ is false
$S_1 \wedge S_2$ is true iff	$S_1$ is true <b>and</b> $S_2$ is true
$S_1 \vee S_2$ is true iff	$S_1$ is true <b>or</b> $S_2$ is true
$S_1 \Rightarrow S_2$ is true iff	$S_1$ is false <b>or</b> $S_2$ is true
i.e., is false iff	$S_1$ is true <b>and</b> $S_2$ is false
$S_1 \Leftrightarrow S_2$ is true iff	$S_1 \Rightarrow S_2$ is true <b>and</b> $S_2 \Rightarrow S_1$ is true

Simple recursive process evaluates an arbitrary sentence, e.g.,

$\neg P_{1,2} \wedge (P_{2,2} \vee P_{3,1}) = \text{true} \wedge (\text{false} \vee \text{true}) = \text{true} \wedge \text{true} = \text{true}$

# Truth tables for connectives

$P$	$Q$	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
false	false	true	false	false	true	true
false	true	true	false	true	true	false
true	false	false	false	true	false	false
true	true	false	true	true	true	true

## Wumpus world sentences

Let  $P_{i,j}$  be true if there is a pit in  $[i,j]$ .

Let  $B_{i,j}$  be true if there is a breeze in  $[i,j]$ .

$$\neg P_{1,1}$$

$$\neg B_{1,1}$$

$$B_{2,1}$$

“Pits cause breezes in adjacent squares”

## Wumpus world sentences

Let  $P_{i,j}$  be true if there is a pit in  $[i,j]$ .

Let  $B_{i,j}$  be true if there is a breeze in  $[i,j]$ .

$$\neg P_{1,1}$$

$$\neg B_{1,1}$$

$$B_{2,1}$$

“Pits cause breezes in adjacent squares”

$$B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$$

$$B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$$

“A square is breezy if and only if there is an adjacent pit”

## Truth tables for inference

$B_{1,1}$	$B_{2,1}$	$P_{1,1}$	$P_{1,2}$	$P_{2,1}$	$P_{2,2}$	$P_{3,1}$	$R_1$	$R_2$	$R_3$	$R_4$	$R_5$	$KB$
false	false	false	false	false	false	false	true	true	true	true	false	false
false	false	false	false	false	false	true	true	true	false	true	false	false
.	.	.	.	.	.	.	.	.	.	.	.	.
false	true	false	false	false	false	false	true	true	false	true	true	false
false	true	false	false	false	false	true	true	true	true	true	true	<u>true</u>
false	true	false	false	false	false	false	true	true	true	true	true	<u>true</u>
false	true	false	false	false	false	true	true	true	true	true	true	<u>true</u>
false	true	false	false	false	false	false	true	true	true	true	true	<u>true</u>
false	true	false	false	true	false	false	true	false	false	true	true	false
.	.	.	.	.	.	.	.	.	.	.	.	.
true	true	true	true	true	true	true	false	true	true	false	true	false

Enumerate rows (different assignments to symbols),  
if  $KB$  is true in row, check that  $a$  is too

# Inference by enumeration

Depth-first enumeration of all models is sound and complete

```

function TT-Entails?(KB, a) returns true or false
  inputs: KB, the knowledge base, a sentence in propositional logic
         a, the query, a sentence in propositional logic
  symbols ← a list of the proposition symbols in KB and a
  return TT-Check-All(KB, a, symbols, [])

function TT-Check-All(KB, a, symbols, model) returns true or false
  if Empty?(symbols) then
    if PL-True?(KB, model) then return PL-True?(a, model)
    else return true
  else do
    P ← First(symbols); rest ← Rest(symbols)
    return TT-Check-All(KB, a, rest, Extend(P, true, model)) and
           TT-Check-All(KB, a, rest, Extend(P, false, model))
  
```

$O(2^n)$  for  $n$  symbols; problem is **co-NP-complete**

# Logical equivalence

Two sentences are **logically equivalent** iff true in same models:

$a \equiv \beta$  if and only if  $a \models \beta$  and  $\beta \models a$

$(a \wedge \beta) \equiv (\beta \wedge a)$	commutativity of $\wedge$
$(a \vee \beta) \equiv (\beta \vee a)$	commutativity of $\vee$
$((a \wedge \beta) \wedge \gamma) \equiv (a \wedge (\beta \wedge \gamma))$	associativity of $\wedge$
$((a \vee \beta) \vee \gamma) \equiv (a \vee (\beta \vee \gamma))$	associativity of $\vee$
$\neg(\neg a) \equiv a$	double-negation elimination
$(a \Rightarrow \beta) \equiv (\neg \beta \Rightarrow \neg a)$	contraposition
$(a \Rightarrow \beta) \equiv (\neg a \vee \beta)$	implication elimination
$(a \Leftrightarrow \beta) \equiv ((a \Rightarrow \beta) \wedge (\beta \Rightarrow a))$	biconditional elimination
$\neg(a \wedge \beta) \equiv (\neg a \vee \neg \beta)$	De Morgan
$\neg(a \vee \beta) \equiv (\neg a \wedge \neg \beta)$	De Morgan
$(a \wedge (\beta \vee \gamma)) \equiv ((a \wedge \beta) \vee (a \wedge \gamma))$	distributivity of $\wedge$ over $\vee$
$(a \vee (\beta \wedge \gamma)) \equiv ((a \vee \beta) \wedge (a \vee \gamma))$	distributivity of $\vee$ over $\wedge$



## Validity and satisfiability

A sentence is **valid** if it is true in **all** models,

e.g., *True*,  $A \vee \neg A$ ,  $A \Rightarrow A$ ,  $(A \wedge (A \Rightarrow B)) \Rightarrow B$

Validity is connected to inference via the **Deduction Theorem**:

$KB \models a$  if and only if  $(KB \Rightarrow a)$  is valid

A sentence is **satisfiable** if it is true in **some** model

e.g.,  $A \vee B$ ,  $C$

A sentence is **unsatisfiable** if it is true in **no** models

e.g.,  $A \wedge \neg A$

Satisfiability is connected to inference via the following:

$KB \models a$  if and only if  $(KB \wedge \neg a)$  is unsatisfiable

i.e., prove  $a$  by *reductio ad absurdum*

# Proof methods

Proof methods divide into (roughly) two kinds:

## Application of inference rules

- Legitimate (sound) generation of new sentences from old
- **Proof** = a sequence of inference rule applications
  - Can use inference rules as operators in a standard search alg.
- Typically require translation of sentences into a **normal form**

## Model checking

truth table enumeration (always exponential in  $n$ )  
improved backtracking, e.g., Davis–Putnam–Logemann–Loveland  
heuristic search in model space (sound but incomplete)  
e.g., min-conflicts-like hill-climbing algorithms

# Resolution

Conjunctive Normal Form (CNF—universal)

conjunction of disjunctions of literals

causes

E.g.,  $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$


Resolution inference rule (for CNF): complete for propositional logic

$$\frac{\cancel{J} \vee \dots \vee \cancel{J}, \quad m_1 \vee \dots \vee m_n}{\cancel{J} \vee \dots \vee \cancel{J_{i-1}} \vee \cancel{J_{i+1}} \vee \dots \vee \cancel{J_k} \vee m_1 \vee \dots \vee m_{j-1} \vee m_{j+1} \vee \dots \vee m_n}$$

where  $\cancel{J}$  and  $m_j$  are complementary literals. E.g.,

$$\frac{P_{1,3} \vee P_{2,2}, \quad \neg P_{2,2}}{P_{1,3}}$$

Resolution is sound and complete for propositional logic



## Conversion to CNF

$$B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$$

1. Eliminate  $\Leftrightarrow$ , replacing  $a \Leftrightarrow b$  with  $(a \Rightarrow b) \wedge (b \Rightarrow a)$ .

$$(B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})) \wedge ((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1})$$

2. Eliminate  $\Rightarrow$ , replacing  $a \Rightarrow b$  with  $\neg a \vee b$ .

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1})$$

3. Move  $\neg$  inwards using de Morgan's rules and double-negation:

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge ((\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1})$$

4. Apply distributivity law ( $\vee$  over  $\wedge$ ) and flatten:

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$$

# Resolution algorithm

Proof by contradiction, i.e., show  $KB \wedge \neg a$  unsatisfiable

```

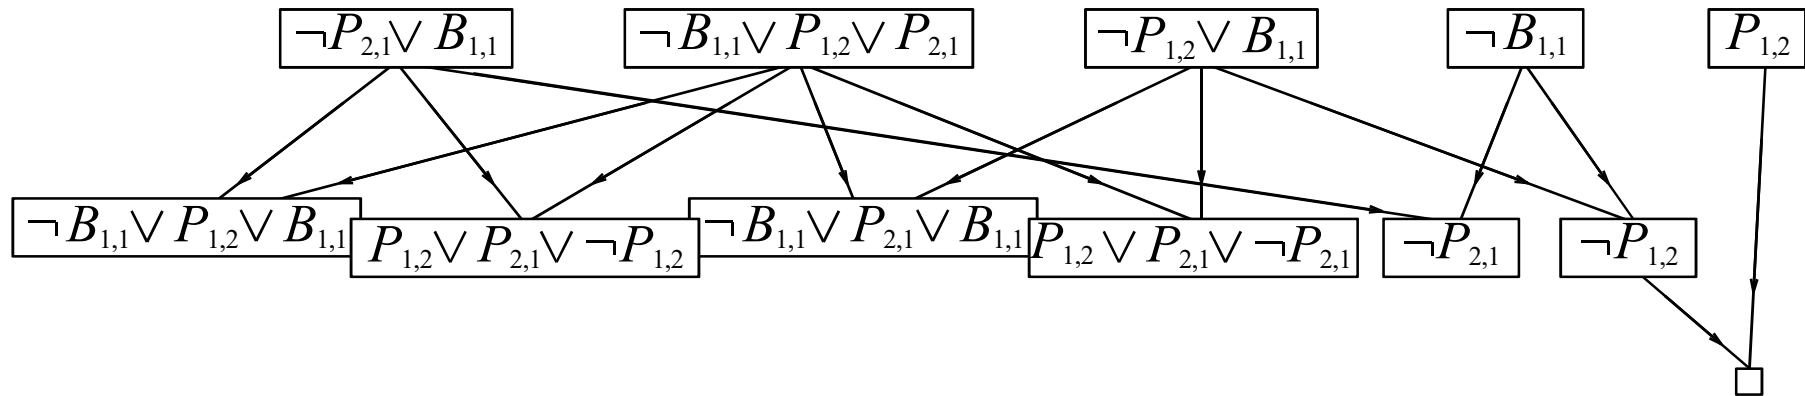
function PL-Resolution( $KB, a$ ) returns true or false
  inputs:  $KB$ , the knowledge base, a sentence in propositional logic
          $a$ , the query, a sentence in propositional logic

   $clauses \leftarrow$  the set of clauses in the CNF representation of  $KB \wedge \neg a$ 
   $new \leftarrow \{ \}$ 
  loop do
    for each  $C_i, C_j$  in  $clauses$  do
       $resolvents \leftarrow$  PL-Resolve( $C_i, C_j$ )
      if  $resolvents$  contains the empty clause then return true
       $new \leftarrow new \cup resolvents$ 
  if  $new \subseteq clauses$  then return false
   $clauses \leftarrow clauses \cup new$ 

```

# Resolution example

$$KB = (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1} \quad a = \neg P_{1,2}$$



# Forward and backward chaining

Horn Form (restricted)

KB = **conjunction** of **Horn clauses**

Horn clause =

- ◆ proposition symbol; or
- ◆ (conjunction of symbols)  $\Rightarrow$  symbol

E.g.,  $C \wedge (B \Rightarrow A) \wedge (C \wedge D \Rightarrow B)$

Modus Ponens (for Horn Form): complete for Horn KBs

$$\frac{a_1, \dots, a_n, \quad a_1 \wedge \dots \wedge a_n \Rightarrow \beta}{\beta}$$

Can be used with **forward chaining** or **backward chaining**.

These algorithms are very natural and run in **linear** time

# Forward chaining

Idea: fire any rule whose premises are satisfied in the *KB*,  
add its conclusion to the *KB*, until query is found

$$P \Rightarrow Q$$

$$L \wedge M \Rightarrow P$$


$$B \wedge L \Rightarrow M$$

$$A \wedge P \Rightarrow L$$

$$A \wedge B \Rightarrow L$$

*A*

*B*





# Forward chaining algorithm

```


function PL-FC-Entails?(KB, q) returns true or false
  inputs: KB, the knowledge base, a set of propositional Horn clauses
         q, the query, a proposition symbol
  local variables: count, a table, indexed by clause, initially the number of premises
                  inferred, a table, indexed by symbol, each entry initially false
                  agenda, a list of symbols, initially the symbols known in KB

  while agenda is not empty do
    p ← Pop(agenda)
    unless inferred[p] do
      inferred[p] ← true
      for each Horn clause c in whose premise p appears do
        decrement count[c]
        if count[c] = 0 then do
          if Head[c] = q then return true
          Push(Head[c], agenda)


  return false

```


# Forward chaining example




# Forward chaining example




# Forward chaining example




# Forward chaining example




# Forward chaining example



# Forward chaining example




# Forward chaining example





# Forward chaining example



## Proof of completeness

FC derives every atomic sentence that is entailed by  $KB$

1. FC reaches a **fixed point** where no new atomic sentences are derived
2. Consider the final state as a model  $m$ , assigning true/false to symbols

3. Every clause in the original  $KB$  is true in  $m$

**Proof:** Suppose a clause  $a_1 \wedge \dots \wedge a_k \Rightarrow b$  is false in  $m$

Then  $a_1 \wedge \dots \wedge a_k$  is true in  $m$  and  $b$  is false in  $m$

Therefore the algorithm has not reached a fixed point!

4. Hence  $m$  is a model of  $KB$

5. If  $KB \models q$ ,  $q$  is true in **every** model of  $KB$ , including  $m$

**General idea:** construct any model of  $KB$  by sound inference, check  $a$

# Backward chaining

Idea: work backwards from the query  $q$ :

to prove  $q$  by BC,

check if  $q$  is known already, or


prove by BC all premises of some rule concluding  $q$

Avoid loops: check if new subgoal is already on the goal stack


Avoid repeated work: check if new subgoal

- 1) has already been proved true, or
- 2) has already failed


# Backward chaining example




# Backward chaining example




# Backward chaining example



# Backward chaining example




# Backward chaining example







# Backward chaining example




# Backward chaining example




# Backward chaining example




# Backward chaining example



# Backward chaining example



# Backward chaining example



## Forward vs. backward chaining

FC is **data-driven**, cf. automatic, unconscious processing,  
e.g., object recognition, routine decisions

May do lots of work that is irrelevant to the goal

BC is **goal-driven**, appropriate for problem-solving,  
e.g., Where are my keys? How do I get into a PhD program?

Complexity of BC can be **much less** than linear in size of KB

# Effective Propositional Model Checking

Davis–Putnam algorithm with three improvements over TT-ENTAILS

- Early termination: detect T/F
- Pure symbol heuristic: same sign in all clauses
- Unit Clause heuristic: clause with one literal

Local search algorithms such as hill-climbing & simulated annealing.

In recent years, there has been a great deal of experimentation to find a good balance between greediness and randomness.

WALKSAT: On every iteration, the algorithm picks an unsatisfied clause and picks a symbol in the clause to flip.



## Summary

Logical agents apply **inference** to a **knowledge base** to derive new information and make decisions

Basic concepts of logic:

- **syntax**: formal structure of **sentences**
- **semantics**: **truth** of sentences wrt **models**
- **entailment**: necessary truth of one sentence given another
- **inference**: deriving sentences from other sentences
- **soundness**: derivations produce only entailed sentences
- **completeness**: derivations can produce all entailed sentences

Wumpus world requires the ability to represent partial and negated information, reason by cases, etc.

Forward, backward chaining are linear-time, complete for Horn clauses  
Resolution is complete for propositional logic. Propositional logic lacks expressive power

Local search methods (WALKSAT) find solutions (sound but not complete).